

Sourcing under Sanctions: Judicial Urgency and Pharmaceutical Procurement Costs^{*}

Darcio Genicolo-Martins^{a,*}, Paulo Furquim de Azevedo^a

^a*Inspere Institute of Education and Research, Rua Quatá 300, São Paulo, SP 04546-042, Brazil*

Abstract

Courts can force governments to deliver medicines, but legal urgency changes how the state buys them. We study 479,330 pharmaceutical purchase-offer-item observations from São Paulo’s electronic procurement platform during 2009–2019. Urgent procurement is 5.4% costlier than ordinary procurement, attracts 5.4% fewer bidders, and succeeds 2.1 pp more often. To separate pricing from sourcing, we compare court-mandated urgent purchases with an administrative urgent channel that lacks judicial sanctions. Because administrative cases are selected and larger, we combine selection bounds, within firm-buyer-item comparisons, and sourcing tests. Manski-Lee bounds place the litigated-over-administrative price gap at [15.9%, 21.1%], compared with a naive 29.5% gap. Yet the same firm selling the same item to the same buyer shows no broad markup in deep urgent markets ($\hat{\beta} = 0.035$, SE 0.041). The markup channel reappears in thinner and earlier subsamples. The main margin is sourcing: administrative orders are 3.3 times larger, winner sets overlap weakly (mean Jaccard 0.268), and modal winners differ in 70.2% of item-buyer pairs. One-sided sanctions secure delivery, but they impair demand aggregation and supplier matching.

Keywords: public procurement, health litigation, accountability, sourcing

JEL: D44, D73, H51, H57, I18, K41

1. Introduction

Courts can make governments deliver public goods and services. That power matters when administrative delay would otherwise deny treatment, income support, schooling, or other statutory entitlements (Glaeser and Shleifer, 2003; Ferraz, 2009; Wang, 2015). But compelled delivery

^{*}Short paper. This version: May 2026. All errors are our own.

^{*}Corresponding author.

Email addresses: darciogm1@insper.edu.br (Darcio Genicolo-Martins), paulofa1@insper.edu.br (Paulo Furquim de Azevedo)

also changes the production problem faced by the state. A procurement office that must deliver a medicine under judicial deadline and personal sanctions cannot always wait for aggregation, batching, or regular supplier competition. The question is therefore not whether delivery has value. The question is which procurement margin the state gives up when delivery is enforced under sanctions.

Higher prices under judicial urgency are empirically ambiguous. They may reflect same-firm markups: incumbent suppliers observe a sanctioned buyer and charge a delivery-risk premium. They may instead reflect sourcing: the state buys smaller lots, loses scale, and matches with a different set of suppliers. This distinction connects accountability models of bureaucratic pressure with evidence on passive waste and procurement capacity ([Prendergast, 2007](#); [Bandiera et al., 2009](#); [Best et al., 2023](#)). Standard procurement data often cannot distinguish these margins because prices, quantities, suppliers, and buyer urgency move together. A high price paid under litigation can look like a markup even when no supplier charges more conditional on selling the same item to the same buyer.

This paper separates pricing from sourcing in São Paulo’s pharmaceutical procurement system. The state buys medicines on BEC, an electronic procurement platform. Right-to-health litigation creates urgent court-mandated purchases in which judges order delivery and may impose personal sanctions on officials for noncompliance. Tender notices can reveal the court-order status of the purchase, making sanction exposure observable to suppliers. These purchases therefore combine legal urgency with one-sided accountability: officials face sanctions for nondelivery, while suppliers do not face an equivalent sanction for refusing to supply.

The setting also contains a comparison rarely available in procurement data. Some patients obtain medicines through an administrative-request channel rather than through litigation. These administrative purchases are urgent and selected; they are not random controls. Their value is narrower and more useful: they provide the closest feasible urgent-procurement comparison because they share the main execution constraints of urgent procurement but lack judicial sanctions. Since administrative purchases differ in selection and scale, the design does not rest on a single comparison. We bound selection, estimate same-firm pricing within firm-buyer-item triples, and measure supplier-set reallocation directly.

We begin with the broader urgent-procurement margin. Relative to ordinary purchases of comparable items, urgent purchases have 5.4% higher negotiated prices, 5.4% fewer bidders, and a

2.1 pp higher tender-success rate. These estimates motivate the mechanism analysis rather than carry the paper’s main causal claim. They show the basic tradeoff: urgency is associated with higher cost and thinner competition, while delivery succeeds more often.

Within urgent procurement, court-mandated purchases are costlier than administrative urgent purchases. The naive litigated-over-administrative gap is 29.5%. Because the administrative channel screens cases, we use Manski-Lee trimming bounds. The selection-corrected gap lies between 15.9% and 21.1%. Wild cluster inference supports the preferred under-the-gun estimate ($p = 0.0080$). The bounds are deliberately modest: they do not claim random assignment to sanction exposure, and they leave the administrative selection process visible.

The central pricing test asks whether the same firm charges a sanction-exposed buyer more for the same item. In firm-buyer-item triples observed under both urgent regimes, the administrative coefficient is 0.035 with standard error 0.041 on 4,573 observations. That estimate is statistically indistinguishable from zero. The null is strongest in deeper markets: above-median quantity purchases and SUS-formulary items yield coefficients of -0.005 and -0.001. The null is not universal. In below-median quantity purchases and in the earlier period, the coefficients are +0.066 and +0.117, respectively. In deep repeated urgent markets, the sanction-related cost margin does not appear as a broad same-firm markup; in thinner markets, supplier leverage can reappear.

The sourcing evidence explains where the cost margin operates. Administrative urgent orders are 3.3 times larger than litigated urgent orders, creating a mechanical bulk-discount component. The decomposition in Figure 1 separates the observed administrative-minus-litigated price gap from the quantity component, the within-firm component, and the residual composition component. The residual is not, by itself, proof of sourcing. Direct winner-switching evidence supports that interpretation: among 2,134 item-buyer pairs observed under both regimes, mean winner-set Jaccard similarity is 0.268, 48.5% have no winner overlap, and the modal winner differs in 70.2% of pairs.

The contribution is fourfold. First, the paper separates same-firm pricing from supplier-set sourcing in a setting where the two margins are observationally equivalent in standard procurement comparisons. Second, it opens the procurement mechanism behind right-to-health litigation: the fiscal cost is not only a court-versus-bureaucracy problem, but a sourcing problem created by legal urgency. Third, it documents a judicial-enforcement form of passive waste. One-sided sanctions secure delivery but weaken the routines through which public buyers aggregate demand and match

with suppliers. Fourth, the policy implication follows from the mechanism. The relevant margin is not only contract design or price caps; it is preserving demand aggregation and sourcing efficiency under legal urgency.

The paper proceeds as follows. Section 2 describes the legal and procurement setting. Section 3 explains the data, classifier, and samples. Section 4 sets out the empirical comparisons. Sections 5 and 6 report the estimates and robustness. Section 7 discusses interpretation and limits.

2. Institutional Background

2.1. Right-to-health litigation and one-sided sanctions

Brazilian courts can order state health agencies to provide medicines to individual plaintiffs. In São Paulo, these orders often arrive with short delivery deadlines and penalties for noncompliance. The sanctions are one-sided in the procurement sense. Public officials must deliver the medicine or face legal consequences; suppliers can decide whether to bid, at what price, and under which delivery conditions. This asymmetry matters because the public buyer must meet a hard delivery obligation in a market where supplier participation remains voluntary.

The court order is not hidden from the market. Procurement notices may reference the judicial order or the administrative process generated by it. Bidders can therefore infer when an order is court-mandated. The empirical question is what this information changes. It may allow suppliers to charge more to a sanctioned buyer. It may also force the buyer into smaller and less coordinated sourcing episodes.

2.2. Pharmaceutical procurement on BEC

São Paulo’s Department of Health buys medicines through BEC, the state’s electronic procurement platform. The platform records item identifiers, reference prices, negotiated prices, quantities, buyers, participating firms, winners, procurement modality, and tender outcomes. These records allow us to follow ordinary purchases, administrative urgent purchases, and litigated urgent purchases using the same procurement infrastructure.

Ordinary procurement can aggregate recurring demand and use planned purchasing routines. Urgent procurement compresses that timeline. The distinction is central to the paper: litigation

does not merely add a legal label to an otherwise identical purchase. It changes the time available to consolidate demand and can change which suppliers find the tender attractive.

2.3. The administrative-request channel

The administrative-request channel processes some patient demands outside court. It is not random assignment away from litigation. Administrative cases pass through screening and differ in scale, item mix, and expected feasibility. For this reason, administrative urgent purchases are not treated as a clean counterfactual.

They are nevertheless informative. They share the urgency of individualized medicine delivery but do not carry judicial sanctions against officials. The comparison between administrative and litigated urgent purchases therefore isolates the closest available institutional contrast around sanction exposure, while requiring explicit treatment of selection and scale. The empirical design uses that contrast only together with selection bounds, within-firm pricing tests, and direct sourcing measures.

3. Data, Classification, and Samples

3.1. Procurement and litigation data

The procurement data cover 479,330 pharmaceutical purchase-offer-item observations in BEC for 2009–2019. The raw BEC file contains 51,059 purchase orders, 33,144 item identifiers, and 100 observed purchasing units. We link these records to SES/SP litigation and administrative-process information to classify purchases into ordinary, administrative urgent, and litigated urgent regimes.

Price regressions use winning bids because transaction prices are observed for accepted offers. The main winners-only price sample contains 196,883 observations. Other outcomes, such as tender success, use the relevant purchase-offer-item or tender-level sample described in each table. [Table 1](#) reconciles the samples used in the paper.

3.2. Classification of purchase regimes

Purchase regimes are classified using structured links and tender-notice text. The classifier combines machine-learning predictions, JSON fields, regex markers, and position-based checks. It

uses 18 judicial patterns, 12 administrative patterns, and 10 boilerplate filters. The voting threshold is 2.

Classifier validation is complete. Against 179,148 ground-truth purchase orders, the classifier reaches 98.6% exact agreement. Per-class F1 is 0.93 for judicial purchases and 0.96 for administrative purchases; the macro-F1 over those two urgent classes is 0.94. The held-out machine-learning component yields F1 scores of 0.93 and 0.94. Online Appendix A reports the validation table and confusion counts.

Residual classification error should attenuate the central urgent-regime contrasts if it moves purchases across administrative and litigated categories without being correlated with unobserved prices conditional on the controls. The main threat is systematic misclassification of one regime into the other among high-cost items. The ground-truth validation reduces that concern, but the interpretation remains tied to the classified regimes rather than to an unobserved legal concept.

3.3. Samples and identifying variation

The analysis uses several nested and partially overlapping samples. The full BEC pharmaceutical sample establishes coverage. The analysis sample applies the core item and variable restrictions. The winners-only sample is used for negotiated-price regressions. The urgent panel restricts to administrative and litigated urgent purchases. The both-regime urgent sample contains cells or pairs observed under both administrative and litigated urgent procurement. The firm-buyer-item triple sample is the most restrictive pricing sample: it requires the same firm to sell the same item to the same buyer under both urgent regimes.

Table 1: Samples and identifying variation.

Sample	Size	Role in the paper
Full BEC pharmaceutical sample	479,330 POI observations	Universe of classified pharmaceutical procurement records.
Analysis sample	226,305 observations	Core item and variable restrictions for ordinary-versus-urgent comparisons.
Winners-only price sample	196,883 winning bids	Negotiated-price regressions use accepted winning bids.
Urgent panel	56,803 winning bids	Administrative-versus-litigated urgent comparison.
Both-regime urgent cells	5,581 cells	Cells with both administrative and litigated urgent purchases.
Both-types singleton cells	11,972 cells	Urgent cells with only one observed regime; used for representativeness checks.
Firm-buyer-item triple sample	4,573 observations; 1,206 triples	Same firm, same buyer, same item pricing test.

Notes: POI denotes purchase-offer-item. “Cells” follow the cell definition used by the corresponding generated representativeness and bounds scripts. The triple sample is a strict subset of urgent winning bids and identifies within-supplier pricing, not supplier reallocation.

3.4. Descriptive facts

The descriptive means already show why sample labels matter. Winner observations include 343,393 ordinary, 16,985 administrative, and 46,212 litigated purchases. Average quantities are much larger in the administrative urgent channel than in the litigated channel, while litigated urgent purchases have higher average log negotiated prices. These differences motivate the explicit decomposition of scale and sourcing rather than a simple administrative-versus-litigated comparison.

4. Empirical Framework

4.1. Separating pricing and sourcing margins

Let a purchase price reflect both the price charged by the winning firm and the process that determines which firm wins. A court order can raise costs through a same-firm markup, through reduced scale, through lower participation, or through a different supplier match. The empirical framework separates these channels with three objects: urgent-versus-ordinary estimates, administrative-versus-litigated urgent estimates, and within firm-buyer-item comparisons.

The within-firm test is the pricing test. It estimates whether the same supplier charges a different price to the same buyer for the same item when the purchase is administrative rather than litigated. Because this comparison conditions on the winning firm, it cannot measure supplier reallocation. That is the point: if the overall gap remains while the within-firm gap disappears, the cost margin is not a broad same-firm markup.

4.2. *Urgent-vs-ordinary procurement*

The first specification compares urgent and ordinary purchases:

$$y_{igbt} = \beta U_{igbt} + \alpha_i + \gamma_t + \delta_b + \varepsilon_{igbt}, \quad (1)$$

where y_{igbt} is a procurement outcome for item i , supplier or purchase observation g , buyer b , and year t ; U_{igbt} indicates urgent procurement; and the fixed effects absorb item, year, and buyer differences. This estimate is descriptive of the urgent cost margin, not the core sanction estimate.

4.3. *Administrative-vs-litigated urgent procurement*

Within urgent procurement, we estimate

$$y_{igbt} = \theta A_{igbt} + \alpha_i + \gamma_t + \delta_b + \varepsilon_{igbt}, \quad (2)$$

where A_{igbt} equals one for administrative urgent purchases and zero for litigated urgent purchases. Negative estimates in price regressions mean administrative purchases are cheaper, or equivalently that litigated purchases are more expensive. Because administrative purchases are selected and larger, θ is not interpreted alone. The paper reports selection bounds and decomposes the scale channel.

4.4. *Selection bounds*

Administrative requests are screened by the health agency. If screened cases are more feasible or less costly, naive comparisons overstate the sanction-related gap. We therefore use Manski-Lee trimming bounds. In strata where administrative observations exceed litigated observations, we trim the administrative price distribution by the differential selection proportion. Trimming the

high tail and low tail gives bounds on the litigated-over-administrative gap. The preferred bounds are reported as percentage price gaps to avoid sign ambiguity.

The bounds require monotone selection into the administrative channel within trimming strata. This is weaker than random assignment but still substantive: it assumes that the screened administrative cases form an ordered subset relative to the potential litigated cases within the stratum. Online Appendix C gives the formal trimming logic and reports the trimming proportions.

5. Results

5.1. Urgent procurement is costlier and less competitive

Urgent procurement differs from ordinary procurement on the three margins expected in a compressed sourcing process. Table 2 reports the compact estimates. Negotiated prices are higher by 5.4%, reference prices are higher by 2.7%, and the number of bidding firms falls by 5.4%. Tender success rises by 2.1 pp. The success result is consistent with urgency reallocating administrative effort toward completion, even as competition weakens.

Table 2: Urgent procurement and procurement outcomes.

Outcome	Effect	SE	Interpretation
Log negotiated price	0.053	0.016	5.4% higher prices
Log reference price	0.027	0.014	2.7% higher reference prices
Log number of bidding firms	-0.056	0.014	5.4% fewer bidders
Tender success	0.021	0.006	2.1 pp higher success

Notes: Compact presentation of existing urgent-versus-ordinary estimates under the preferred item, year, and PBU fixed-effects specification. Price estimates use winning bids. The table motivates the mechanism analysis and is not the core sanction-exposure design.

5.2. The under-the-gun gap is selection-bounded

Table 3 compares the naive under-the-gun estimate with Manski-Lee bounds. The coefficient is reported as the administrative-minus-litigated log price coefficient, while the percentage column reports the litigated-over-administrative price gap. The naive gap is 29.5%. After trimming for selection, the gap is bounded between 15.9% and 21.1%.

The bound is the right object for the main paper because the administrative channel is selected. A parametric Heckman correction is not a credible substitute here: the available exclusion is weak,

Table 3: Manski-Lee bounds on the under-the-gun coefficient.

Specification	log negotiated price		implied		N
	Coef. on Admin	SE	Lit-vs-Admin (%)		
Naive UTG (item + year + PBU FE)	-0.259	0.092	29.530		61620
Lee lower bound (admin top-trim)	-0.192	0.095	15.924		45624
Lee upper bound (admin bottom-trim)	-0.148	0.097	21.123		45624

Trimming applied within $item \times year \times PBU$ strata where the admin sample exceeds the litigated count, by the differential proportion. Lower bound trims the top of the admin price distribution; upper bound trims the bottom. Standard errors clustered at PBU level.

the fixed effects absorb much of the useful variation, and the resulting correction is non-informative. Online Appendix C reports it as a diagnostic rather than as a preferred estimate.

5.3. Same firm, same buyer, same item: no broad markup in deep markets

Table 4 reports the within firm-buyer-item test. In the baseline triple sample, the administrative coefficient is 0.035 with standard error 0.041. This is a precise test of broad same-firm markups in the observed deep urgent-market sample: the same firm does not charge a higher price to the sanction-exposed buyer once the firm, buyer, and item are held fixed.

Table 4: Within firm-buyer-item null: robustness across subsamples.

Subsample	$\hat{\beta}_{\text{Admin}}$	SE	N
All triples (baseline)	0.035	0.041	4,573
Above-median quantity	-0.005	0.044	2,188
Below-median quantity	0.066***	0.025	2,027
SUS-formulary items	-0.001	0.026	2,540
Non-formulary items	0.101	0.064	2,033
Earlier period	0.117***	0.037	2,553
Later period	-0.038	0.052	1,492

Notes: Each row reports the within firm-buyer-item triple Admin coefficient on log negotiated price, with FBI and year fixed effects, PBU-clustered SEs. Subsamples are constructed from the triple sample (1,206 triples). Significance: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

The heterogeneity matters. The null holds in above-median quantity purchases and SUS-formulary items, where repeated demand and broader supplier bases make sourcing deeper. The coefficient becomes positive in below-median quantity purchases and in the earlier period. The interpretation is therefore not that markups never occur. It is that in deep repeated urgent markets, the sanction-related cost margin does not appear as a broad same-firm markup. In thinner markets, supplier leverage can reappear.

5.4. Fragmented sourcing: scale loss and supplier reallocation

The remaining evidence points to sourcing. Administrative urgent orders are 3.3 times larger than litigated urgent orders, and the bulk-discount elasticity is -0.329. Figure 1 shows the pricing-versus-sourcing reconciliation. The observed administrative-minus-litigated gap is decomposed into the mechanical quantity component, the within-firm component, and a residual supplier-composition component.

Table 5 reports that direct evidence. Among item-buyer pairs observed under both urgent regimes, mean winner-set Jaccard similarity is 0.268. 48.5% of pairs have no winner overlap, and




Figure 1: Pricing versus sourcing decomposition. The figure separates the observed administrative-minus-litigated log price gap from the mechanical quantity component, the within firm-buyer-item component, and the residual composition component. The residual should not be read alone as proof of sourcing; Table 5 provides the direct winner-switching evidence.

the modal winner differs in 70.2% of pairs. Conditional on the same firm, the markup is absent in deep markets; unconditionally, the winning supplier often changes. That combination is the mechanism.

Table 5: Winner switching across regimes for the same item-buyer pair.

Item-buyer pairs (both regimes)	2,134
Mean distinct winners, admin	2.05
Mean distinct winners, litigated	1.99
Mean Jaccard similarity, winner sets	0.268
Pairs with any winner overlap (%)	51.5
Pairs with NO winner overlap (%)	48.5
Pairs with same modal winner (%)	29.8

Notes: Sample of buyer×item pairs with at least one administrative and one litigated urgent purchase. Jaccard similarity is the ratio of the cardinality of the intersection to the cardinality of the union of the two regimes’ winning-firm sets. “Modal winner” is the most frequent winning firm within each regime for a given pair. Source: BEC-SP, 2009–2019.

6. Falsification and Robustness

6.1. Placebo on never-litigated items

The placebo uses items never subject to litigation. If the urgent-price premium reflected generic time trends or the procurement platform alone, the placebo should reproduce the main effect. It does not. The negotiated-price placebo coefficient is -0.020 with standard error 0.032.

Table 6: Placebo Test: Items Never Subject to Litigation

	bid_price_log		bid_price_ref_log	
	(1)	(2)	(3)	(4)
Urgent Purchase	-0.0204 (0.0316)	0.0513*** (0.0153)	-0.0306 (0.0447)	0.0305** (0.0141)
Observations	39,283	196,883	46,874	226,278
R ²	0.83364	0.86796	0.81121	0.85183
Within R ²	7.03×10^{-6}	0.00022	1.73×10^{-5}	7.43×10^{-5}
item_id fixed effects	✓	✓	✓	✓
year_n fixed effects	✓	✓	✓	✓
pbu_id fixed effects	✓	✓	✓	✓

Placebo sample: items with zero litigated purchases across 2009-2019. Main sample: items with at least one litigated and one ordinary purchase. DV: log price. Item + Year + PBU FE. SE clustered at PBU level.

6.2. Wild cluster inference

Inference is clustered at the purchasing-unit level in the main specifications. For the under-the-gun contrast, a Rademacher wild cluster bootstrap with 999 replications gives $p = 0.0080$ for

the preferred specification and a confidence interval of $[-0.469, -0.039]$ in log points. The tighter item-by-year-month specification gives $p = 0.0390$.

Table 7: Wild cluster bootstrap inference on the UTG coefficient.

Specification	$\hat{\beta}$	Asy. SE	t	p_{boot}	95% CI (log-points)
Preferred (item + year + PBU)	-0.259	0.092	-2.80	0.008	$[-0.469, -0.039]$
Tightest (item $\times$ year-month + PBU)	-0.283	0.116	-2.44	0.039	$[-0.549, -0.001]$

Notes: Rademacher wild cluster bootstrap on the Admin coefficient with $B = 999$ replicates and restricted-null residuals; clustering at the PBU level. Negative coefficients mean litigated purchases are more expensive than administrative ones, in log-points. Source: `analysis/44_wild_bootstrap.R`.

6.3. Dynamic evidence as diagnostic

Dynamic evidence is supportive but not strong enough to carry identification by itself. The BJS event-study estimate at first exposure is 0.052 with standard error 0.018, and the estimate five periods after exposure is 0.147 with standard error 0.026. Honest-DiD sensitivity does not survive deviations at the observed maximum pre-period scale (no). We therefore treat the event study as a diagnostic for timing, not as the primary research design.

7. Policy Interpretation and Limits

7.1. What the evidence shows

The evidence shows a procurement mechanism. Court mandates secure delivery, but they can impair the sourcing margin through which the state procures the mandated good. In deep urgent pharmaceutical markets, the cost margin appears through fragmented sourcing, lost scale, and supplier-set reallocation rather than through a broad same-firm markup.

7.2. What the evidence does not show

The paper does not estimate patient health benefits, and it does not evaluate the social value of right-to-health litigation. It studies São Paulo pharmaceuticals during 2009–2019. The urgent-panel estimates tilt toward items observed often enough to support within-regime comparisons, and the firm-buyer-item triple sample is more selected still. The procurement-cost calculation in the appendix is not a full welfare estimate and should not be read as the value of eliminating sanctions.

7.3. Restoring aggregation under judicial urgency

The policy implication is to preserve delivery while restoring aggregation. Framework agreements, pooled urgent procurement, stronger administrative screening capacity, and pre-contracted suppliers for recurring medicines target the margin identified here. They do not ask courts to reduce access. They reduce the procurement penalty created when each order is sourced as a fragmented emergency.

8. Conclusion

Judicial enforcement changes public procurement through more than prices. In São Paulo's urgent pharmaceutical markets, court-mandated delivery raises costs relative to the closest administrative urgent channel, but the same supplier does not broadly charge more to the same buyer for the same item in deep markets. The gap instead appears through smaller orders and different winning suppliers.

That distinction matters for policy. If the problem were only same-firm markups, price caps or contract clauses would be the natural response. If the problem is sourcing, the institutional response is different: preserve legal delivery while rebuilding aggregation, pre-contracting, and supplier matching under urgency. The evidence supports the latter interpretation for deep repeated urgent markets, while leaving room for supplier leverage in thinner ones.

References

- Bandiera, O., Prat, A., Valletti, T., 2009. Active and passive waste in government spending: Evidence from a policy experiment. *American Economic Review* 99, 1278–1308. doi:[10.1257/aer.99.4.1278](https://doi.org/10.1257/aer.99.4.1278).
- Best, M.C., Hjort, J., Szakonyi, D., 2023. Individuals and organizations as sources of state effectiveness. *American Economic Review* 113, 2121–2167. doi:[10.1257/aer.20191598](https://doi.org/10.1257/aer.20191598).
- Ferraz, O.L.M., 2009. The right to health in the courts of Brazil: Worsening health inequities? *Health and Human Rights* 11, 33–45.
- Glaeser, E.L., Shleifer, A., 2003. The rise of the regulatory state. *Journal of Economic Literature* 41, 401–425. doi:[10.1257/002205103765762725](https://doi.org/10.1257/002205103765762725).

- Prendergast, C., 2007. The motivation and bias of bureaucrats. *American Economic Review* 97, 180–196. doi:[10.1257/aer.97.1.180](https://doi.org/10.1257/aer.97.1.180).
- Wang, D.W.L., 2015. Right to health litigation in Brazil: The problem and the institutional responses. *Human Rights Law Review* 15, 617–641. doi:[10.1093/hrlr/ngv025](https://doi.org/10.1093/hrlr/ngv025).