

Sourcing under Sanctions: Judicial Urgency and Pharmaceutical Procurement Costs^{*}

Darcio Genicolo-Martins^{a,*}, Paulo Furquim de Azevedo^a

^a*Inspere Institute of Education and Research, Rua Quatá 300, São Paulo, SP 04546-042, Brazil*

Abstract

Court mandates secure delivery, but they may change how governments procure the mandated goods. We study São Paulo pharmaceutical procurement on BEC, covering 479,330 purchase-offer-item observations during 2009–2019. Higher prices under judicial urgency can reflect incumbent suppliers charging same-firm markups or fragmented sourcing that changes scale and supplier matching. We use an administrative urgent channel as the closest feasible comparison: it lacks judicial sanctions but is selected and larger, so we combine Lee selection bounds, within firm-buyer-item comparisons, and sourcing tests. The bounds place the litigated-over-administrative price gap between 15.9% and 21.1%. Within the same firm, buyer, and item, the administrative coefficient is 0.035 (SE 0.041), providing little evidence of a broad same-firm markup in deep repeated urgent markets. The within-firm gap reappears only in the earlier period, while order-size differences operate through scale. The sourcing margin is substantial: administrative orders are 3.3 times larger, and modal winners differ in 70.2% of item-buyer pairs. One-sided sanctions secure delivery, but they change how the state buys: the policy margin is preserving aggregation and supplier matching under legal urgency.

Keywords: public procurement, health litigation, accountability, sourcing

JEL: D44, D73, H51, H57, I18, K41

1. Introduction

Court mandates can correct nondelivery while changing the state’s production problem. When a judge orders a public agency to provide a medicine, the relevant choice is no longer only whether

^{*}Short paper. This version: May 2026. All errors are our own.

^{*}Corresponding author.

Email addresses: darciogm1@insper.edu.br (Darcio Genicolo-Martins), paulofa1@insper.edu.br (Paulo Furquim de Azevedo)

the patient receives treatment. The agency must also source the good under legal urgency, with accountability concentrated on nondelivery. A procurement office facing short deadlines and one-sided sanctions cannot always wait for aggregation, batching, or regular supplier competition. The public-economics question is therefore which procurement margin the state gives up when delivery is compelled under sanctions.

Higher prices under judicial urgency have two distinct interpretations. They may reflect same-firm markups: suppliers see that a buyer is sanctioned and charge more for the same item. They may instead reflect sourcing reallocation: the state buys smaller lots, loses scale, attracts fewer bidders, or matches with different suppliers. Standard procurement comparisons have trouble separating these mechanisms because prices, quantities, suppliers, and urgency move together. A court-mandated purchase can look expensive even if no incumbent supplier raises its price conditional on selling the same item to the same buyer.

This distinction connects several literatures. Research on public procurement documents large price dispersion, passive waste, and the role of buyer capacity ([Bandiera et al., 2009](#); [Best et al., 2023](#); [Bosio et al., 2022](#); [Decarolis et al., 2020](#)); work on bureaucracy and accountability shows that public agents respond to one-sided pressure and distorted incentives ([Prendergast, 2007](#); [Bandiera et al., 2021](#)); and work on court efficiency studies how courts shape procurement performance ([Coviello et al., 2018](#)). A separate literature on right-to-health litigation studies access, priority setting, and the fiscal pressure of judicialized health claims ([Ferraz, 2009](#); [Wang, 2015](#); [Biehl et al., 2009](#)). We link these by opening a different court-procurement channel: court-mandated urgent delivery under one-sided compliance pressure. The object is not to restate the aggregate fiscal pressure of litigation, but to identify whether the procurement-cost margin comes from pricing by the same suppliers or from the sourcing process that legal urgency creates.

The setting is São Paulo pharmaceutical procurement. The state buys medicines through BEC, an electronic procurement platform that records items, prices, quantities, buyers, suppliers, modalities, and tender outcomes. Right-to-health litigation creates urgent court-mandated purchases. Court orders often impose short delivery deadlines and penalties for noncompliance; procurement notices may reveal the judicial order or the administrative process generated by it. The sanction is asymmetric in the procurement market. Officials must deliver or face legal consequences, while suppliers remain voluntary bidders who choose whether and how to participate.

São Paulo also has an administrative request channel. Some patients obtain medicines outside litigation, generating urgent procurement without judicial sanctions. This comparison is valuable not because it randomizes sanction exposure, but because it is the closest feasible urgent-procurement comparison. Administrative urgent purchases share the compressed execution environment, yet they are selected and larger. The empirical design therefore combines three pieces of evidence: Lee bounds for selection, within firm-buyer-item comparisons for same-firm pricing, and direct sourcing tests for scale and supplier-set reallocation.

We first document the broader urgent-procurement margin. Relative to ordinary purchases of comparable items, urgent purchases have 5.4% higher negotiated prices, 5.4% fewer bidders, and a 2.1 pp higher tender-success rate. These estimates establish the procurement environment rather than the paper’s sanction design: urgency raises costs and weakens competition even as the state completes tenders more often.

The under-the-gun comparison is within urgent procurement. Court-mandated purchases are costlier than administrative urgent purchases, with a naive litigated-over-administrative gap of 29.5%. Because administrative requests are screened and selected, we trim the overrepresented administrative group within item-year-PBU strata. The Lee bounds place the selection-corrected litigated-over-administrative gap between 15.9% and 21.1%, and wild cluster inference supports the preferred estimate ($p = 0.0080$). The design does not claim random assignment to sanctions. It asks how much of the observed gap remains after disciplining the most direct urgent comparison for selection.

The pricing test then asks whether the same firm charges a sanctioned buyer more for the same item. In firm-buyer-item triples observed under both urgent regimes, the administrative coefficient is 0.035 with standard error 0.041. Deep markets are repeated urgent settings with enough scale or standardized demand to support within-triple comparisons, mainly above-median quantity and SUS-formulary purchases. In those subsamples the same-firm markup is absent: the corresponding coefficients are -0.005 and -0.001. The result is not a claim that same-firm pricing is everywhere zero. In below-median quantity purchases and in the earlier period, the coefficients are +0.066 and +0.117. These margins differ in kind: the quantity dimension tracks scale rather than same-firm pricing, while the earlier-period gap is a genuine within-firm difference confined to the earlier part of the sample. In deep repeated urgent markets the sanction-related cost margin does not appear

as a broad same-firm markup; the residual within-firm gap is confined to the earlier period.

The sourcing evidence shows where the margin operates. Administrative urgent orders are 3.3 times larger than litigated urgent orders, so fragmented court-mandated buying gives up scale. Figure 1 decomposes the observed administrative-minus-litigated price gap into a mechanical quantity component, a within-firm component, and a residual composition component. The residual is not, by itself, proof of sourcing. The direct evidence is supplier reallocation: among 2,134 item-buyer pairs observed under both urgent regimes, mean winner-set Jaccard similarity is 0.268, 48.5% have no winner overlap, and the modal winner differs in 70.2% of pairs. Conditional on the same firm, the broad markup is weak in deep markets; unconditionally, the state often buys from a different supplier set.

The contribution is to identify the procurement mechanism behind a legal-enforcement margin. Standard procurement comparisons confound incumbent pricing with supplier-set reallocation; this paper separates them and shows that, in deep repeated urgent markets, court-mandated delivery raises costs through fragmented sourcing rather than a broad same-firm markup. It thereby reframes right-to-health litigation as a question of public-sector production under legal constraint—how court orders change the way the state buys, not only how much it spends. The pattern is a judicial-enforcement form of passive waste: one-sided sanctions secure delivery but weaken the routines that aggregate demand and match buyers with suppliers. The implication is not to weaken access to medicines, but to build procurement capacity that preserves aggregation and supplier matching under legal urgency.

The paper proceeds as follows. Section 2 describes the setting, Section 3 explains the data and samples, Section 4 sets out the empirical comparisons, Sections 5 and 6 report results and robustness, and Section 7 discusses interpretation and limits.

2. Institutional Background

2.1. Right-to-health litigation and one-sided sanctions

Brazil’s constitutional right to health gives courts a central role in disputes over access to medicines (Ferraz, 2009; Wang, 2015; CNJ/INSPER, 2019). In these cases, courts can order state health agencies to provide a specific medicine to an individual plaintiff. In São Paulo, such orders

often arrive with short delivery deadlines and may attach daily fines or other compliance measures, exposing the public administration or responsible officials to consequences for noncompliance. The institutional fact that matters for procurement is not the legal doctrine itself, but the hard delivery obligation it creates.

The sanctions are one-sided in the procurement sense. The public buyer must make the medicine available or face legal consequences; suppliers remain voluntary participants. They can decide whether to bid, which delivery conditions to accept, and what price to offer. Judicial urgency therefore shifts pressure onto the buyer without creating a symmetric obligation on the supply side. A procurement office under this constraint may have less time to consolidate orders, wait for recurring demand, or search for a broader supplier set.

Court-mandated status is also not necessarily hidden from the market. Procurement notices may reference the judicial order or the administrative process generated by it, and the classifier used in this paper exploits those text and structured-record markers. When the notice reveals the court-related status, bidders can infer that the purchase is tied to a judicial mandate. The empirical question is what this information and deadline pressure change: suppliers may charge more to a sanctioned buyer, or the buyer may instead be forced into smaller, less coordinated sourcing episodes.

2.2. Pharmaceutical procurement on BEC

São Paulo’s Department of Health buys medicines through BEC, the state’s electronic procurement platform. BEC records item identifiers, reference prices, negotiated prices, quantities, buyers, participating firms, winning suppliers, procurement modality, and tender outcomes. These records make the empirical design possible because ordinary purchases, administrative urgent purchases, and litigated urgent purchases are observed within the same procurement infrastructure.

The relevant transaction unit depends on the outcome. Price regressions use accepted winning bids, where negotiated prices are observed. Other outcomes use the purchase-offer-item or tender-level records appropriate to the table. Section 3 and Online Appendix A describe the classification and sample construction; the institutional point here is that BEC links the legal regime to item-level procurement behavior.

Ordinary pharmaceutical procurement can plan recurring demand, combine needs across re-

quests, and offer suppliers a predictable tender. Urgent procurement compresses that timeline: to deliver a specific medicine quickly, the buyer may procure a smaller lot, repeat purchases more often, or use a thinner supplier pool. Litigation therefore does not merely add a legal label to an otherwise identical purchase; it changes the time available to aggregate demand and the set of suppliers for whom the tender is attractive.

2.3. The administrative-request channel

The administrative-request channel processes some patient demands outside court. Patients can seek medicines through the health administration, and accepted requests can generate urgent, individualized procurement without a judicial order. These purchases share the operational pressure of individualized delivery, but they do not carry the court-sanction regime described above.

This channel is useful precisely because it is imperfect. It is not random assignment away from litigation. Administrative cases pass through screening and can differ from litigated cases in item mix, feasibility, and scale. The comparison therefore does not recover what a litigated purchase would have cost without sanctions; it is the closest feasible urgent-procurement comparison, creating urgent pharmaceutical procurement outside the court-sanction regime.

This institutional structure motivates the empirical strategy. Administrative urgent purchases share generic urgency but differ in selection and scale, so the design (Sections 3 and 4) classifies regimes within BEC and then handles selection and scale directly, rather than asking the administrative channel to do more than the institution can support.

3. Data, Classification, and Samples

3.1. Procurement and litigation data

The data cover São Paulo pharmaceutical procurement through BEC during 2009–2019. The empirical file contains 479,330 purchase-offer-item observations. BEC records item identifiers, buyers, suppliers, reference prices, negotiated prices, quantities, participating firms, winners, procurement modality, and tender outcomes. We link these procurement records to SES/SP litigation and administrative-process markers to classify purchases into ordinary, administrative urgent, and litigated urgent regimes.

The unit distinction is central. The classifier operates upstream at the purchase-order/tender-notice level. After regime labels are linked to BEC item records, the empirical analysis is conducted at the purchase-offer-item level. Price regressions use accepted winning bids because transaction prices are observed for accepted offers. Other outcomes, such as tender success, use the broader purchase-offer-item or tender-level sample specified in each table.

3.2. Classification of purchase regimes

Purchase regimes are classified using structured links and tender-notice text. The classifier combines machine-learning predictions, JSON fields, regex markers, and position-based checks. It uses 18 judicial patterns, 12 administrative patterns, and 10 boilerplate filters. The voting threshold is 2.

Classifier validation is complete and is conducted at the purchase-order/tender-notice level. Against 179,148 ground-truth purchase orders, the classifier reaches 98.6% exact agreement. Per-class F1 is 0.93 for judicial purchases and 0.96 for administrative purchases; the macro-F1 over those two urgent classes is 0.94. The held-out machine-learning component yields F1 scores of 0.93 and 0.94. Online Appendix A reports the validation table and classifier-error diagnostics.

Remaining misclassification would attenuate regime contrasts unless it is systematically correlated with residual prices within controls; the more serious threat is systematic misclassification between regimes among high-cost items. The validation evidence reduces that concern, though the interpretation remains tied to classified procurement regimes rather than to an unobserved legal concept.

3.3. Samples and identifying variation

The paper uses several samples because each empirical comparison asks a different question. The full BEC pharmaceutical file gives the linked POI-level coverage. The analysis sample applies core item and variable restrictions for ordinary-versus-urgent comparisons. The winners-only price sample contains 196,883 accepted winning bids and is used when negotiated price is the outcome.

Within urgent procurement, the urgent panel contains administrative and litigated winning bids and supports the under-the-gun comparison. Both-regime urgent cells are item-by-year-month cells containing both administrative and litigated purchases; singleton urgent cells contain only

one urgent regime and are used to assess representativeness. The firm-buyer-item triple sample is narrower by design. It requires the same supplier to sell the same item to the same buyer under both urgent regimes, so it identifies within-supplier pricing. It does not measure supplier-set reallocation, which requires the winner-switching evidence in Section 5. Online Appendix A reports the full sample-construction table alongside the classifier universe.

3.4. Descriptive facts

The descriptive facts show why raw administrative-versus-litigated means are not a design. Administrative urgent purchases are larger than litigated urgent purchases, while litigated purchases have higher average log negotiated prices. These differences combine selection, scale, and sanction exposure, so we do not interpret raw means causally. The design instead bounds selection, decomposes scale, and tests same-firm pricing and supplier-set reallocation directly.

4. Empirical Framework

4.1. Separating pricing and sourcing margins

Observed procurement prices combine the price charged by the winning supplier with the process that determines which supplier wins. Judicial urgency can therefore affect costs through at least three margins. First, the same supplier may charge a different price to the same buyer for the same item when the purchase is court-mandated. Second, urgent litigated purchases may be smaller and lose bulk discounts. Third, compressed timelines and fragmented demand may change participation and the winning supplier set.

Standard price regressions conflate these channels. The framework treats the observed urgent price gap as an accounting object with a scale component, a within-firm pricing component, and a supplier-composition component. This is not a structural model: the residual composition component is interpreted only together with direct winner-switching evidence. The within firm-buyer-item comparison is the same-firm pricing test. Because that test conditions on the winning supplier, it cannot measure supplier-set reallocation. Quantity controls are also channel diagnostics. Since quantity is itself a mechanism through which judicial urgency may affect prices, specifications that condition on quantity are not the preferred estimates of the total effect.

4.2. Urgent-vs-ordinary procurement

The first comparison estimates the broad urgent-procurement margin:

$$y_{oibt} = \beta \text{Urgent}_{oibt} + \alpha_i + \gamma_t + \delta_b + \varepsilon_{oibt}, \quad (1)$$

where o indexes the purchase-offer-item observation, i the item, b the buyer or purchasing unit, and t the year. The dependent variable is log negotiated price, log reference price, log number of bidders, or tender success. The fixed effects absorb item, time, and buyer differences.

Negotiated-price regressions use accepted winning bids; tender-success specifications use the broader outcome sample. Standard errors are clustered by purchasing unit unless otherwise stated. The coefficient β establishes the cost and competition margin associated with urgent procurement within items, years, and buyers. It does not isolate judicial sanction exposure.

4.3. Administrative-vs-litigated urgent procurement

The central contrast is within urgent procurement:

$$y_{oibt} = \theta \text{Admin}_{oibt} + \alpha_i + \gamma_t + \delta_b + \varepsilon_{oibt}, \quad (2)$$

where Admin_{oibt} equals one for administrative urgent purchases and zero for litigated urgent purchases. Throughout the under-the-gun analysis, coefficients are administrative minus litigated. In price regressions, a negative θ means litigated purchases are more expensive. Percentage gaps are reported in the reader-facing litigated-over-administrative direction.

The administrative channel is selected and larger, so θ alone is not a stand-alone sanction effect. We use the comparison as one part of the design: selection bounds discipline the urgent contrast, within firm-buyer-item tests isolate same-firm pricing, and sourcing measures document supplier-set reallocation.

For the same-firm pricing test, the estimating comparison can be written as

$$y_{ofibt} = \kappa \text{Admin}_{ofibt} + \mu_{fib} + \gamma_t + \varepsilon_{ofibt}, \quad (3)$$

where f indexes the supplier and μ_{fib} is a firm-by-item-by-buyer fixed effect. This specification

uses observations where the same firm sells the same item to the same buyer under both urgent regimes. It tests whether the same supplier prices differently across regimes. By construction, it cannot estimate the effect of changing which supplier wins.

4.4. Selection bounds

Administrative requests are screened before purchase. Screening likely admits cases that are administratively feasible, easier to source, or less costly. That selection can make administrative purchases cheaper even without judicial sanctions, so naive administrative-versus-litigated comparisons may overstate the sanction-related price gap.

We therefore use Lee trimming bounds, implemented within strata (Lee, 2009). The procedure trims the overrepresented administrative group within item $\times$ year $\times$ PBU strata. Trimming the high and low tails of the administrative price distribution gives lower and upper bounds for the litigated-over-administrative gap. Selection is bounded under a monotonicity restriction; it is not eliminated. Within a trimming stratum, screened administrative observations are assumed to form an ordered selected subset relative to the potential litigated comparison. This is weaker than exogenous assignment but still substantive, and the paper does not claim point identification of a counterfactual with sanctions removed.

The preferred object is therefore a bounded litigated-over-administrative price gap, combined with mechanism evidence. Online Appendix B reports trimming rates, alternative trimming strata, and the non-informative parametric selection diagnostic.

5. Results

5.1. Urgent procurement is costlier and less competitive

Urgent procurement differs from ordinary procurement in the direction expected when sourcing is compressed. Panel A of Table 1 reports the compact estimates. Negotiated prices are higher by 5.4%, reference prices are higher by 2.7%, the number of bidding firms falls by 5.4%, and tender success rises by 2.1 pp. Urgency appears to help the state complete purchases while weakening the competitive conditions under which those purchases are sourced. These estimates establish the urgent-procurement margin, not judicial sanction exposure; the remaining results move inside urgent procurement and separate pricing from sourcing.

Table 1: Urgent procurement margins and under-the-gun bounds.

Panel A. Urgent versus ordinary procurement

Outcome	Effect	SE	Interpretation
Log negotiated price	0.053	0.016	5.4% higher prices
Log reference price	0.027	0.014	2.7% higher reference prices
Log number of bidding firms	-0.056	0.014	5.4% fewer bidders
Tender success	0.021	0.006	2.1 pp higher success

Panel B. Under-the-gun: administrative vs. litigated urgent gap

Specification	Admin coef.	SE	Gap (%)	N
Naive UTG: item + year + PBU FE	-0.259	0.092	29.5%	61,620
Lee lower bound: admin top-tail trim	-0.192	0.095	15.9%	45,624
Lee upper bound: admin bottom-tail trim	-0.148	0.097	21.1%	45,624

Notes: POI denotes purchase-offer-item. *Panel A* reports urgent-versus-ordinary estimates under item, year, and PBU fixed effects; negotiated-price and reference-price estimates use accepted winning bids, tender success uses the broader tender/POI sample, and standard errors are clustered by PBU. *Panel A* establishes the urgent-procurement margin and is not the sanction-exposure design. *Panel B* reports the administrative-versus-litigated urgent gap; coefficients are administrative minus litigated log negotiated price, so negative coefficients mean litigated purchases are more expensive, while the percentage column is the reader-facing litigated-over-administrative gap. Lee trimming is applied within item $\times$ year $\times$ PBU strata where administrative observations exceed litigated observations; the mean trimming rate is 26.9% and the maximum is 100.0%.

5.2. The under-the-gun gap is selection-bounded

Panel B of Table 1 compares the naive under-the-gun estimate with Lee bounds. Negative coefficients mean litigated purchases are more expensive: the coefficient is administrative minus litigated log price, while the percentage column reports the same comparison in the reader-facing litigated-over-administrative direction. The naive gap is 29.5%. After trimming the selected administrative group, the bounded gap is 15.9% to 21.1%.

The bounded interval, not the naive coefficient, is the disciplined result. It does not eliminate selection; it asks how much of the administrative advantage could be explained by monotone screening within item-by-year-by-PBU strata. Even under that correction, litigated urgent procurement remains more expensive. A parametric Heckman correction is non-informative in this design and is noted only as a diagnostic in Online Appendix B. The next question is whether the gap reflects the same firms charging more under court pressure.

5.3. Same firm, same buyer, same item: no broad markup in deep markets

Table 2 reports the within firm-buyer-item test. The sample compares the same firm selling the same item to the same buyer under both urgent regimes. This isolates same-firm pricing, while intentionally conditioning away supplier reallocation. We use “deep markets” descriptively for repeated urgent settings with greater scale or more standardized demand, proxied by above-median quantity, SUS-formulary status, and later-period procurement. In the baseline triple sample, the administrative coefficient is 0.035 with standard error 0.041. Because the coefficient is administrative minus litigated, the same firm charges statistically indistinguishable prices across the two urgent regimes.

Table 2: Within firm-buyer-item null: robustness across subsamples.

Subsample	$\hat{\beta}_{\text{Admin}}$	SE	N
All triples (baseline)	0.035	0.041	4,573
Above-median quantity	-0.005	0.044	2,188
Below-median quantity	0.066***	0.025	2,027
SUS-formulary items	-0.001	0.026	2,540
Non-formulary items	0.101	0.064	2,033
Earlier period	0.117***	0.037	2,553
Later period	-0.038	0.052	1,492

Notes: Each row reports the within firm-buyer-item triple Admin coefficient on log negotiated price, with FBI and year fixed effects, PBU-clustered SEs. Coefficients are administrative minus litigated; negative values mean litigated purchases are more expensive within the same firm, buyer, and item, while values near zero indicate no detectable within-triple price difference. Subsamples are constructed from the triple sample (1,206 triples). Significance: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

The heterogeneity is as important as the baseline null. The coefficient remains near zero in above-median quantity purchases and SUS-formulary items, where repeated demand and broader supplier bases make sourcing deeper. It becomes positive in below-median quantity purchases and in the earlier period. These margins differ in kind. Holding firm, buyer, and item fixed, the within-triple log-quantity coefficient is -0.259 (SE 0.074): the quantity gradient is the bulk-discount channel, not same-firm pricing. The earlier-period gap, by contrast, is a genuine within-firm difference, confined to the earlier part of the sample. The result is not a claim that markups are absent everywhere. It is that in deep repeated urgent markets, the sanction-related cost margin does not appear as a broad same-firm markup; the residual within-firm gap is confined to the earlier

period.

5.4. Fragmented sourcing: scale loss and supplier reallocation

The remaining evidence points to sourcing. Administrative urgent orders are 3.3 times larger than litigated urgent orders, and the bulk-discount elasticity is -0.329. Fragmented court-mandated buying therefore gives up scale. Figure 1 reconciles the observed administrative-minus-litigated gap with a mechanical quantity component, the within-firm pricing component, and a residual supplier-composition component.




Figure 1: Pricing versus sourcing decomposition. Components are administrative minus litigated. The figure separates the observed price gap from the mechanical quantity component, the within firm-buyer-item pricing component, and a residual composition component. The residual is a reconciliation residual and should not be read alone as proof of sourcing; Table 3 provides direct winner-switching evidence.

Table 3 reports that direct evidence. Among item-buyer pairs observed under both urgent regimes, mean winner-set Jaccard similarity is 0.268. 48.5% of pairs have no winner overlap, and the modal winner differs in 70.2% of pairs. Conditional on the same firm, prices are statistically indistinguishable in deep markets; unconditionally, the winning supplier often changes. That combination is the empirical signature of fragmented sourcing.

Table 3: Winner switching across urgent regimes within item-buyer pairs.

Item-buyer pairs (both regimes)	2,134
Mean distinct winners, admin	2.05
Mean distinct winners, litigated	1.99
Mean Jaccard similarity, winner sets	0.268
Pairs with any winner overlap (%)	51.5
Pairs with NO winner overlap (%)	48.5
Pairs with same modal winner (%)	29.8
Pairs with different modal winner (%)	70.2

Notes: Unit is a buyer×item pair with at least one administrative and one litigated urgent purchase. Jaccard similarity is the ratio of the cardinality of the intersection to the cardinality of the union of the two regimes’ winning-firm sets. “Modal winner” is the most frequent winning firm within each regime for a given pair. Source: BEC-SP, 2009–2019.

6. Falsification and Robustness

6.1. Placebo on never-litigated items

The first check uses items never observed in litigation. If the main urgent-price pattern reflected generic procurement-platform dynamics or broad time trends, administrative urgency among never-litigated items should reproduce it. It does not: the negotiated-price placebo coefficient is -0.020 with standard error 0.032. This is a falsification check for that specific alternative explanation, not a stand-alone identification design. Online Appendix C reports the full placebo table.

6.2. Wild cluster inference

The main specifications cluster standard errors by purchasing unit. Because PBU clusters are few and uneven, Online Appendix C reports Rademacher wild-cluster inference for the under-the-gun contrast. The preferred specification rejects a zero gap ($p = 0.0080$), and the tighter item-by-year-month specification gives $p = 0.0390$. Wild bootstrap inference supports the preferred estimate, but it does not remove the selection concern handled by the bounds.

6.3. Dynamic evidence as diagnostic

The dynamic evidence is useful for timing, not for the primary identification strategy. We use the BJS event-study estimator and Honest-DiD sensitivity diagnostics (Borusyak et al., 2024; Rambachan and Roth, 2023). In the BJS event study, the first post-exposure estimate is 0.052

with standard error 0.018, and the estimate five periods after exposure is 0.147 with standard error 0.026. The first post-period estimate does not survive Honest-DiD deviations at the observed maximum pre-period scale. Because the dynamic design is sensitive to plausible deviations from parallel trends, we use it as diagnostic evidence. The main claims rest on the selection bounds, the within firm-buyer-item pricing test, and the sourcing evidence.

Additional appendix diagnostics bound the within-firm null, adjust the heterogeneity splits for multiple testing, and benchmark winner switching against within-regime supplier turnover.

7. Policy Interpretation and Limits

7.1. What the evidence shows

The evidence identifies a procurement mechanism, not a general verdict on courts. Court mandates secure delivery, but one-sided legal urgency can weaken the sourcing margin through which the state obtains the mandated medicine. In deep repeated urgent pharmaceutical markets, the procurement-cost margin appears through fragmented sourcing, lost scale, and supplier-set reallocation rather than a broad same-firm markup; a residual within-firm gap is confined to the earlier part of the sample. The mechanism is a judicial-enforcement form of passive waste: the state delivers, but it buys under conditions that make aggregation and supplier matching harder.

7.2. What the evidence does not show

The evidence does not estimate patient health benefits, the social value of right-to-health litigation, or a structural counterfactual in which sanctions are removed. The setting is São Paulo pharmaceutical procurement during 2009–2019. The urgent-panel estimates tilt toward medicines observed often enough to support administrative-versus-litigated comparisons, and the firm-buyer-item triple sample is more selected still. The appendix reports a fiscal procurement-cost calculation, not a full welfare estimate. It excludes health benefits, search costs, compliance benefits, and other welfare components, and it should not be read as the value of scaling back legal enforcement.

7.3. Restoring aggregation under judicial urgency

The policy margin is to preserve delivery while preventing legal urgency from becoming procurement fragmentation. Because the measured margin is lost scale and supplier matching rather than

contract language or price caps, the relevant instruments are those that rebuild aggregation under urgency. Framework agreements and pooled urgent procurement restore scale; pre-contracted suppliers for recurring medicines support supplier matching; stronger administrative screening widens an urgent channel without court sanctions for admissible cases; and inventories or protocols for recurrent litigated medicines let the state anticipate and aggregate demand rather than source isolated emergencies. Related evidence on hospital group purchasing shows why such aggregation can matter for procurement costs (Lin and Wang, 2025). The paper does not estimate these tools directly; it points to them because the mechanism it measures is aggregation and supplier matching.

8. Conclusion

What does the state give up when courts force urgent delivery? In São Paulo pharmaceutical procurement, the main loss in deep repeated urgent markets is not a broad same-firm markup. It is sourcing efficiency. Court-mandated purchases are costlier than the closest administrative urgent comparison, but the same supplier selling the same item to the same buyer does not broadly charge more under the litigated regime. The gap instead appears through smaller orders and different winning suppliers.

That distinction is the paper’s contribution. Standard procurement comparisons can confound incumbent pricing with supplier-set reallocation. Separating the two shows how one-sided legal urgency can create passive waste: delivery is secured, but the routines that aggregate demand and match public buyers with suppliers deteriorate. The result is not that same-firm pricing is everywhere zero: a residual within-firm gap persists in the earlier part of the sample.

The policy implication is to preserve delivery while rebuilding sourcing capacity under legal urgency: if the margin is sourcing, price caps or contract language alone miss the problem, while instruments that restore aggregation address it more directly. The results are specific to São Paulo pharmaceuticals during 2009–2019 and do not estimate patient benefits or full welfare. They identify the procurement margin that legal urgency puts under strain: not only how much the state pays, but how the state is forced to buy.

References

- Bandiera, O., Best, M.C., Khan, A.Q., Prat, A., 2021. The allocation of authority in organizations: A field experiment with bureaucrats. *Quarterly Journal of Economics* 136, 2195–2242. doi:[10.1093/qje/qjab029](https://doi.org/10.1093/qje/qjab029).
- Bandiera, O., Prat, A., Valletti, T., 2009. Active and passive waste in government spending: Evidence from a policy experiment. *American Economic Review* 99, 1278–1308. doi:[10.1257/aer.99.4.1278](https://doi.org/10.1257/aer.99.4.1278).
- Best, M.C., Hjort, J., Szakonyi, D., 2023. Individuals and organizations as sources of state effectiveness. *American Economic Review* 113, 2121–2167. doi:[10.1257/aer.20191598](https://doi.org/10.1257/aer.20191598).
- Biehl, J., Petryna, A., Gertner, A., Amon, J.J., Picon, P.D., 2009. Judicialisation of the right to health in Brazil. *The Lancet* 373, 2182–2184. doi:[10.1016/S0140-6736\(09\)61172-7](https://doi.org/10.1016/S0140-6736(09)61172-7).
- Borusyak, K., Jaravel, X., Spiess, J., 2024. Revisiting event-study designs: Robust and efficient estimation. *Review of Economic Studies* 91, 3253–3285. doi:[10.1093/restud/rdae007](https://doi.org/10.1093/restud/rdae007).
- Bosio, E., Djankov, S., Glaeser, E.L., Shleifer, A., 2022. Public procurement in law and practice. *American Economic Review* 112, 1091–1117. doi:[10.1257/aer.20200738](https://doi.org/10.1257/aer.20200738).
- CNJ/INSPER, 2019. Judicialização da saúde no Brasil: Perfil das demandas, causas e propostas de solução. Conselho Nacional de Justiça, Brasília.
- Coviello, D., Moretti, L., Spagnolo, G., Valbonesi, P., 2018. Court efficiency and procurement performance. *Scandinavian Journal of Economics* 120, 826–858. doi:[10.1111/sjoe.12225](https://doi.org/10.1111/sjoe.12225).
- Decarolis, F., Giuffrida, L.M., Iossa, E., Mollisi, V., Spagnolo, G., 2020. Bureaucratic competence and procurement outcomes. *Journal of Law, Economics, and Organization* 36, 537–597. doi:[10.1093/jleo/ewaa004](https://doi.org/10.1093/jleo/ewaa004).
- Ferraz, O.L.M., 2009. The right to health in the courts of Brazil: Worsening health inequities? *Health and Human Rights* 11, 33–45.
- Lee, D.S., 2009. Training, wages, and sample selection: Estimating sharp bounds on treatment effects. *Review of Economic Studies* 76, 1071–1102. doi:[10.1111/j.1467-937X.2009.00536.x](https://doi.org/10.1111/j.1467-937X.2009.00536.x).
- Lin, H., Wang, Y., 2025. The value of group purchasing: Evidence from the U.S. hospital industry. *Journal of Public Economics* 247, 105380. doi:[10.1016/j.jpubeco.2025.105380](https://doi.org/10.1016/j.jpubeco.2025.105380).
- Prendergast, C., 2007. The motivation and bias of bureaucrats. *American Economic Review* 97, 180–196. doi:[10.1257/aer.97.1.180](https://doi.org/10.1257/aer.97.1.180).
- Rambachan, A., Roth, J., 2023. A more credible approach to parallel trends. *Review of Economic Studies* 90, 2555–2591. doi:[10.1093/restud/rdad018](https://doi.org/10.1093/restud/rdad018).
- Wang, D.W.L., 2015. Right to health litigation in Brazil: The problem and the institutional responses. *Human Rights Law Review* 15, 617–641. doi:[10.1093/hrlr/ngv025](https://doi.org/10.1093/hrlr/ngv025).