Screening Under Incomplete Observability: Equilibrium Loser-Side Participation as an Award-Layer Signal of Cartel Activity

Darcio Genicolo-Martins[1,*], Paulo Furquim de Azevedo[1]

[1] *INSPER, São Paulo, Brazil*

Abstract

Cartel-screening methods in the empirical literature presuppose bid-level observability. In many enforcement environments only the contract-award layer survives, and the question of what collusion-relevant information remains under that coarsening is open. We study the screening value of endogenous loser-side participation. A separating-equilibrium framework identifies $\log(1 + \text{tenders_count})$ together with the wins $= 0$ filter as a sufficient ranking statistic for cover-bidder type given award-record data; the median $+ 1.5 \times$IQR binary rule is its information-coarsening. In São Paulo's Bolsa Eletrônica de Compras (2009–2019), 2,735 of 16,843 persistent zero-win firms cross this rule. Their presence associates with a $+3.6\% - +7.7\%$ higher conditional log negotiated unit price across estimators. Forcing overlap on observables reverses the sign (overlap-cell ATT -9.72%; propensity-score-trimmed -30.67%). We read the reversal as the signature of a screening-value object rather than a treatment-effect object: overlap restrictions strip the endogenous sorting that generates the screen's economic content. Against CADE's cartel records the construct reaches AUC 0.748 on a contemporaneous benchmark restricted to four pre-2020 cases and 0.864 under temporal holdout against cobidder labels; against direct defendants in the broader BEC firm universe AUC is 0.49. The exercise locates the construct as cartel-adjacency evidence under coarsened observability, separates an award-layer screening stage from a bid-layer forensic stage,

*Corresponding author.

Email address: darciogm1@insper.edu.br (Darcio Genicolo-Martins)

This version: May 2026. We thank seminar participants at INSPER for comments on earlier drafts; all remaining errors are our own.

and frames the heterogeneity across buyer size as variation in the detection regime rather than as identification of an institutional channel.

Keywords: screening under incomplete observability, cover bidding, cartel adjacency, award-layer enforcement, separating equilibrium, participation-based statistics
JEL No.: D44, D73, H57, K21, L41

1. Introduction

The empirical literature on cartel detection has converged on methods that read collusion off the distribution of submitted bids: variance, kurtosis, skewness, or normalized spread within a tender are informative about coordinated bidding when the bid-level record is preserved (Porter and Zona, 1993; Bajari and Ye, 2003; Imhof, 2019; Wallimann et al., 2023). Real enforcement environments rarely match this informational benchmark. Many electronic procurement systems archive only the contract-award layer—winner identity, participant identity, item code, and negotiated price—and export bid-by-bid records to operational logs that competition authorities and audit courts query slowly, partially, or not at all. The screening methods designed for the bid layer become inoperative in the layer that survives.

The question we ask is conceptual rather than operational: what collusion-relevant information remains when enforcement observes only the award layer? An answer cannot reproduce the bid-level moment battery on coarsened data. It must instead identify a statistic generated by the collusive arrangement itself that is preserved when bid amounts are dropped. We argue that endogenous loser-side participation is such a statistic. Cartels with cover bidders generate persistent zero-win participation footprints that are visible in award records, and the cumulative exposure of those footprints carries information about cover-bidder type that bid-distribution moments do not exhaust.

Theory. A separating-equilibrium framework with cover bidders and genuine entrants (Online Appendix A) makes this concrete. Two results anchor the empirical strategy. First, in equilibrium cover bidders satisfy $\Pr(\text{win} \mid C) = 0$: the wins $= 0$ filter is a

separating choice rather than a descriptive coincidence. Second, given award-record data, the log-exposure statistic $\log(1 + \text{tenders_count})$ is a sufficient ranking statistic for cover-bidder type under a monotone-likelihood-ratio Poisson approximation; the binary median $+ 1.5 \times$IQR rule is its information-coarsening. The framework treats endogenous sorting into cartel-affected environments as constitutive of the screening object, not as a confound to be matched away. This distinction governs the interpretation of every empirical exercise that follows.

Empirical evidence. We test the framework on São Paulo's Bolsa Eletrônica de Compras (BEC, 2009–2019), an electronic platform that preserves the award envelope for 1,654,401 tender-items across 1,308 purchasing units and two procurement modalities. Among the 41,444 active firms, 16,843 have a zero win rate across the full sample; 2,735 of them cross the participation threshold the framework identifies, and we call these *frequent losers*. The empirical content of the paper is structured by the screening logic above. The continuous loss-intensity statistic strictly dominates the binary rule in discrimination (DeLong et al. 1988 test, $p < 10^{-3}$), consistent with the framework's identification of $\log(1 + \text{tc})$ as the primitive and the binary rule as its operational coarsening. Frequent-loser presence associates with a $+3.6\% - +7.7\%$ higher conditional log negotiated unit price across four estimators (cross-fit, OLS, CEM, IPW) on the broad sample. Once the comparison is restricted to items where treated and control overlap on observables the conditional coefficient reverses sign: overlap-cell ATT -9.72%, propensity-score-trimmed -30.67%. We read the reversal as the empirical signature of a screening-value object rather than a treatment-effect object: the overlap restriction strips the endogenous sorting that, under the framework, generates the screen's economic content. The broad-sample estimate is not a defended causal effect; it is a measurement of the loser-side participation footprint where the cartel deploys it.

Validation and scope. Against CADE's procurement-cartel adjudication portfolio the construct discriminates 193 adjudicated cobidder firms inside the always-loser stratum at firm-level AUC 0.924 in-sample, 0.864 under temporal holdout (train 2009–2016, test 2017–2019), and 0.748 on a contemporaneous benchmark restricted to four cases

adjudicated before 2020 (95% DeLong CI $[0.713, 0.783]$). Cobidders are the structurally appropriate validation object under coarsened observability: the award-record envelope cannot adjudicate cartel membership, and participation alongside adjudicated cartel members is the label that the screening statistic is designed to recover. Against the 47 direct CADE defendants active in the broader BEC firm universe the construct discriminates at AUC 0.49. The asymmetry between cobidder and direct-defendant discrimination is the design's empirical signature: direct defendants are typically frequent winners and lie outside the construct's target stratum. The price imprint varies across procuring-unit-size quartiles by a factor of 12.6×, evidence of heterogeneity in the detection regime that we interpret as scope information rather than as identification of the institutional channel.

Contribution. The contribution is at the conjunction of three literatures. Bid-coordination tests (Bajari and Ye, 2003) require a researcher-supplied partition between colluders and a competitive fringe; the construct supplies a candidate *ex ante* firm-level flag computable from award records. Bid-distribution screens (Imhof, 2019; Wallimann et al., 2023) require bid microdata; the construct runs on the layer that survives the operational coarsening. Empirical work on cover bidding has been retrospective (Porter and Zona, 1999; Asker, 2010); the construct shifts cover-bidder identification into a prospective screening stage. The cumulative claim is structural: under incomplete observability the appropriate enforcement architecture separates an award-layer screening stage that produces candidate flags from a bid-layer forensic stage that adjudicates suspected coordination. The two stages answer different questions and require different statistics; the construct fills the first.

What this paper does and does not do. We do not claim a uniform treatment effect of cover bidding on prices, identification of an oversight channel through which buyer size moderates the participation footprint, or recovery of cartel membership at the firm level. We claim that the participation primitive identified by the framework carries collusion-relevant information that survives the coarsening from the bid layer to the award layer; that the screening-value interpretation makes the sign reversal under overlap restrictions diagnostic rather than confounding; and that the construct discriminates a

structurally appropriate cartel-adjacency object at rates well above random matching at comparable participation volume. The construct integrates with bid-distribution methods rather than replacing them (§9.4).

The remainder of the paper proceeds in five blocks. Section 2 situates the contribution against the bid-coordination, bid-distribution, and cover-bidding literatures, and against work on enforcement design under incomplete information. Section 3 describes the BEC data layers and the CADE adjudication portfolio. Sections 4–5 state the participation primitive, the binary operationalization, and the descriptive identifying stance, alongside the two design-based strategies (RDD, DiD) that return null. Sections 6–7 report the CADE validation and the within-item price association, with §7.1 treating the sign reversal under overlap restrictions as the screening-value-vs-treatment-effect distinction. Sections 8–9 develop the modal informativeness contrast, the buyer-size heterogeneity, the leakage and adversarial-adaptation audits, and the screening-vs-forensic-stage comparison against the Imhof–Wallimann pipeline. Sections 10–11 catalog scope limitations and conclude.

2. Related Literature

The paper sits at the conjunction of four literatures, organized by the data layer at which each operates. Bid-coordination tests read collusion off bid-function exchangeability conditional on a researcher-supplied partition between colluders and a competitive fringe. Bid-distribution screens read collusion off within-tender moments of submitted bids. Cover-bidding studies identify the loser side of collusive arrangements retrospectively, after enforcement has produced the partition. Work on enforcement design under incomplete information asks how detection regimes should be structured when monitoring is imperfect. The construct's positioning is determined by the informational layer at which each literature operates: the first two require bid-level observability that is often absent; the third requires the post-adjudication record; the fourth motivates the architectural separation between screening and forensic stages that the empirical exercise instantiates.

Bid-coordination tests presuppose a partition. Bajari and Ye (2003) establish the canonical framework for detecting collusion through exchangeability of bid functions and conditional independence of residuals, conditional on a researcher-supplied partition between suspected colluders and a competitive fringe. The papers that build on Bajari–Ye inherit the constraint: Porter and Zona (1993) and Porter and Zona (1999) detect bid rotation conditional on assumed cartel membership; Baldwin et al. (1997) use post-prosecution records to identify timber-auction colluders; Conley and Decarolis (2016) infer bidder groups from co-bidding patterns—the closest precedent in data parsimony, recovering group structure rather than firm-level flags; Clark et al. (2021) and Schurter (2020) study cartel dynamics with ex-post-known membership. The partition is constructed *ex post* or by assumption in every case. The construct supplies a candidate *ex ante* firm-level flag computable from award-record data, which bid-coordination tests can then take as input where microdata exist.

Bid-distribution screens require the bid layer. Imhof (2019); Imhof et al. (2018); Huber and Imhof (2019) and Wallimann et al. (2023) train machine-learning classifiers on tender-level bid statistics—within-tender coefficient of variation, kurtosis, skewness, normalized spread—reaching AUC above 0.85 against ground-truth European cartel cases. Abrantes-Metz et al. (2006) pioneer the variance screen; Chassang and Ortner (2022) provide moment-based theoretical foundations; Harrington (2008) surveys the field. Every implementation operates at the tender level and requires per-bidder bid amounts. Real enforcement environments often preserve only the award envelope; the construct runs on that envelope and integrates with bid-distribution screens at the forensic stage where microdata are recoverable (§9.4).

Cover-bidding identification has been retrospective. Porter and Zona (1999); Pesendorfer (2000); Asker (2010) and Marshall and Marx (2012) document cover bidders in prosecuted cartels through case records (Ohio milk, Italian highway, NYC stamps, and the sequence reviewed by Marshall and Marx); Athey et al. (2011) compare cover-bidding incentives across auction formats theoretically. In every empirical case the cover bidders are identified *ex post*—through leniency testimony, post-conviction records, or judicial discovery—and

after the cartel has been adjudicated. We are not aware of a published prospective identification of cover-bidder candidates from public award-record data. The construct shifts the post-adjudication observation into a pre-investigation screening stage.

Enforcement design under incomplete information. A separate literature studies how detection regimes should be structured when monitoring is imperfect. **?** and **?** establish the optimal-deterrence baseline under costly enforcement; Harrington (2008) surveys cartel-screening design in this register; **?** treats the optimal sequencing of investigative stages under information asymmetry between agency and target. Two implications travel to our setting. First, when bid-level observability is absent, the appropriate screening object is whatever statistic the collusive arrangement generates that survives the data-coarsening step—a question of identifiability under monitoring constraints rather than of detection power against a fixed informational benchmark. Second, screening stages and forensic stages answer different questions: screening allocates investigative attention; forensic analysis adjudicates a specific allegation. The architectural separation between an award-layer screening stage and a bid-layer forensic stage, which the empirical exercise instantiates, is the natural implication of this distinction.

Procurement institutions and the observability environment. A growing empirical literature documents how procurement institutions shape what enforcement can observe. Coviello and Gagliarducci (2017) show that longer political tenure worsens procurement outcomes in Italy; Decarolis et al. (2020) show that discretionary procedures and competition limits jointly predict corruption risk; and Decarolis and Giuffrida (2017) documents the procurement-officer channel. We treat these institutional features as variation in the monitoring regime that the construct's signal is read against. Heterogeneity in the construct's informativeness across buyer-size quartiles (§7.3) is interpreted in this register rather than as identification of an oversight channel.

Synthesis. The four literatures define the informational architecture in which the contribution sits. Bid-coordination tests require a partition; bid-distribution screens require the bid layer; cover-bidding identification requires the post-adjudication record; enforcement-design work asks what survives when observability is incomplete. The construct addresses

the conjunction: it identifies a participation primitive that survives data coarsening, supplies a candidate firm-level flag that bid-coordination tests can take as input, recasts cover-bidder identification as prospective rather than retrospective, and instantiates the screening–forensic architectural separation under the data-layer constraint that real enforcement faces. It does not dominate any individual benchmark; it operates where each cannot and integrates with each where it does.

3. Institutional Setting and the Observability Architecture

The empirical exercise requires three institutional facts about the data layer at which BEC enforcement operates and the source of the validation labels we use. We catalog these here and refer the reader to standard treatments of the underlying procurement statutes for further detail.

Two procurement modalities, two informational regimes. Brazilian public procurement during our 2009–2019 sample is governed by Lei 8.666/93 and Lei 10.520/2002. Two modalities dominate BEC and differ in observability and monitoring structure. *Convite* (Lei 8.666/93, Art. 22) is a sealed-bid invitation procedure with a statutory three-bidder quorum requirement and a low value cap. *Pregão* (Lei 10.520/2002) is a reverse electronic auction in which bidders submit progressively lower offers in real time, with no minimum-bidder requirement and a public bidding interface. The two modalities therefore expose different bidding behaviors to public visibility: convite preserves only the final award record; pregão preserves a richer real-time interaction whose footprint is observable to other bidders during the auction itself. We treat the modal contrast as variation in the observability environment that the construct's signal is read against, not as the mechanism the construct exploits. The empirical exercise in §8.2 reads modal heterogeneity as informativeness across detection regimes; that the price imprint and the cobidder discrimination are sharper in pregão is a property of the regime, not a falsification of an alternative mechanism. The 2021 reform (Lei 14.133/2021) postdates our sample and consolidates pregão-style auctions as the default, shifting the modal mix toward the regime in which the construct is most informative.

The data layer that survives. The construct operates on the data layer that BEC's analytical warehouse exports: winner identity, participant identity, item code, and negotiated price. Bid-by-bid records exist in the operational system but are not routinely exported to the layer that competition authorities and audit courts query.[2] The asymmetry is not Brazil-specific: EU member-state platforms preserve bid microdata heterogeneously, and the U.S. federal system (FPDS-NG) preserves awards but not bid history at the contracting-officer layer. The data-layer mismatch between the bid-distribution-screen literature and the operational records that survive at most enforcement agencies—which the construct's screening problem is built around—is therefore a feature of the architecture, not a parochial Brazilian artifact.

Adjudication portfolio for validation. The Conselho Administrativo de Defesa Econômica (CADE) is the federal competition authority that adjudicates bid-rigging cases under Lei 12.529/2011, with leniency provisions (Art. 86) generating the disclosure trail the cartel-detection literature uses as ground truth. CADE's procurement-cartel portfolio relevant to our sample comprises 12 adjudicated cases covering 65 firm-defendants (47 active in BEC during the sample window), with final decisions spanning 2015–2025. Eight of the twelve were decided after our 2019 sample cut. The temporal structure makes the validation in §6 partly prospective: the construct's discrimination against the adjudicated population is computed on participation patterns observable years before enforcement adjudicated those patterns. The interval between operational cartel activity and final adjudication—4 to 16 years from process opening to ruling in our portfolio—is the relevant benchmark for the screening-stage interpretation of the contribution: an award-layer screen that flags candidate environments at t accelerates the path to forensic investigation relative to the post-adjudication identification on which the prior literature relies.

The administrative-audit institutions surrounding BEC (TCE-SP, CGE) supervise

[2] In TCE-SP audits we examined, requesting bid-level microdata from the BEC operational system is a multi-step process that delays investigations by weeks; the award-record layer returns results in minutes.

procurement legality and maintain debarred-firm registries; we do not require them as part of the empirical strategy. We treat them as part of the broader monitoring environment within which the construct's screening signal is read.

4. Data and Frequent Losers Definition

4.1. BEC Platform

The data come from the Bolsa Eletrônica de Compras (BEC), São Paulo's centralized electronic procurement platform, operated by the State Treasury since 2002. BEC serves 1,308 purchasing units (PBUs)—state agencies, public universities, hospitals, schools—and clears two procurement modalities during our sample: *convite* (sealed-bid invitation procedure with a statutory three-bidder requirement, Lei 8.666/93) and *pregão* (electronic reverse auction with no minimum-bidder requirement, Lei 10.520/2002). The two modalities provide a within-platform institutional contrast that the empirical analysis exploits in §7.1 and §8.2; the construct itself does not depend on the convite minimum-bidder rule (§3). Our 2009–2019 window comprises 4,533,754 contract-record items across 22 calendar semesters, generated from 39,961,357 individual bids placed by 41,444 distinct firms (Table B.1). The contract-record layer preserves participants, winners, negotiated prices, reference prices, and procurement metadata; the bid-by-bid log is retained in the operational system but exported to the analytical warehouse only on administrative request, an envelope asymmetry central to the construct's portability claim (§10.3).

The modal mix shifts over the sample window. Pregão auctions expand from a minority share at the start of the sample toward near-universal coverage by 2019, consistent with the 2021 consolidation of pregão-style electronic auctions as the institutional default for Brazilian procurement (§3). Convite remains a meaningful share in the small-cap range and in early years, providing the modal variation that powers the falsification check of §7.1 and the discrimination comparison of §8.2.

4.2. Frequent Losers Definition

Behavioral anchor. Sustained participation with a zero win rate is rationally dominated under standard expected-profit competition. With positive per-bid cost $c > 0$, the

expected gain per bid for a firm Bayesian-updating its win probability after n consecutive losses falls below c once $n > \bar{V}m/c - 2$, where $\bar{V}$ is the mean procurement-order value and m is the gross margin (Online Appendix A). At BEC parameters ($\bar{V}$ = R$86,000 from the sample mean, $m = 0.10$ from documented procurement-supplier margins), this threshold is approximately $n = 15$ for c = R$500 and $n = 41$ for c = R$200. The plausible cost range is anchored on documented BEC bid-preparation effort (one to two staff-hours plus documentation compliance, $c \geq$ R$200); at the FL threshold of 14 zero-win participations expected profit is already negative under that anchor, and at the 90th-percentile FL participation count of 50+ zero-win tenders cumulative direct cost exceeds R$10,000. Sustained participation past this point is consistent with one of three readings: a non-pecuniary objective (cover bidding, capacity signaling, or eligibility maintenance); systematically incorrect beliefs that fail to update on the loss record; or zero-win BEC participation offset by wins in unobserved tenders elsewhere. The construct flags the participation footprint these readings share, and we organize the empirical work around the cover-bidding reading without claiming it as the unique interpretation (§10.2).

Two-step classification. We define frequent losers as the intersection of two conditions.

1. **Always-losers**: firms with a zero win rate across all 2009–2019 BEC tenders. The pool comprises 16,843 firms (41% of all BEC participants), reflecting low electronic participation costs, small firms bidding beyond their operational capacity, and the broader compositional features of decentralized procurement (Coviello and Gagliarducci, 2017; Decarolis and Giuffrida, 2017). Frequent losers are a small extreme subset of this population.
2. **IQR threshold**: among always-losers, those with participation counts above median + 1.5 × IQR of the always-loser participation distribution.[3] The threshold yields 14 or more tenders with zero wins, corresponding to a statistical cutoff of 13.5. The threshold was fixed *ex ante* on the participation distribution and the

[3]We use median + 1.5× IQR rather than the Tukey rule (Q3 + 1.5× IQR) because the participation distribution is heavily right-skewed; the adopted rule yields a stricter threshold. The Tukey alternative flags 1,981 firms with AUC 0.834 ([0.804, 0.863]), against 0.924 ([0.921, 0.926]) under the adopted rule. Robustness to the multiplier choice between 1.0× and 3.0× is reported in §9.1.

behavioral anchor above, before any examination of CADE labels. The definition yields 2,735 frequent-loser firms.

The treatment variable losers_{igt} equals one if tender-item i in item group g at time t has at least one frequent-loser firm among its participants, and zero otherwise.

The economic construct is broader than the binary rule. The continuous loss-intensity statistic $\log(1+\text{tenders_count})$ is the primitive that Proposition 2 identifies as a sufficient ranking statistic for cover-bidder type given award-record data; the binary 14-tender rule is the information-coarsening of this statistic under the framework's separating equilibrium (Online Appendix A). The horse race in §8.1 establishes that the continuous statistic strictly dominates the binary in discrimination (DeLong $p = 2e - 05$). We treat the binary rule as the deployable operational implementation and the continuous statistic as the discriminating instrument.

Strict temporal validation. The classification pools participation over the full 2009–2019 window, which maximizes within-stratum discrimination but does not isolate the deployable projection. When we estimate the threshold only on the 2009–2016 training window, the cutoff falls to 7 and the firm-level AUC is 0.767 ([0.734, 0.800]). Frozen against 2017–2019 items, the binary indicator drops to 0.565, while the continuous loss-intensity statistic retains 0.770 (Table 2). The temporal-holdout audit in §9.3 reports the same gap on operational metrics: in-sample precision overstates the deployable projection by approximately 50% at top-500, and the temporal-holdout AUC of 0.864 is the conservative discriminating reference we adopt for operational claims. The binary rule is a deployable triage trigger; the continuous statistic is the more persistent identifier of the loser-side participation footprint.

4.3. Sample Construction

The participation distribution among always-losers is heavily right-skewed (Figures B.1–B.2). The regression sample retains convite and pregão tenders with a recorded winner and item types with at least one frequent-loser participation during the sample window. The last restriction conditions on the treatment in the sense that item types that never

Table 1: Loser-Side Concentration as Concept, FL as Operational Rule

Measure	Editorial role	Threshold	Full-sample AUC	Strict firm AUC	Strict item AUC
Continuous log(1 + tc)	Latent construct	—	0.939	0.750	0.770
Binary FL: median + 1.5 IQR	Operational rule	13.5	0.924	0.767	0.565
Alternative: Q3 + 1.5 IQR	Tukey alternative	18.5	0.834	—	—
Alternative: median + 1.0 IQR	Looser rule	10.0	0.902	—	—

Notes: The strict-train threshold estimated only on 2009–2016 is 7.0, far below the full-sample paper-rule threshold of 13.5. The revised paper attaches theory to the continuous loser-side concentration object and treats the binary FL rule as an administrative compression rather than a theoretically unique boundary.

Table 2: Strict Temporal Holdout with Threshold Frozen on 2009–2016

Scope	Score	Threshold	AUC	95% CI	Positives / N
Firm-level (AL)	Binary FL	7.0	0.767	$[0.734, 0.800]$	193/21,819
Firm-level (AL)	$\log(1 + \text{tc})$	—	0.750	$[0.706, 0.795]$	193/21,819
Item-level 17–19	Binary FL	7.0	0.565	$[0.564, 0.566]$	5,299/1,009,729
Item-level 17–19	$\log(1 + \text{tc})$	—	0.770	$[0.764, 0.776]$	5,299/1,009,729

Notes: Train-window threshold (2009–2016 AL): 7.0. Full-sample reference: 13.5. Binary FL threshold estimated only on 2009–2016; item-level rows apply the frozen threshold to 2017–2019 items.

attract frequent losers anywhere in 2009–2019 are excluded; the restriction ensures controls come from item markets where the construct can plausibly fire. The unrestricted sample returns a coefficient within 0.5 percentage points of the restricted-sample baseline (§9), so the conditioning is empirically innocuous.

The resulting regression sample contains 1,654,401 tender-items across 18,783 item types, defined at the finest product-code level available in BEC (not aggregated to two-digit groups). Of the 18,783 item types, 15,101 (82%) have both frequent-loser-present and frequent-loser-absent tenders within sample, contributing to within-item identification under the fixed-effects specification of Equation (1). Frequent-loser firms participate in 4.8% of tender-items.

Table 3 reports unconditional descriptive statistics by frequent-loser presence. The unconditional negotiated-price ratio between frequent-loser-present and frequent-loser-absent tenders—R$140,673 versus R$12,465, a factor of approximately eleven—reflects the sorting of frequent losers toward larger contracts: frequent-loser-present tenders attract on average 9.06 bidders versus 4.90 in frequent-loser-absent tenders. Under within-item, within-year, within-procuring-unit fixed effects, the raw 2.43-log-point gap compresses to roughly 0.05 log points (Table B.2); the regressions compare similar tenders within tight product-code categories rather than across categories. Table 3 establishes the magnitude of the selection problem the fixed-effects structure is designed to absorb, and the comparison to Table B.2 shows the absorption is substantial.

4.4. CADE Validation Portfolio

The construct's external validation rests on the procurement-cartel adjudication portfolio of CADE, the federal competition authority (§3). The portfolio relevant to our sample comprises 12 adjudicated cases covering 65 firm-defendants in total, with final decisions spanning 2015–2025. Of the 65 firm-defendants, 47 match BEC participants by CNPJ root, and these 47 firms constitute the *direct-defendant* ground-truth population. Eight of the twelve cases were adjudicated after 2019, so the construct's discrimination against the adjudicated population was computed on participation patterns observable years before enforcement confirmed those patterns—a prospective character we exploit in

Table 3: Descriptive Statistics: Tenders With vs. Without Frequent Losers

	With FL			Without FL		
	Mean	SD	N	Mean	SD	N
Negotiated price	140,673.37	3,148,893.84	79,452	12,465.30	484,167.26	1,574,949
Log negotiated price	4.92	3.87	79,452	2.49	2.59	1,574,949
Number of firms	9.06	6.35	79,456	4.90	3.17	1,574,991
Log number of firms	1.99	0.67	79,456	1.39	0.65	1,574,991
Number of bids	18.56	24.55	79,456	9.58	13.19	1,574,991
Non-FL firms	7.82	6.01	79,456	4.90	3.17	1,574,991
FL count per tender	1.25	0.75	79,456	0.00	0.00	1,574,991

Notes: Sample restricted to item types with $\geq$1 FL tender. FL = frequent losers defined by median + 1.5×IQR threshold.

§6.

The validation operates against two distinct populations, and the distinction is central to the construct's claim. The *direct defendants* are the 47 firms named in CADE rulings as cartel members; in adjudicated cases these firms are typically frequent winners, the operational core of the collusive arrangement. The *cobidders* are the 193 firms that participated alongside the direct defendants in the same BEC tender-items, inside the always-loser stratum. Cobidders are cover-bidder *candidates* by structure: they share tenders with adjudicated cartel members without ever winning. The construct's target is the cobidder population, not the direct-defendant population. The AUC differential—firm-level 0.924 (95% CI [0.921, 0.926]) against cobidders versus 0.49 against direct defendants—is the empirical signature of this design (§6, §10.2).

5. Empirical Strategy

5.1. The Empirical Object: Screening Value Under Coarsened Observability

The empirical object is not a treatment effect of frequent-loser presence on prices. It is the screening value of endogenous loser-side participation given award-record data. The distinction is conceptual rather than rhetorical and governs the interpretation of every estimate in the paper.

A treatment-effect framing would ask: holding observable features fixed, what is the average price difference between items into which frequent losers enter and items into which they do not? Under that framing the relevant estimator restricts attention to environments in which treated and control items genuinely overlap on observables, and any estimate computed off that support is a confounded approximation of the within-support average. The screening framing reverses this. Under the separating equilibrium of Online Appendix A, frequent-loser participation is generated by the cartel's deployment problem; the cartel chooses where and how often to deploy cover bidders as a function of detection costs and item-specific incentives. Endogenous sorting into the items where the cartel deploys is constitutive of the screening object: the participation footprint the construct flags is, by construction, the footprint where the deployment optimum places it. The screening question asks how informative this footprint is about the underlying collusive environment when only award-layer outcomes are observable. The relevant estimator integrates over the distribution of items into which frequent losers actually enter, not over a counterfactual support reweighted to remove the sorting.

We therefore distinguish two empirical objects, both reported, neither claimed to be the other.

Object 1: broad-sample association (β). The within-cell conditional mean difference in log negotiated price between tender-items that include at least one frequent-loser participant and otherwise comparable items in the same product code, year, and procuring unit. We estimate this object on 1,654,401 tender-items and bracket it with four estimators (OLS, cross-fit, CEM, IPW). β is informative about the loser-side participation footprint where it is deployed; it is the screening-value object the framework's participation primitive is designed to recover. We do not claim β as a causal effect of frequent-loser presence and we do not defend it against the standard treatment-effect-identification critique. Under the screening framing, the question of whether β would survive a counterfactual that strips the deployment sorting is not the relevant question.

Object 2: overlap-restricted association (β^{ov}). The conditional mean difference in log negotiated price computed on the subsample of items where treated and control pop-

ulations genuinely overlap on observables. We estimate this object on the overlap cell and on the propensity-score-trimmed support (§7.1, Table 10). β^{ov} is the within-support comparison a treatment-effect framing would target; it is informative about the price differential at the items where the two populations look alike on observables, holding the participation footprint fixed at that subsample. We report it as a different empirical object from β, not as a sharper version of β.

The two objects answer different questions. The screening framing privileges β because the screening statistic is read against the cartel's deployment distribution, not against a counterfactual support. The treatment-effect framing privileges β^{ov} because identification is local to the subsample on which the treated-and-control populations are comparable. The empirical fact that $\beta > 0$ on the broad sample while $\beta^{\text{ov}} < 0$ on the overlap restriction (Section 7.1) is the empirical signature of this distinction: the screen's economic content is generated by the deployment sorting that the overlap restriction strips. We do not interpret the sign reversal as a robustness failure of the broad-sample estimate, because the broad-sample estimate is not a treatment-effect claim. We interpret it as the diagnostic that locates the construct as a screening-value object.

Two design-based strategies that return null. A sharp regression discontinuity at the convite/pregão statutory value caps is gated by the absence of bunching: McCrary-style density tests around the $R\$80{,}000$ cap return a left/right ratio of 0.94 and a log-density discontinuity of -0.063 (Table B.5), with the convite share dropping cleanly from 69% below to 59% above. Local first-stage power is too weak to support a fuzzy RDD at standard bandwidths. A difference-in-differences design using Decreto 9.412/2018's 2018 raise of the convite cap fails parallel trends at the available pre-treatment window. We report both failures openly. They do not constrain the screening interpretation, which does not require a causal identification of β, but they do constrain alternative framings of the empirical object that would.

5.2. Reduced-Form Specification

The reduced-form specification is

$$y_{igt} = \beta \cdot \text{losers}_{igt} + \mathbf{x}'_{igt}\boldsymbol{\delta} + \alpha_g + \lambda_t + \gamma_k + \varepsilon_{igt} \tag{1}$$

where y_{igt} is the outcome for tender-item i in item group g at time t and procuring unit k; losers_{igt} indicates the presence of at least one frequent-loser participant; $\mathbf{x}_{igt}$ includes a convite indicator and any specification-specific controls; α_g, λ_t, and γ_k are item, year, and procuring-unit fixed effects; and ε_{igt} is the disturbance, clustered at the item level. The headline outcome is $\log p_{\text{negotiated}}$; auxiliary outcomes (number of bidders, number of bids, number of non-frequent-loser bidders) are reported in §7.1 and the appendix.

We estimate Equation (1) under four specifications. The general within-item-and-year baseline drops the procuring-unit fixed effect; the within-procuring-unit specification adds it. Two modal subsamples, pregão-only and convite-only, decompose the headline along the observability-regime contrast set up in §3. The estimation sample comprises 1,654,401 tender-items (§4.3).

Three complementary estimators bracket Equation (1). Cross-fitting defines the frequent-loser list on odd years, estimates the equation on even years, and reverses the assignment, removing within-sample classification noise. Coarsened exact matching (CEM) and inverse-probability weighting (IPW) adjust for selection on observables (year, procedure type, item group, procuring-unit-size quartile)—under the screening framing these estimators bracket the broad-sample object β, not a treatment-effect target. A leave-one-out instrumental-variable specification uses the supply of frequent losers in other procuring units in the same product market and year as instrument; we read $\hat{\beta}_{\text{IV}}$ as a measurement-error diagnostic for the binary indicator's coarsening of the underlying continuous loss-intensity statistic, not as a causal estimate. The four estimators (cross-fit, IPW, OLS, CEM) jointly bound the headline range $+3.6\% - +7.7\%$ used throughout the paper (§7.1).

Standard errors are clustered at the item level as the baseline choice. Robustness to procuring-unit clustering and to two-way Cameron et al. (2011) clustering is reported in

§9.1.

5.3. Identification Audits as Screening-Value Diagnostics

The screening framing reorganizes the role of the four identification audits reported in §9.2. Each audit answers a distinct question about the screening object rather than a question about the credibility of a causal estimate.

The robustness values of Cinelli and Hazlett (2020) ($RV_{q=1} = 17.5\%$) and Oster (2019) ($\hat{\delta} = 261.64$) bound the broad-sample association against selection on observables already absorbed by item, year, and procuring-unit fixed effects. They calibrate how much unobserved-confounder variance would have to load on observables to explain β away. Under the screening framing this is the bound on selection-on-observables explanations of the screen's signal.

Strict-overlap matching delivers the second empirical object β^{ov} (Table 10). Its role is not to sharpen β but to instantiate the alternative empirical object whose contrast with β identifies the screening-value interpretation.

The leakage audit (Table C.2) decomposes the in-sample item-level discrimination AUC into a structural component and a tautology component, by holding cobidder firms out of score construction within cross-validation folds and by evaluating under temporal holdout. It is a generalization diagnostic for the screening statistic, not an identification test for β.

The sub-threshold-always-loser placebo (Table 11) confirms that the participation primitive operates through the threshold-crossing population the construct identifies, rather than through generic always-loser supply. It is a specification diagnostic for the participation primitive, not an identification test for β.

The four audits are reported in §9.2 and the identification scope they delineate is consolidated in §10.1.

5.4. Heterogeneity Specifications

We characterize heterogeneity in β along two dimensions. Both are descriptive specifications of Equation (1) estimated within sub-samples; we do not interpret either as institutional identification.

The first dimension is procuring-unit size, which the framework of Online Appendix A maps to the principal detection cost θ_k. We split the sample into quartiles of procuring-unit size (annual contract volume in BEC) and re-estimate Equation (1) within each quartile. Robustness to alternative size measures (cumulative item count, headcount of distinct purchasers) is reported alongside the headline. Under the screening framing, this specification reads the construct's informativeness across detection regimes (§7.3); the gradient is heterogeneity in the screening object across the monitoring environment, not identification of a causal channel.

The second dimension is the cell-heterogeneity grid in winner concentration and repeated-pair density. For each frequent-loser firm j, we compute the Herfindahl–Hirschman Index of winning firms across all tenders in which j participates, $\text{HHI}_j = \sum_w s_{wj}^2$, where s_{wj} is the share of j's tenders won by firm w. We classify j as high-HHI if HHI_j exceeds the sample median and low-HHI otherwise; orthogonally, j is high-pairs if it shares at least two tenders with at least one repeat partner among winners, and low-pairs otherwise. The 2×2 cross-classification yields four cells, and we estimate Equation (1) with cell interactions to characterize how the screening signal varies across cells. The specification anchors the cell-heterogeneity discussion of §8.3; we report the cell interpretation as descriptive scope of the screening object, not as a positive test of cartel-rotation predictions.

6. Validation: Cartel Adjacency Under Coarsened Observability

The validation asks the structurally appropriate question for a screening statistic computed on award-record data: does the construct concentrate on populations of firms with a higher prior probability of operating inside cartel-affected environments than random selection at comparable participation volume? The question is not whether the construct identifies cartel members. Award-record data does not contain the information required for membership identification, and the screening framing of §5.1 does not require it. We treat the cobidder population—firms that participated alongside CADE direct defendants—as the empirical content the screening statistic is designed to recover, and report against direct-defendant labels separately to bound the construct's scope.

6.1. Validation populations

Table 4: Validation populations and object discipline

Population	Size	Definition and role
All BEC participants	41,444	Universe of firms bidding in BEC, 2009–2019.
Always-losers	16,843	Firms with zero wins across the sample; stratum within which the screening statistic operates.
Frequent losers	2,735	Always-losers above the median $+1.5\times$IQR participation threshold; the flagged population.
CADE direct defendants in BEC	47	Firms named in CADE rulings; typically frequent winners; outside the screening statistic's target population.
Cobidders in the always-loser stratum	193	Always-losers that participated alongside CADE direct defendants; the cartel-adjacency label the screening statistic is designed to recover.

Notes: The screening statistic operates within the always-loser stratum, not across the broader BEC universe. Participation alongside an adjudicated cartel member is consistent with cover bidding, capacity-transfer motives, market intelligence, or bystander participation in cartel-affected markets. Strong AUC against this label corroborates that the screening statistic concentrates on cartel-relevant environments; it does not adjudicate membership.

6.2. Conservative benchmark

The primary benchmark restricts the ground truth to four CADE cases adjudicated before 2020 (coleta de lixo, aquecedores solares, trens-metrôs, tecnologia-informação 2019-10): 30 firm-defendants, all active in BEC, with 210 always-losers co-bidding with at least one of them, of which 108 are frequent losers. The participation-stratified co-bidding rate is 3.95% against a baseline of 1.24% ($p < 0.001$, 1,000 permutations), an excess ratio of 3.2×. The ROC AUC against the contemporaneous ground truth is 0.748 (95% DeLong CI $[0.713, 0.783]$).

This is the primary benchmark for the screening interpretation. The construct uses no information from cases adjudicated post-sample and the labels predate estimation: existence and membership were established by enforcement before the screening statistic was built. Within the always-loser stratum, frequent losers co-occur with adjudicated cartel members at rates above random matching at comparable participation volume.

The excess is what the screening framing predicts under the separating equilibrium of Online Appendix A: where the cartel deploys cover bidders, the loser-side participation footprint concentrates.

6.3. Prospective benchmark and within-firm enrichment

A more permissive benchmark uses CADE's full procurement-cartel portfolio (12 cases, 65 firm-defendants, of which 47 are active in BEC). Among 2,735 frequent losers, 193 (7.1%) co-participate with at least one of the 47; the participation-stratified excess ratio is 3.5× ($p < 0.001$). Firm-level AUC is 0.924 (95% CI [0.921, 0.926]) in-sample and 0.864 ([0.858, 0.870]) under temporal holdout (train 2009–2016, test 2017–2019). The benchmark treats post-sample adjudications as informative labels for cartel involvement during the BEC window—a reasonable but non-trivial assumption that the conservative benchmark in §6.2 avoids. We report it as complementary, not as the primary anchor.

A within-firm exercise on the same portfolio: 3 of 7 always-loser direct defendants (43%) cross the frequent-loser threshold against a population baseline of 16%. Table 5 lists them.

Table 5: Direct CADE defendants flagged by the construct

Firm	Cartel	CADE process	Tenders	Wins
Sol Tecnologia	Solar water heaters	08012.001273/2010-24	84	0
Nova Esperança Locadora	School transportation	08700.005876/2019-85	97	0
Jofran Comércio	Trash bags (Op. Colludium)	08700.005789/2015-02	65	0

These are the exception within a direct-defendant population that is otherwise frequent-winner-heavy; the AUC against direct defendants in the broader BEC firm universe remains 0.49. The asymmetry between cobidder discrimination (0.924 in-sample, 0.864 under temporal holdout) and direct-defendant discrimination (0.49) is the design's empirical signature, not a failure of validation. The screening statistic recovers loser-side participation footprints; direct defendants are by construction the winner side of the same arrangements and lie outside the target population.

6.4. Robustness to the enforcement record

Dropping all 31,447 CADE-involved tender-items and re-estimating the within-PBU baseline yields $\hat{\beta} = 0.062$ ($N = 1{,}622{,}954$), virtually unchanged from 6.4%. The screening signal in β operates independently of the enforcement record. The residual sample either contains undetected cartels that the screening statistic flags before enforcement reaches them, or higher-price environments correlated with frequent-loser presence through channels orthogonal to documented collusion. Both readings are consistent with the screening interpretation, and neither relies on the within-sample CADE adjudications for its content.

6.5. What the validation does and does not show

The validation supports four claims. First, the screening statistic concentrates on cobidder-labeled firms inside the always-loser stratum at a 3.2× contemporaneous excess (AUC 0.748) and a 3.5× prospective excess (AUC 0.924 in-sample, 0.864 under temporal holdout). Second, the price imprint β is insensitive to removing CADE-involved tender-items, which means the screening signal is not driven by the enforcement record itself. Third, the asymmetry between cobidder and direct-defendant discrimination delimits the scope: the screening statistic recovers cartel-adjacent environments, not cartel members. Fourth, the labels are indirect by construction—a feature of working at the award layer rather than the bid layer. The validation corroborates the screening interpretation of the construct; it does not, and under coarsened observability cannot, adjudicate the legal status of any individual firm.

7. Main Results

This section delivers the screening-value evidence in three moves. First, the broad-sample association β across four estimators (§7.1). Second, the overlap-restricted association β^{ov} and the sign reversal as the diagnostic that locates the construct as a screening-value object rather than a treatment-effect object (§7.2). Third, the heterogeneity of the screening signal across detection regimes (§7.3).

7.1. Broad-Sample Conditional Association (β)

Table 6 reports OLS estimates of Equation (1). Frequent-loser presence associates with a higher log negotiated unit price in all four specifications: 6.8% within item and year (6.4% adding procuring-unit fixed effects), and 9.3% in pregão against 3.8% in convite. The comparison set is the 1,654,401 tender-items in the same product-code × year × procuring-unit cells without frequent-loser participants.

Table 6: Negotiated Prices (log): Effect of Frequent Losers

	(1) General	(2) General	(3) Pregão	(4) Convite
FL presence	0.0677***	0.0636***	0.0933***	0.0382**
	(0.0230)	(0.0215)	(0.0255)	(0.0186)
Convite	−0.0090**	0.0095***		
	(0.0037)	(0.0035)		
Observations	1,654,401	1,654,401	546,549	1,107,852
R^2	0.879	0.886	0.882	0.893
Item FE	YES	YES	YES	YES
Year FE	YES	YES	YES	YES
PBU FE	NO	YES	YES	YES

Notes: Standard errors clustered at the item level in parentheses. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$. Columns (3) and (4) restrict the sample to pregão and convite procedures.

The pregão–convite contrast is informative on whether the construct's signal could be mechanically driven by the convite minimum-bidder rule (Lei 8.666/93, Art. 22, § 3): a quorum-filler reading would predict a sharper imprint in convite, but the data show the opposite (+9.59% in pregão vs. +3.92% in convite; Table B.4). The signal is therefore not reducible to the quorum rule. We read the asymmetry as informativeness across observability regimes (§8.2), not as a positive test of any specific institutional channel.

Cross-fit estimation yields 3.6%, CEM +0.077, and IPW 0.055. A leave-one-out IV returns 0.194 ($p < 0.05$, first-stage $F = 396$), the direction expected when the binary indicator coarsens the underlying continuous loss-intensity statistic (§8.1). The four broad-sample estimators jointly bound the conditional range +3.6% − + 7.7%.

Under the screening framing of §5.1, β is the empirical content the construct is designed to recover: the conditional association at the items frequent losers do participate in, integrated over the cartel's deployment distribution. We do not present β as a treatment-effect estimate. The relevant question for what follows is not whether β would survive a counterfactual that reweights items to remove the deployment sorting, but whether the deployment sorting itself is part of the screening-value content.

7.2. Overlap-Restricted Association (β^{ov}) and the Screening-Value Diagnostic

Restricting the comparison to cells in which frequent-loser-present and -absent items genuinely overlap on observables produces a reversed conditional coefficient: overlap-cell ATT −9.72%, propensity-score-trimmed ATT −30.67% (Table 10).

Table 7: Item-Level Scope Checks: Conditional Association Under Stricter Overlap

Specification	Coefficient	SE	p-value	N
baseline_fe	+0.064	0.021	0.003	1,653,658
overlap_cell_att	−0.097	0.015	< 0.001	1,517,066
overlap_ref_att	−0.097	0.015	< 0.001	1,434,636
ps_att_trimmed	−0.307	0.020	< 0.001	400,687

Notes: All rows remain descriptive. The overlap designs tighten comparability by forcing treated and untreated items into the same observed pre-treatment cells; the propensity-score row further trims to common support. None of these designs eliminates selection on unobservables, so the revised paper should call the resulting gap a conditional price association rather than a causal effect.

This is a different empirical object. β^{ov} is the within-support comparison at the items where treated and control populations look alike on observables, computed by holding the participation footprint fixed at that subsample. It is the estimand a treatment-effect framing would target. β is the broad-sample association at the items where frequent losers actually deploy, integrated over the deployment distribution. The two objects answer different questions; the sign reversal is the empirical signature that they are different.

Why we read the reversal as a screening-value diagnostic. Under the separating equilibrium of Online Appendix A, frequent-loser participation is generated by the cartel's

deployment problem. The cartel's choice of where to place cover bidders is informative: the deployment sorting is not a confound to be matched away but the equilibrium response to incentives we cannot observe directly. Three pieces of empirical evidence are consistent with this reading rather than with a treatment-effect-that-fails-on-overlap reading.

(i) The broad-sample association moves with cartel adjacency. The construct discriminates the 193 adjudicated CADE-cobidder firms inside the always-loser stratum at AUC 0.924 in-sample, 0.864 under temporal holdout, and 0.748 on the contemporaneous benchmark of §6. Under a confound-only reading of β, the construct's labels would not concentrate on adjudicated cartel members; the empirical concentration is inconsistent with that reading.

(ii) The continuous loss-intensity statistic strictly dominates the binary rule in discrimination of the same labels (DeLong et al. 1988 test, $p < 10^{-3}$; §8.1). The participation primitive is ranked, not categorical; the binary rule is the information-coarsening of a deeper signal. A pure-confound reading does not predict the dose-response.

(iii) Cinelli and Hazlett (2020)'s $RV_{q=1} = 17.5\%$ and Oster (2019)'s $\hat{\delta} = 261.64$ bound the broad-sample association against selection on observables of the kind already absorbed by the fixed effects. They do not bound the deployment sorting that the overlap restriction surfaces, by construction; under the screening framing, that sorting is the object, not the threat.

We therefore report β and β^{ov} as paired empirical objects. The construct's headline range is $+3.6\% - +7.7\%$ on the broad sample. The overlap-restricted coefficient (-9.72% to -30.67%) is the within-support comparison at the cells where treated and control populations are observable comparables; it is reported in parallel, not as a sharper version of the headline. Whether a referee privileges β or β^{ov} depends on which empirical object the reader treats as primary; we have made our reasons explicit and report both.

7.3. Heterogeneity Across Detection Regimes

Splitting the sample into quartiles of procuring-unit size and re-estimating the within-item specification produces a monotonic gradient: +21.4% in the smallest quartile, +9.8%, +4.5%, and +1.7% in the largest, all significant at conventional levels with item-clustered

standard errors. The 12.6× ratio is the largest source of heterogeneity in the data and is preserved across alternative measures of buyer size (annual contract volume, cumulative item count, headcount of distinct purchasers; Table B.3).

The framework in Online Appendix A predicts a gradient in the same direction through the detection-cost comparative static $\partial m^*/\partial \theta_k < 0$. We read the gradient as heterogeneity in the screening object across the monitoring environment: the construct is more informative where the detection regime exposes a thinner audit footprint. Buyer size proxies for several correlated institutional features—procurement-officer tenure (Coviello and Gagliarducci, 2017), internal-audit infrastructure, item composition, and discretionary procedure use (Decarolis et al., 2020; Decarolis and Giuffrida, 2017)—none of which can be separated with the variation we have. We do not identify a causal channel. The implication for the contribution is that the screening signal varies systematically with features of the enforcement environment, in the direction the framework's deployment problem predicts.

8. Heterogeneity of the Screening Signal

The screening object reads against the framework's deployment problem. This section reports two pieces of within-evidence that discipline the framing: first, that the continuous loss-intensity statistic strictly dominates the binary rule in discrimination, consistent with Proposition 2's identification of $\log(1 + \text{tenders_count})$ as the sufficient ranking statistic (§8.1); and second, that the modal contrast between pregão and convite reads as variation in the observability environment within which the screening signal is recovered (§8.2). §8.3 catalogs framework predictions that the data do not adjudicate under disciplined estimation; we do not weight these against the screening interpretation because the screening framing does not require their identification.

8.1. Loss Intensity Dominates the Threshold Cutoff

The deployable binary rule that defines frequent losers (more than 13.5 tender-items, zero wins) is the information-coarsening of the participation primitive identified by the framework, not a primitive in itself. Proposition 2 identifies $\log(1 + \text{tenders_count})$ as a

sufficient ranking statistic for cover-bidder type given award-record data; the binary rule is the information-coarsening of the latent posterior under the separating equilibrium of Online Appendix A. The horse race in Table 8 confirms the prediction empirically.

Table 8: Horse Race: FL14 Binary vs Continuous Loss Intensity (Same Sample)

Specification	Coef on FL14	Coef on log(1+ max tc)	Convite	N
(1) FL14 only	$+0.0653^{***}$	–	$+0.171^{***}$	1,653,658
	(0.0216)		(0.013)	
(2) Continuous only	–	$+0.0188^{***}$	$+0.169^{***}$	1,653,658
		(0.0055)	(0.013)	
(3) Both, additive	-0.0746^{*}	$+0.0349^{***}$	$+0.169^{***}$	1,653,658
	(0.0383)	(0.0108)	(0.013)	
Firm-level AUC against 193 *CADE-cobidder labels (always-loser pool,* $N = 16{,}843$*)*				
FL14 binary	AUC = 0.911 [0.898, 0.925]			
log(1+tenders_count)	AUC = 0.939 [0.932, 0.946]			
DeLong test (continuous vs FL14)	$Z = -4.30$, $p = 1.7 \times 10^{-5}$			

Notes: Outcome is $\log p_{\text{negotiated}}$ within the v13 panel of 1,653,658 items. Item-level treatment is $\mathbf{1}$[any FL14 participant] in column (1) and the log of the maximum loss-intensity score among always-loser participants in column (2). Column (3) includes both. The FL14 coefficient flips sign when continuous intensity is added to the specification: this is the disciplinary diagnostic that the binary discretization at 14 tenders is an operationalization of the underlying continuous signal, not an economically primitive threshold. We report the binary specification as a deployable implementation rule and the continuous specification as the substantive object. SEs clustered at item level. DeLong tests use the firm-level AUC reported in the bottom panel. *** $p < 0.01$, ** $p < 0.05$, * $p < 0.1$.

On a harmonized same-sample basis, the continuous $\log(1 + \text{tenders_count})$ score discriminates the 193 adjudicated CADE-cobidder firms inside the always-loser stratum at AUC 0.939 (95% CI [0.932, 0.946]), against AUC 0.911 ([0.898, 0.925]) for the binary FL14 indicator (DeLong et al. 1988 test, $Z = -4.30$, $p = 2e - 05$). The continuous statistic is strictly more informative.

Standalone price coefficients align with the discrimination result: FL14 binary +0.0653 ($p < 0.001$), continuous $\log(1 + \max \text{tc})$ +0.0188 ($p < 0.001$); Table 8, columns 1–2. When both are included additively (column 3), the FL14 coefficient flips sign (-0.0746, $p < 0.05$) while the continuous coefficient nearly doubles (+0.0349, $p < 0.001$). The flip is what the framework predicts: once the continuous primitive is in the specification, the binary indicator captures tail residuals against the continuous dosage rather than independent information about cover-bidder type. The dose-response pattern is consistent with the screening interpretation and inconsistent with a pure-confound reading of the

broad-sample association (§7.2).

We adopt the binary FL14 rule as the operational implementation throughout the paper because it is administratively interpretable; we adopt the continuous statistic as the discriminating instrument because it is what the framework identifies as sufficient. Both are reported. The headline range $+3.6\% - +7.7\%$ in §7.1 is computed on the binary; the AUC headline 0.924 (95% CI [0.921, 0.926]) reported in §6.3 is binary firm-level discrimination on the broader sample.

8.2. Modal Contrast as Observability-Regime Variation

The screening signal is sharper in pregão than in convite, along two dimensions. On the price-imprint side, the within-modal estimates of §7.1 place the binary specification at $+9.59\%$ ($p < 10^{-3}$) inside pregão versus $+3.92\%$ ($p = 0.037$) inside convite, a ratio of $2.45\times$. On the discrimination side, the modal-by-modal AUC against the 193 adjudicated CADE-cobidder firms is 0.952 in pregão primary auctions versus 0.816 in convite primary auctions, with a bootstrap difference of -0.136 ($p \approx 0$). Both dimensions point in the same direction.

The two BEC modalities expose different observability regimes (§3): pregão records a real-time electronic auction with public bidding interaction, convite a sealed-bid invitation procedure with a statutory three-bidder quorum requirement. A naive reading of the modal contrast would invoke convite's minimum-bidder rule as the mechanism through which the screening signal arises, in which case the screening should fire more cleanly in convite, not less. The data flatly reverse this prediction: the screening signal is sharper where the minimum-bidder rule is absent. We do not interpret the reversal as a positive test of any specific institutional theory. We read it as variation in the screening signal across observability regimes.

The reading is consistent with two compatible mechanisms that the design cannot separate: pregão's open electronic interaction may expose loser-side participation footprints more visibly than convite's sealed-bid envelope, or convite's smaller and more homogeneous bidder pools may compress the discriminating gap. The compatibility is itself informative for the JLEO frame: the screening object's informativeness varies with

the observability regime in which it is computed, and the regime that consolidates as the institutional default (Lei 14.133/2021, the post-sample reform that §3 catalogs) is the regime in which the screening signal is sharpest. Whether this is read as strengthening the construct's portability or as scope-restricting it depends on which observability regime a candidate jurisdiction operates under; both readings are consistent with the screening-stage interpretation.

The convite-modal AUC carries a width caveat the pregão-modal estimate does not: only 6 adjudicated CADE-cobidder firms fall in the within-convite stratum, so confidence intervals on the modal AUC difference are wider than the point estimate suggests. We report the point estimate because it aligns in direction with the price-imprint evidence in §7.1, on a much larger sample.

8.3. Predictions the Data Do Not Adjudicate

The framework generates two finer-grained predictions that the data do not adjudicate under disciplined estimation. We report them here for transparency. They are not load-bearing for the screening interpretation, which rests on the participation primitive (§8.1) and the observability-regime variation (§8.2), not on the cell-level structure of cartel rotation.

The first prediction concerns the HHI–repeated-pair-density grid over which cover-bidder rotation should be most pronounced (low HHI, low pairs). The data place the largest positive conditional price gap exactly in the Low-HHI $\times$ Low-pairs cell, in the predicted direction; but this cell is also the largest cell by item count, so the configuration is observationally consistent with both the framework's prediction and a generic compositional explanation. The audit of the diagonally opposite High-HHI $\times$ High-pairs cell (Table 9) returns a negative price coefficient (-0.062 to -0.119 depending on period and modality), concentrated in the 2009–2013 sub-period and in convite, which we read as period and modal compositional drivers rather than as refutation of the rotation prediction.

The second prediction concerns a discrete first-time-FL imprint as a firm crosses the threshold. The estimate sits in a $+0.062$ to $+0.100$ range across three specifications:

Table 9: Audit of the Negative High-HHI x High-Pairs Cell

Dimension	Group	FL coefficient	SE	p-value
modality	Convite	−0.062	0.024	0.009
modality	Pregao	−0.065	0.032	0.041
period	2014-2016	+0.009	0.033	0.796
period	2017-2019	−0.055	0.090	0.540
period	2009-2013	−0.119	0.025	< 0.001
direct_cade_item	0	−0.080	0.025	0.001

Notes: Each row re-estimates the within-item price association inside the problematic High-HHI x High-pairs cell. The purpose is not to rescue the mechanism, but to document where the sign flip lives. The full CSV additionally reports decompositions by buyer, item group, PBU size, and tender-value quartile.

unconditional event-study +0.100 ($p = 0.019$, $N = 9{,}118$); coarsened exact matching +0.085 ($p = 0.076$, $N = 1{,}930$); propensity-score matching +0.062 ($p = 0.312$, $N = 2{,}200$). Conditional matching compresses the magnitude substantially without inverting the direction. We adopt the propensity-score estimate as the conservative reference and report the full specification grid in Online Appendix C.

Neither prediction is load-bearing. Where the predictions track the data, they do so observationally; where the data do not adjudicate, we report the gap as scope rather than as mechanism evidence.

9. Robustness

This section stress-tests the within-item conditional price gap of §7.1 and the procuring-unit-size oversight gradient of §7.3 against four distinct concerns. §9.1 addresses sensitivity to the design choices embedded in the construct (the IQR threshold multiplier, the clustering of standard errors, and the binary-vs-continuous specification of loss intensity). §9.2 reports identification audits (Cinelli and Hazlett 2020 robustness value, strict-overlap matching, leakage diagnostics, and a placebo on random firm assignment). §9.3 compares in-sample operational metrics to a temporal-holdout audit, the most consequential disclosure for deployment claims. §9.4 compares the construct's discrimination to bid-distribution detectors trained on bid microdata, treating the gap as informational

complementarity rather than dominance.

9.1. Threshold and Specification Sensitivity

The headline price gap and its oversight gradient hold across reasonable variations in the design choices that define the construct. We document three: the IQR threshold multiplier, the clustering of standard errors, and the binary-versus-continuous specification of loss intensity.

The IQR multiplier governs the bright-line cutoff that separates frequent losers from the broader always-loser stratum. Our headline multiplier is 1.5, which yields a tender-count threshold of 13.5 and a flagged population of 2,735 firms. Table D.1 re-estimates the within-PBU specification under five multipliers ranging from 1.0× to 3.0× IQR. The price coefficient is positive and significant in every row, declining monotonically from +0.079 at 1.0× (threshold 10, 3,442 firms) to +0.050 at 3.0× (threshold 24, 1,456 firms). The decline is a feature, not a defect: as the cutoff tightens, the binary indicator filters fewer firms with high loss-intensity, and §8.1 established that the underlying primitive is the continuous $\log(1 + \text{tenders_count})$ statistic. The binary at 1.5× is the operational sweet spot between coverage and signal density; the construct neither evaporates at conservative thresholds nor inflates at permissive ones.

The clustering of standard errors does not change the point estimate—which is mechanical—but it does change inference. Table D.2 reports three regimes: item-level clustering (baseline, SE 0.0215), procuring-unit-level clustering (SE 0.0144, tighter than baseline), and two-way Cameron et al. (2011) clustering on item and procuring unit jointly (SE 0.0240, looser than baseline). The coefficient is significant at $p < 0.01$ under all three regimes; the baseline item-level cluster is the conservative middle choice.

The binary-versus-continuous specification of loss intensity is discussed in detail in §8.1; we summarize the implication for robustness here. The continuous $\log(1+\text{tenders_count})$ score discriminates the same adjudicated CADE-cobidder population at AUC 0.939 (95% CI [0.932, 0.946]) versus 0.911 ([0.898, 0.925]) for the binary FL14 indicator (DeLong $Z = -4.30$, $p = 2e - 05$). Standalone price coefficients in Table 8—FL14 binary +0.0653 ($p < 0.001$) and continuous +0.0188 ($p < 0.001$)—point in the same direction. The

headline price gap is not an artifact of binary discretization.

9.2. Identification Audits

The construct's identification status is descriptive throughout the paper. This subsection collects four audits that calibrate the strength of the descriptive claim under specific identification concerns: omitted-variable sensitivity, common-support restriction, leakage, and a placebo on a non-cover-bidding channel.

The Cinelli and Hazlett (2020) robustness value reported in §7.1 is $RV_{q=1} = 17.5\%$: an unobserved confounder would need to explain at least 17.5% of the residual variation in both frequent-loser presence and log price to drive the coefficient to zero. The Oster (2019) bounds in Table C.1 agree directionally. The full-sample $\hat{\delta} = 261.64$ implies that unobservables would need to be more than two hundred times as important as the observables already captured by item, year, and procuring-unit fixed effects to overturn the result; the pregão subsample yields an even larger $\hat{\delta} = 634.00$. The convite subsample carries a smaller margin ($\hat{\delta} = 43.62$, $\beta^*(\delta = 1) = -0.1395$)— consistent with §8.2's finding that the construct's signal-to-noise is highest in pregão environments and that the convite estimate is the noisier modal slice. Both sensitivity tools agree: the broad-sample association is unlikely to be explained away by selection on observables of the kind already accounted for.

The strict-overlap matching exercise complements the sensitivity bounds with a different identification angle. Where Cinelli and Oster ask how strong unobserved selection would need to be to overturn the result, strict-overlap matching restricts comparisons to the cells in which frequent-loser items and non-frequent-loser items genuinely overlap on observables. Table 10 reports four specifications. The baseline within-item-and-PBU coefficient is +0.064 (n = 1,653,658, identical to the headline). Restricting to overlap cells—requiring frequent-loser and non-frequent-loser items to share the same observed pre-treatment cell—returns −0.097 (n = 1,517,066). Trimming further to common-support under propensity-score restriction returns −0.307 (n = 400,687).

The reversal under strict overlap restricts comparisons to the items frequent losers *would* have entered with high probability, removing precisely the systematic selection

Table 10: Item-Level Scope Checks: Conditional Association Under Stricter Overlap

Specification	Coefficient	SE	p-value	N
baseline_fe	+0.064	0.021	0.003	1,653,658
overlap_cell_att	−0.097	0.015	< 0.001	1,517,066
overlap_ref_att	−0.097	0.015	< 0.001	1,434,636
ps_att_trimmed	−0.307	0.020	< 0.001	400,687

Notes: All rows remain descriptive. The overlap designs tighten comparability by forcing treated and untreated items into the same observed pre-treatment cells; the propensity-score row further trims to common support. None of these designs eliminates selection on unobservables, so the revised paper should call the resulting gap a conditional price association rather than a causal effect.

differences between selected and unselected items that the broad-sample association captures (§7.1). Both signs describe the same population from different angles, neither carrying causal weight on its own; we report the four specifications jointly so the reader can locate the construct's descriptive claim within the appropriate scope.

The leakage audit in Table C.2 addresses a distinct concern: that the in-sample item-level AUC of 0.995 (reference row) is partially tautological, because cobidders are by construction always-loser firms with elevated participation counts. Audit 1 holds the score fixed and swaps the label to direct-CADE defendants, returning 0.506—random discrimination, confirming the structural scope of the construct (cover bidders, not ringleaders). Audit 2 isolates the structural component by withholding cobidder firms from score construction within five cross-validation folds and re-evaluating on items containing held-out cobidders: AUC 0.891 (95% CI $[0.887, 0.894]$). Audit 3 evaluates the construct under temporal holdout (train 2009–2016, test 2017–2019): AUC 0.864 (95% CI $[0.858,\ 0.870]$). The screen retains discriminating power above 0.85 under both audits. The drop of 0.10–0.13 is the pure-leakage component, and we adopt the temporal-holdout AUC of 0.864 as the conservative discriminating reference for the operational-metrics audit in §9.3.

A placebo specification using the supply of sub-threshold always-losers (firms with persistent zero wins but below the IQR cutoff) as an instrument completes the audit

set. The first-stage F-statistic is 33, comfortably above the Stock and Yogo (2005) weak-instrument threshold, but the 2SLS coefficient is 5.24 (Table 11, below)—an implausible magnitude that signals the placebo instrument operates through channels other than cover bidding. The construct's price imprint draws on the above-threshold cover-bidder population we identify, not on generic always-loser supply.

Table 11: Placebo IV: Sub-Threshold Always-Losers as Instrument

	First Stage FL presence	2SLS Log Price
Sub-threshold supply (LOO)	0.000040***	
	(0.000007)	
F-statistic	33.0	
FL presence (instrumented)		5.2445***
		(1.3200)

Notes: Placebo instrument uses supply of sub-threshold always-losers (below the IQR cutoff). The sub-threshold first stage ($F = 33$) exceeds the Stock and Yogo (2005) 10% maximal IV size critical value of 16.38 for a single instrument. The implausibly large 2SLS coefficient (5.24) is therefore not attributable to weak-instrument bias. Rather, it reflects the absence of a genuine cover-bidding mechanism: sub-threshold always-losers do not drive procurement price anomalies, confirming that the main instrument's price effect operates through the cover-bidding channel rather than through correlated demand shocks. SE clustered at item level.

9.3. Operational Metrics: In-Sample vs Temporal Holdout

The deployment claim rests on operational metrics computed under the temporal-holdout regime, not under in-sample evaluation. The construct's discriminating performance under temporal holdout is firm-level AUC 0.864 (95% CI [0.858, 0.870]) against the 193 adjudicated CADE-cobidder labels, with the score constructed using only 2009–2016 participation counts and evaluated on items in 2017–2019. This is the operationally honest projection for a screening tool that is built on past participation data and deployed against a future case-load. We report it as the operational reference throughout the paper.

Table E.1 reports precision, recall, and lift at five top-k cuts under both in-sample and temporal-holdout evaluation. In the deployment-relevant region, top-100 has precision 0.070 (lift 6.1×), top-500 has precision 0.070 (lift 6.1×), and top-1,000 has precision 0.066 (lift 5.8×). At top-500, the screen flags 35 adjudicated cobidder firms out of 193 in the pool (recall 18%); at top-1,000 the recall rises to 34%. The construct's signal sustains through the deployable region, with lifts in the 5.8× to 6.6× range across cuts. The in-sample column over-states precision by approximately 50% at top-500 (0.132 in-sample versus 0.070 holdout, retention 53%). The inflation arises because roughly 47% of the top-500 in-sample ranking comes from 2017–2019 participation, after CADE adjudications were already underway for some of the cartels, so the in-sample score is partly retrospective rather than purely prospective. We disclose the gap explicitly and operate on the temporal-holdout column.

Three operational implications follow. First, the construct is deployable: discriminating AUC 0.864 and lifts of 5.8× to 6.6× in the top-100 to top-1,000 region are operationally meaningful for a screening tool that runs on contract-award records alone, at near-zero marginal data cost. Second, the temporal holdout is the honest reference: any operational claim in this paper that requires a precision or lift number is computed on the holdout column, and the in-sample column is reported for comparison and transparency. Third, periodic rescoring matters: an oversight body running the screen in production should refresh the score window as new years of contract-award data arrive, both to dampen drift and to keep the prospective component dominant in the deployment ranking.

9.4. Horse Race Against Bid-Distribution Detectors

The construct's competitive position vis-à-vis bid-distribution detectors that require bid microdata is the final calibration. The relevant comparison is at fixed *data cost*, not at fixed sample, because the two pipelines operate on different data envelopes: the construct uses award-record fields only (winner identity, participant identity, item identifier), while the bid-distribution pipeline requires per-bid offer values.

Table 12 reports the comparison on the full always-loser pool (N = 16,779 firms, 193 adjudicated cobidder labels). The construct discriminates at AUC 0.903 (95% CI

[0.884, 0.923]) under the binary FL14 indicator and 0.884 (95% CI [0.860, 0.908]) under the continuous log(1 + tenders_count) score. The full Imhof–Wallimann pipeline—five within-tender bid-distribution features (coefficient of variation, skewness, kurtosis, spread, second-low ratio) (Imhof, 2019; Wallimann et al., 2023)—reaches AUC 0.888 (95% CI [0.865, 0.911]) on the same target. The two pipelines reach comparable discrimination on different data envelopes. We do not claim that the construct outperforms the bid-distribution pipeline; we claim that it reaches comparable accuracy on a substantially thinner data envelope.

Table 12: Full Imhof–Wallimann Pipeline vs Loser-Side Concentration: Discrimination at Different Data Costs

Model	Features required	AUC	95% CI
Participation-only (award-record level)			
FL flag (binary)	participation count > 14, win flag	0.903	[0.884, 0.923]
log(1+tenders_count)	participation count, win flag	0.884	[0.860, 0.908]
Bid-distribution (per-bid microdata)			
Imhof CV only (v13 strawman)	per-bid offer values	0.585	[0.553, 0.616]
Imhof full pipeline	5 within-tender features (CV, skew, kurt, spread, second-low)	0.888	[0.865, 0.911]
Combined			
FL flag + Imhof full	participation count + 5 bid features	0.955	[0.943, 0.967]
log(1+tenders_count) + Imhof full	participation count + 5 bid features	0.962	[0.954, 0.969]

Notes: 5-fold cross-validated AUC against 193 CADE-cobidder labels within the always-loser pool ($N = 16{,}779$). Random forests with 500 trees, otherwise default tuning. Bid-level features are aggregated to the firm level by averaging the per-tender within-tender statistic across each firm's tenders. The two top blocks compare alternative single-construct screens operating on different data envelopes; the bottom block reports the combined classifiers. The combined model gains 5–7 AUC points over either single construct, indicating non-redundant signal in the participation-only and bid-distribution information sources. The participation-only screen achieves discrimination comparable to the full Imhof–Wallimann pipeline at substantially lower data cost, which is the operational claim of the paper. We do not claim the participation-only screen *outperforms* the bid-distribution pipeline; the relevant comparison is at fixed data cost, where the two pipelines reach similar accuracy and the combination dominates.

The combined specifications in the bottom panel of Table 12 address the integration question. FL14 combined with the full Imhof pipeline reaches AUC 0.955 (95% CI [0.943, 0.967]); continuous loss intensity combined with the full Imhof pipeline reaches AUC 0.962 (95% CI [0.954, 0.969]). Both combined classifiers gain 5–7 AUC points over either single construct, indicating non-redundant signal in the participation-only and bid-distribution information sources. Restricting attention to the same-sample audit in Table E.2, where all detectors evaluate the identical subset of firms with bid microdata available, the construct adds +0.035 AUC over the Imhof full pipeline alone (DeLong

$p = 0.014$), and the combined model adds +0.096–+0.098 AUC over Imhof full alone (DeLong $p < 0.001$).

The construct's competitive position is therefore informational complementarity, not dominance. It reaches comparable discrimination at lower data cost when bid microdata are unavailable, and it adds non-redundant signal when bid microdata are available. The administrative pathway in §11 integrates the construct with the bid-distribution stage rather than substituting for it: the **Screen** stage runs on award records, the **Forensic** stage deploys bid-distribution methods, and the gain from running both is identifiable in the data.

10. Limitations

The empirical content of the paper has a deliberately narrow scope. This section consolidates the limitations along three dimensions: what identification can deliver in §10.1, what the construct flags and what it does not in §10.2, and how the empirical claim travels beyond São Paulo's electronic procurement platform in §10.3.

10.1. Identification scope

Identification is descriptive throughout the paper. Two design-based strategies that would extract a causal estimate from the data return null at the available bandwidths: a sharp regression discontinuity at the convite/pregão statutory caps, and a difference-in-differences exploiting the 2018 cap raise. The first is gated by the absence of a McCrary-test-passing density around the institutional thresholds in the BEC sample; the second is gated by the absence of a clean pre-trend in the relevant outcome series. We report the failures openly and treat the empirical content as a descriptive association rather than as a causal estimate.

The strict-overlap matching exercise of §7.1, audited in §9.2, sets the limit on the descriptive claim. The within-PBU coefficient 6.4% on the broad sample reverses to −9.72% under overlap-restricted matching and to −30.67% under propensity-score trimming (Table 10). Neither the broad-sample association nor the overlap-restricted inversion carries causal weight on its own; both are conditional descriptions of the same population

at different scopes. The sensitivity bounds (Cinelli and Hazlett (2020) robustness value 17.5%, Oster (2019) $\hat{\delta} = 261.64$) discipline the broad-sample association against selection on observables already absorbed by item, year, and procuring-unit fixed effects, but do not extend to the selection on unobservables that overlap-restricted matching surfaces.

The implication is concrete. The $+3.6\% - +7.7\%$ headline range is a conditional descriptive association across estimators, not an estimate of cover-bidding's causal effect on prices. The 12.6× procuring-unit-size gradient is observationally consistent with the framework's comparative static $\partial m^*/\partial \theta_k < 0$, without identifying the causal channel. For the construct's intended use—prioritizing oversight case-load—this descriptive scope is sufficient: a screen needs to flag environments at higher prior risk than random selection would, on data the oversight body already holds.

10.2. Construct scope

The construct flags loser-side participation patterns, not cartel members. Inside the always-loser stratum the construct discriminates the 193 adjudicated CADE-cobidder firms at firm-level AUC 0.924 (95% CI [0.921, 0.926]) in sample and 0.864 (95% CI [0.858, 0.870]) under temporal holdout (§9.3). Against the 47 direct CADE defendants in the broader BEC firm universe, AUC is 0.49. The asymmetry is by design: direct defendants are typically frequent winners, while the construct operates on participation footprints with persistent zero wins. A deploying oversight body that wants to identify cartel ringleaders directly will need a different tool; the construct flags the loser side, and the AUC differential between the two populations is what that scope looks like in the data.

The participation footprint the construct flags is observationally compatible with at least three non-competitive readings, and our evidence does not adjudicate among them at the firm level. The working reading we adopt throughout the paper is cover bidding under cartel rotation, derived from the framework in Online Appendix A. Two competing readings remain: a *capacity-transfer* reading, in which firms maintain bidding presence to signal capacity availability for subcontracting or for affiliate consortia; and a *market-intelligence* reading, in which firms participate to observe prices, identify procurement

schedules, or build operational priors on competitor behavior. These readings are not mutually exclusive, and the same firm can have multiple non-competitive motives. The price imprint of §7.1 and the oversight gradient of §7.3 are observationally consistent with all three. We adopt the cover-bidding reading because the framework derives from it and because the strongest heterogeneity dimensions—the modal asymmetry between pregão and convite (§8.2) and the procuring-unit-size oversight gradient (§7.3)—match its comparative statics most directly. The adoption is a working interpretation, not an identification claim.

Two further mechanism predictions that the framework generates remain undiscriminated by the data. The cell-heterogeneity prediction across HHI and repeated-pair density places the largest positive price gap in the Low-HHI $\times$ Low-pairs cell, in the predicted direction (§8.3); but this cell is also the largest cell by item count, leaving the configuration observationally compatible with both the structural prediction and a generic compositional explanation. The first-time-FL prediction is in the predicted direction throughout (+0.062 to +0.100 across specifications), with conditional matching compressing the magnitude substantially without inverting it (§8.3). Neither mechanism prediction is refuted; neither is cleanly confirmed.

The construct's behavior under deliberate adversarial adaptation is bounded but not invariant. We simulate four strategic adaptations holding the construct fixed (Online Appendix G, Table F.1): periodic rotation of cover bidders, occasional wins to evade the bright-line filter, CNPJ splitting, and threshold-aware participation capping. The construct is resilient to rotation and occasional wins (AUC essentially unchanged), degrades modestly under CNPJ splitting (-3.7 points), and is substantially vulnerable to threshold capping (-12.7 points). Combining threshold capping with rotation drives AUC from 0.928 to 0.632. The construct is therefore informative absent adaptation and degrades predictably under specified attacks; the threshold rule is the most exposed margin, and operational lifetime in any deployment requires periodic recalibration and integration with bid-distribution methods where those data are available.

10.3. External validity and portability

The empirical evidence in this paper comes from a single jurisdiction and a single time window: São Paulo's Bolsa Eletrônica de Compras, 2009–2019. The construct's portability beyond this envelope rests on three structural features of the procurement environment, all explicit in BEC and identifiable in any candidate jurisdiction before deployment: a commodity-heavy item composition that supports within-product-code price comparisons; an electronic platform with publicly auditable participant lists, so that participation counts can be computed without privileged data access; and an institutional asymmetry between the winner and loser sides of the bidding pool that makes persistent zero wins informative about strategic non-competitive participation rather than about randomized matching outcomes. These are diagnostic conditions, not optional ones.

Settings that this paper does not test directly are several. The federal Brazilian procurement registry (ComprasNet) and other state-level registries (Banco do Brasil, individual UF platforms) preserve the same award-record envelope as BEC and are accessible through the same legal regime; replication is technically straightforward and substantively pending. Non-commodity settings (services, infrastructure, public works) are within the framework's logical scope but outside the empirical scope of the present work, and the within-item identifying variation that underwrites the headline price gap is not equally available in those settings. Time periods outside the sample window are likewise untested: the BEC sample begins after the platform's operational stabilization and ends before the consolidation of the new federal procurement law, and the construct's stability across regulatory regimes is an open empirical question.

The new federal procurement law (Lei 14.133/2021) consolidates the pregão-style electronic-auction format as the institutional default for Brazilian federal and inter-state procurement. Combined with the modal asymmetry of §8.2—the construct's discrimination is sharper in pregão—the regulatory direction is favorable to the construct's portability rather than adversarial. Portability remains contingent on the three prerequisites above; the contingencies are recognizable, observable, and increasingly satisfied by the regulatory direction Brazilian procurement is taking.

11. Conclusion

The bid-rigging detection literature has been built on the assumption that bid-level data are observable. Real enforcement environments routinely violate that assumption: many electronic procurement systems preserve only the contract-award envelope, and the bid-distribution methods that work on the bid layer become inoperative on the layer that survives. The paper takes this informational mismatch as its starting point and asks what collusion-relevant statistic can be recovered when only the award layer is observable.

Endogenous loser-side participation is such a statistic. A separating-equilibrium framework with cover bidders and genuine entrants identifies $\log(1 + \text{tenders_count})$ together with the wins $= 0$ filter as a sufficient ranking statistic for cover-bidder type given award-record data; the binary rule we deploy is its information-coarsening. The empirical exercise on São Paulo's BEC (2009–2019) finds a broad-sample association of $+3.6\% - +7.7\%$ in the conditional log negotiated unit price across four estimators, discriminating power against adjudicated CADE cobidders at firm-level AUC 0.924 in-sample and 0.864 under temporal holdout, and a 3.2× contemporaneous excess against the conservative pre-2020 benchmark. The continuous primitive strictly dominates the binary rule, consistent with the framework's identification of the loss-intensity statistic. The construct does not discriminate direct CADE defendants in the broader BEC firm universe (AUC 0.49): the asymmetry between cobidder and direct-defendant labels is the design's empirical signature, not a failure.

The screening framing reorganizes how the empirical evidence is read. The broad-sample estimate is informative about the participation footprint where the cartel deploys it, integrated over the deployment distribution; the overlap-restricted estimate (-9.72% to -30.67%) is the within-support comparison at items where treated and control populations overlap on observables. The two are different empirical objects; the sign reversal between them is the diagnostic that locates the construct as a screening-value object rather than a treatment-effect object. The buyer-size gradient (12.6× across procuring-unit quartiles) is read as variation in the screening signal across detection regimes, not as identification of a causal channel through which monitoring infrastructure operates on

prices. Three independent corroborations support the screening reading: cartel-adjacency validation moves with the broad-sample association; the continuous statistic delivers a model-predicted dose-response; sensitivity bounds discipline selection on observables of the kind the fixed effects already absorb.

Two implications for enforcement design follow. First, what survives data coarsening is not a moment of submitted bids but the equilibrium-generated participation footprint that the collusive arrangement leaves on the award layer. The information that the bid-layer literature reads off bid distributions is partially recoverable from a different statistic at the award layer when bid data are unavailable, and the screening object the recovery delivers is a participation statistic rather than a price moment. Second, the architecture of enforcement under incomplete observability separates an award-layer screening stage from a bid-layer forensic stage. The screening stage produces candidate flags; the forensic stage adjudicates suspected coordination on the bid distribution where microdata exist. The two stages answer different questions and require different statistics. The construct fills the first; the bid-distribution literature fills the second; the benchmark in §9.4 confirms that the two operate on the same underlying cases without dominating each other.

The contribution stops at the screening stage by design. The construct does not adjudicate cartel membership and cannot do so under coarsened observability; the cobidder population is the structurally appropriate validation object given the data layer. Adversarial adaptation degrades the construct's discrimination predictably (Online Appendix G), and replication beyond São Paulo remains an open empirical question that the construct's three prerequisites—a winner-loser participation registry, an electronic platform, and an institutional asymmetry between the two sides of the bidding pool—will determine. The conceptual content travels: wherever an enforcement environment exposes the award layer without exposing the bid layer, the question is not whether bid-distribution screens can be retrofitted but what participation primitive the collusive arrangement generates that survives the coarsening, and how the screening stage that reads that primitive should be sequenced before the forensic stage that requires bid-level evidence.

References

Abrantes-Metz, R. M., L. M. Froeb, J. F. Geweke, and C. T. Taylor (2006). A variance screen for collusion. *International Journal of Industrial Organization 24*(3), 467–486.

Asker, J. (2010). A study of the internal organization of a bidding cartel. *American Economic Review 100*(3), 724–762.

Athey, S., J. Levin, and E. Seira (2011). Comparing open and sealed bid auctions: Evidence from timber auctions. *Quarterly Journal of Economics 126*(1), 207–257.

Bajari, P. and L. Ye (2003). Deciding between competition and collusion. *Review of Economics and Statistics 85*(4), 971–989.

Baldwin, L. H., R. C. Marshall, and J.-F. Richard (1997). Bidder collusion at forest service timber sales. *Journal of Political Economy 105*(4), 657–699.

Cameron, A. C., J. B. Gelbach, and D. L. Miller (2011). Robust inference with multiway clustering. *Journal of Business & Economic Statistics 29*(2), 238–249.

Chassang, S. and J. Ortner (2022). Robust screens for non-competitive bidding in procurement auctions. *Econometrica 90*(1), 315–346.

Cinelli, C. and C. Hazlett (2020). Making sense of sensitivity: Extending omitted variable bias. *Journal of the Royal Statistical Society: Series B 82*(1), 39–67.

Clark, R., J.-F. Houde, and J. Kastl (2021). The dynamics of bidder collusion: Evidence from a long-lived cartel. *Journal of Political Economy 129*(8), 2353–2391.

Conley, T. G. and F. Decarolis (2016). Detecting bidder groups in collusive auctions. *American Economic Journal: Microeconomics 8*(2), 1–38.

Coviello, D. and S. Gagliarducci (2017). Tenure in office and public procurement. *American Economic Journal: Economic Policy 9*(3), 59–105.

Decarolis, F., R. Fisman, P. Pinotti, and S. Vannutelli (2020). Rules, discretion, and corruption in procurement: Evidence from italian government contracting. NBER Working Paper 28209.

Decarolis, F. and L. M. Giuffrida (2017). Corruption in procurement. *Working Paper*.

DeLong, E. R., D. M. DeLong, and D. L. Clarke-Pearson (1988). Comparing the areas under two or more correlated receiver operating characteristic curves: a nonparametric approach. *Biometrics 44*(3), 837–845.

Harrington, J. E. (2008). Detecting cartels. In P. Buccirossi (Ed.), *Handbook of Antitrust Economics*, pp. 213–258. MIT Press.

Huber, M. and D. Imhof (2019). Machine learning with screens for detecting bid-rigging cartels. *International Journal of Industrial Organization 65*, 277–301.

Imhof, D. (2019). Detecting bid-rigging cartels with descriptive statistics. *Journal of Competition Law & Economics 15*(4), 427–467.

Imhof, D., Y. Karagök, and S. Rutz (2018). Screening for bid rigging—does it work? *Journal of Competition Law & Economics 14*(2), 235–261.

Marshall, R. C. and L. M. Marx (2012). *The Economics of Collusion: Cartels and Bidding Rings*. Cambridge, MA: MIT Press.

OECD (2019). Government at a glance 2019. Technical report, Organisation for Economic Co-operation and Development, Paris.

Oster, E. (2019). Unobservable selection and coefficient stability: Theory and evidence. *Journal of Business & Economic Statistics 37*(2), 187–204.

Pesendorfer, M. (2000). A study of collusion in first-price auctions. *Review of Economic Studies 67*(3), 381–411.

Porter, R. H. and J. D. Zona (1993). Detection of bid rigging in procurement auctions. *Journal of Political Economy 101*(3), 518–538.

Porter, R. H. and J. D. Zona (1999). Ohio school milk markets: An analysis of bidding. *RAND Journal of Economics 30*(2), 263–288.

Sánchez Graells, A. (2019). 'screening for cartels' in public procurement: Cheating at solitaire to sell fool's gold? *Journal of European Competition Law & Practice 10*(4), 199–211.

Schurter, K. (2020). Identification in auctions with collusion. *Working Paper*.

Stock, J. H. and M. Yogo (2005). Testing for weak instruments in linear IV regression. In D. W. K. Andrews and J. H. Stock (Eds.), *Identification and Inference for Econometric Models: Essays in Honor of Thomas Rothenberg*, pp. 80–108. Cambridge: Cambridge University Press.

Wallimann, H., M. Huber, and D. Imhof (2023). A machine learning approach for flagging incomplete bid-rigging cartels. *Computational Economics 62*, 1669–1720.