

Cheap Signals, Costly Proof: Award-Layer Evidence Triage in Cartel Enforcement

Darcio Genicolo-Martins^{1,*}, Paulo Furquim de Azevedo¹

¹*INSPER, São Paulo, Brazil*

Abstract

Cartel enforcement begins before cartel proof exists. Agencies often observe contract-award records long before they recover the bid-level evidence needed to prove coordination. This paper asks whether that cheap award layer can allocate costly forensic attention. Using São Paulo’s BEC procurement platform (2009–2019) and CADE adjudications, we construct a *frequent-loser* screen from award records alone: among zero-win firms, persistent participation ranks loser-side cartel-adjacency risk. The screen targets forensic priority, not cartel membership. We validate the ranking against adjudication-anchored cobidder labels, discipline it with exposure, timing, leakage, and direct-defendant scope checks, and compare it with bid-distribution forensics on the same target. The award-layer signal adds non-redundant information to bid-layer forensics. Used sequentially, the award-layer ranking reduces the bid-microdata pool for the forensic stage by 83% while recovering 131 of 193 adjudicated cobidders. The results support an enforcement-design claim: cheap administrative records can discipline where proof-producing investigation should begin while liability remains in the richer evidentiary record.

Keywords: cartel screening, public procurement, award-layer data, enforcement triage, incomplete observability, bid rigging

JEL No.: D44, D73, H57, K21, L41

*Corresponding author.

Email address: darciogm1@insper.edu.br (Darcio Genicolo-Martins)

We thank seminar participants at INSPer for comments; all remaining errors are our own.

1. Introduction

Cartel enforcement begins before cartel proof exists. An agency deciding where to investigate rarely starts with the communications, bid histories, leniency testimony, or tender-level conduct evidence that would support liability. It starts with thinner information: administrative records, complaints, procurement anomalies, repeated interactions, and case leads that make some firms and markets worth a closer look. The first legal-economic problem is therefore not proof itself. It is the allocation of proof-producing effort.

That distinction matters in procurement. The empirical cartel-screening literature usually starts from the bid layer: submitted bid amounts, bid ranks, within-tender dispersion, bid histories, and sometimes bidder networks. Classic procurement screens compare bidding patterns under competition and collusion ([Porter and Zona, 1993](#); [Bajari and Ye, 2003](#)); variance and descriptive screens formalize suspicious bid distributions ([Abrantes-Metz et al., 2006](#); [Harrington, 2008](#)) and field-test them in procurement settings ([Imhof et al., 2018](#); [Imhof, 2019](#)). Recent work extends the approach through bidder groups ([Conley and Decarolis, 2016](#)), machine learning ([Huber and Imhof, 2019](#); [Wallimann et al., 2023](#)), and robust screen design ([Chassang and Ortner, 2022](#)). Those objects are close to the evidentiary stage, but they are not always cheap to obtain at scale. Many procurement systems make contract-award records routinely visible while keeping per-bidder bid files in operational systems that require case-specific recovery. An enforcement agency then faces a sequencing problem: before opening the costly bid layer, can it use the cheap award layer to decide where bid-layer forensics should begin?

This paper studies that question on São Paulo’s BEC procurement platform, covering 1,654,401 tender-items from 2009 to 2019. BEC records a routine award layer—participants, winners, item codes, negotiated prices, procuring units, years, and modalities—and also preserves a richer bid layer recoverable through the LANCES export. CADE procurement-cartel adjudications provide external legal anchors. This institutional combination lets us study an earlier stage than most bid-distribution screens: the point at which an agency observes award records but has not yet recovered the bid-level evidence

needed for forensic evaluation.

The award-layer statistic is deliberately simple. Among firms with zero wins, we rank participation intensity using $\log(1 + \text{tenders_count})$ and implement an auditable binary rule that flags *frequent losers*. The economic logic is monotone rather than classificatory. If coordinated bidders use loser-side participation to create the appearance of competition while avoiding the win, persistent zero-win participation can rank loser-side exposure under a transparent behavioral condition. The statistic does not identify a firm type or assign cartel membership. In the paper’s staged design, the rank allocates forensic attention; the richer bid layer remains where tender-level conduct and agreement are evaluated.

That legal scope determines the validation target. Direct CADE defendants are legal membership anchors, but they are not the natural object of a zero-win loser-side screen. Procurement cartels often allocate winning roles, rotate designated winners, and use other firms on the losing side of the tender (Pesendorfer, 2000; Asker, 2010). The same role-allocation logic is central in cartel theory and dynamic case evidence (Marshall and Marx, 2012; Clark et al., 2021). Large-scale procurement evidence shows how such arrangements can appear in administrative bidding records (Kawai and Nakabayashi, 2022). We therefore validate the screen against always-loser firms that bid alongside direct CADE defendants in adjudicated cartel environments. These cobidders define an adjudication-anchored exposure target. Direct defendants anchor legally adjudicated environments; cobidders are the loser-side footprint the award layer can plausibly rank before bid-level proof is recovered.

The empirical design is organized around the threats that would make such a screen uninformative for enforcement. A raw co-bidding association could reflect participation volume alone, so we compare observed concentration with participation- disciplined placebos and exposure-adjusted audits that hold fixed the markets in which firms could meet CADE defendants. A retrospective score could reuse information from the target period, so we report timing exercises that form the score before the evaluation window and distinguish relaxed from strictly ex ante definitions. A loser-side statistic could be overread as a cartel-membership detector, so we test whether it also classifies direct

defendants. It should not.

The results support a limited but operationally useful claim. The frequent-loser ranking concentrates adjudication-anchored cobidders inside the always-loser stratum, including when the score is formed before the target window. It performs poorly against direct CADE defendants, consistent with its loser-side scope rather than with a generic cartel-membership interpretation. Cobidders inside the frequent-loser stratum also look different from other frequent losers: they are deployed more broadly, appear closer to legal cartel anchors, operate in more concentrated product portfolios, and display bid-level patterns consistent with credible losing roles. Those patterns give the adjacency target economic content and make it harder to reduce the result to ordinary high-volume losing alone.

The main institutional result is a cost-of-evidence result from the sequencing exercise. On the same adjudication-anchored target, the award-layer score adds information to an Imhof–Wallimann-style bid-distribution benchmark. The two layers therefore enter at different stages: award records prioritize where to look, and bid records evaluate what is found. Used as a gatekeeper, the award-layer ranking cuts the bid-microdata pool for the forensic stage by 83% while recovering 131 of 193 adjudicated cobidders. Joint scoring with award and bid features is the full-observability upper bound; sequential gatekeeping is the architecture relevant for agencies that must decide where to spend costly bid recovery.

This paper makes three contributions. First, it reframes procurement-cartel screening as an evidence-allocation problem under costly observability. Existing bid-distribution screens show how to evaluate suspicious bidding once bid-level data have been recovered; we study the prior institutional stage, where agencies observe award records but must decide where recovery should start. This turns screening toward enforcement triage and away from treating a statistical flag as proof ([Harrington, 2008](#); [Sánchez Graells, 2019](#)); robust-screen design formalizes the same distinction between flags and evidentiary conclusions ([Chassang and Ortner, 2022](#)). Second, the paper shows that a minimal award-layer statistic—persistent zero-win participation—ranks a legally limited loser-side adjacency target anchored in CADE adjudications. Third, it evaluates a sequential

architecture in which cheap award-layer triage gates costly bid-layer forensics, reducing the evidentiary footprint without moving liability out of the richer case record.

The boundaries are part of the design. The award layer ranks forensic priority, the validation target is adjudication- anchored exposure, and the bid layer remains the proof-producing stage. Price evidence is treated as scope information because the association changes across support restrictions. Strategic response also shapes implementation: a static cutoff invites gaming, while refreshed thresholds, continuous ranks or capacity-constrained queues, bunching checks, and bid-layer follow-up fit the staged enforcement architecture.

The broader lesson is institutional. A legal system that waits for full proof before ranking investigative priorities will spend scarce forensic capacity poorly; a legal system that treats cheap suspicion as proof will overreach. The useful middle ground is a disciplined triage rule: one that is cheap enough to run before proof exists, informative enough to order forensic attention, and legally limited enough to leave liability where it belongs—in the richer evidentiary record.

2. Costly Proof and Observable Awards

The paper’s empirical object is an institutional sequence, not a stand-alone classifier. An enforcement agency first observes a thin administrative record, then decides whether to spend resources recovering richer evidence, and only later can that evidence support legal proof. This section defines the sequence for BEC procurement. Subsection 2.1 states the screening-before-proof problem. Subsection 2.2 separates the award layer from the forensic bid layer. Subsection 2.3 defines the CADE anchors and the loser-side cartel-adjacency target used for validation.

2.1. Screening Before Proof

Antitrust agencies do not choose investigation under full information. They choose it while the facts needed for liability are still costly to obtain. In a procurement-cartel case, proof requires evidence of coordinated conduct: communications, allocation rules, bid rotation, cover bidding, or tender-level patterns inconsistent with independent rivalry.

Brazilian competition law treats bid rigging as an infraction under Lei 12.529/2011; comparable U.S. and European frameworks also place agreement and anticompetitive effect at the center of the case ([Hovenkamp, 2005](#)). A pattern in award records does not contain those legal predicates.

The institutional question is therefore earlier than proof. The Becker–Stigler tradition treats enforcement intensity, investigation, and sanctions as costly instruments ([Becker, 1968](#); [Stigler, 1970](#)). Cartel-detection work makes the same allocation logic operational: when investigation is costly and misconduct is only partially observable, agencies should use cheap screens to decide where expensive inquiry has the highest expected return ([Baker, 2003](#); [Harrington, 2008](#)). Leniency policy illustrates the same enforcement trade-off in cartel cases ([Wils, 2007, 2016](#)). For procurement cartels, the first-stage problem is not whether a cheap statistic can decide liability. It is whether a cheap statistic can order the queue for recovering the evidence that could later support liability.

This is the sense in which the paper studies triage rather than classification. A model that uses submitted bid amounts, bid ranks, and within-tender bid moments assumes the forensic layer has already been opened. That is the right information set for evaluating suspicious bidding, but not for deciding where the first costly bid-file request should be sent. An award-layer screen has value only at this earlier point. Its output is a priority list for proof-producing work, not an accusation.

The legal-economic object is *forensic priority*. A useful triage statistic must be computable from records already available before case-specific bid recovery; it must rank firms or environments where richer cartel evidence is more likely to be worth recovering; and it must remain legally limited, so that administrative suspicion does not substitute for proof. The BEC setting supplies the two data layers and the external adjudication anchors needed to evaluate whether an award-layer statistic can play that role.

2.2. Award and Bid Layers in BEC

São Paulo’s BEC platform is useful because it preserves the two information layers that an enforcement sequence requires. The sample contains 1,654,401 tender-items from 2009 to 2019. Two procurement modalities dominate. *Convite* is a sealed-bid invitation

procedure under Lei 8.666/93 with a statutory three-bidder quorum and a low value cap. *Pregão* is a reverse electronic auction under Lei 10.520/2002, with progressively lower real-time offers and no minimum-bidder rule. The procedures differ, but both generate a routine award record before any case-specific bid recovery occurs. This matters because procurement rules, auction format, and bureaucratic implementation can shape both bidding behavior and observed procurement outcomes. Auction format affects participation and bidding incentives (Athey et al., 2011; Decarolis, 2014); public procurement outcomes also depend on bureaucratic incentives, competence, discretion, and rules (Coviello and Gagliarducci, 2017; Decarolis et al., 2020); rules and discretion create an additional institutional margin (Decarolis et al., 2025).

The *award layer* is that routine administrative envelope. It records participant identity, winner identity, item code, negotiated price, procuring unit, year, and modality. These fields identify who appeared, who won, what was bought, and where a firm repeatedly participated without winning. They do not reveal the bids submitted by each participant, the timing of offers, bid revisions, or within-tender dispersion. The award layer is therefore cheap enough for broad monitoring, but too thin to supply proof of coordination.

The *bid layer* is the forensic-recoverable layer. BEC preserves per-bidder offers, timestamps, and bid-revision sequences in the LANCES export. Those records support bid-distribution forensics: within-tender dispersion, skewness, spread, distance from the winning bid, and related Imhof–Wallimann-style features (Imhof et al., 2018; Imhof, 2019). Machine-learning variants combine these bid moments at scale (Huber and Imhof, 2019; Wallimann et al., 2023). But these features require the bid file to be recovered, processed, and linked. They are not available at the same cost as an award-record query.

The empirical design treats that cost difference as substantive rather than incidental. If bid microdata were already available for every firm at no meaningful cost, the agency could run forensic screens directly. The institutional problem arises because the bid layer is available but costly: feasible for case work and validation, but not the natural first object for system-wide triage.

Figure 1 summarizes the distinction. It is an information-cost diagram, not a detector horse race. Bid-distribution methods remain the natural forensic tools once bid microdata




Figure 1: Information layers in procurement-cartel enforcement. Bid-distribution forensics requires per-bidder bid amounts and within-tender bid moments; the award-layer screen uses participant identity, winner identity, item code, and negotiated price. The question is whether the cheap layer can rank where the costly layer should be opened first.

have been recovered. Award-layer triage asks whether the thinner record can rank where that recovery should start.

2.3. CADE Anchors and Cartel-Adjacency Labels

CADE adjudications supply the legal anchors, not the screen’s labels mechanically. The linked BEC portfolio contains 12 adjudicated procurement-cartel cases with 65 firm-defendants, of which 47 are active in BEC during the sample window. Final decisions span 2015–2025, with long process-to-ruling lags. That timing is central: if a cheap screen is useful, it is because it can move investigative attention before the completed adjudication record exists.

The adjudication record creates two objects with different roles. Direct CADE defendants are legal membership anchors: firms named in adjudicated procurement-cartel cases and matched to BEC participant records. They identify environments where legal process established cartel conduct. They are not, however, the natural target of a zero-win loser-side statistic. Direct defendants may occupy winner-side, allocation, or rotating roles in procurement cartels. A screen built from repeated losing should not be judged by whether it classifies those firms as if it were a generic cartel-membership detector.

The validation target is narrower. A cobidder is an always-loser firm that appears in at least one BEC tender-item with a direct CADE defendant. This definition yields 193 cartel-adjacent loser-side firms in the main validation target. Cobidder status is not legal membership and is not treated as a latent truth about agreement. It is an adjudication-anchored exposure label: the firm is observed on the losing side near a legally established cartel anchor. The label is therefore appropriate for evaluating forensic priority, not liability.

This distinction avoids two errors. The first would be to collapse defendants and cobidders into a single cartel class, overstating what award records can establish. The second would be to require a loser-side screen to recover direct defendants, misstating what the statistic is designed to rank. The relevant question is instead whether persistent zero-win participation concentrates the loser-side firms most plausibly worth bid-layer follow-up around adjudicated cartel environments.

Table 1 fixes the roles used throughout the paper. The outer BEC population defines where procurement participation is observed. The always-loser stratum defines the side of the bidding pool the award layer can rank. Frequent losers are the operational screen. Direct defendants are legal anchors. Cobidders are the loser-side adjacency target for validation.

Table 1: Data layers, populations, and legal-economic roles

Object	Size / layer	Role in the analysis
Award layer	Routine record	Participants, winners, item codes, prices, procuring units, years, and modalities; the cheap triage layer.
Bid layer	LANCES export	Per-bidder offers, timestamps, revisions, and bid moments; the costly forensic layer.
All BEC participants	41,444 firms	Outer firm universe for award-record construction.
Always-losers	16,843 firms	Zero-win stratum in which the award-layer screen ranks loser-side exposure.
Frequent losers	2,735 firms	Operational first-stage flag within always-losers; not a legal category.
Direct CADE defendants in BEC	47 firms	Legal anchors from adjudicated procurement-cartel cases.
Always-loser cobidders	193 firms	Firms losing alongside direct defendants; cartel-adjacent validation target, not legal membership.

The rest of the paper keeps these objects separate. The screen operates inside the always-loser stratum. The validation target is cobidder adjacency. Direct defendants anchor the environments but do not define the screen’s success. Bid-layer forensics remains the proof-producing stage. This separation lets the paper evaluate a legally limited triage architecture without turning suspicion into adjudication.

3. Award-Layer Triage

The award layer observes outcomes before it observes conduct. A losing bid may be sincere, poorly priced, capacity constrained, exploratory, strategic, or coordinated. The screen therefore uses what the award layer can measure at the triage stage: which firms appeared, which firms won, and which firms kept appearing without ever winning. The object is a ranking for forensic priority.

This section defines that ranking and the empirical implications that discipline it. Subsection 3.1 introduces persistent zero-win participation as the award-layer primitive. Subsection 3.2 explains why the continuous ranking can be converted into an operational queue or binary frequent-loser flag. Subsection 3.3 states the tests that the rest of the paper uses to evaluate whether the ranking is informative for enforcement.

3.1. *Persistent Zero-Win Participation*

Let T_i denote the number of BEC tender-items in which firm i participates, and let W_i denote the number it wins. The award-layer screen operates inside the zero-win stratum, $\{i : W_i = 0\}$, which contains 16,843 firms. Within that stratum, the continuous score is

$$s_i = \log(1 + T_i).$$

The log transformation stabilizes the scale but does not change the ordering. The economic object is the rank: among firms the award layer observes only as losers, which ones appear persistently enough to warrant bid-layer follow-up?

The maintained condition is monotonic. If some losing firms are cartel-deployed cover participants, they should appear more persistently in the zero-win stratum than ordinary losing firms, because their role is useful only through repeated deployment in cartel-targeted tenders. The appendix formalizes this condition as a monotone-likelihood-ratio ranking result. The main text uses it more modestly: persistent zero-win participation is a cheap statistic that can order suspicion under a transparent behavioral assumption and direct attention to the bid layer where conduct can be evaluated.

The zero-win restriction is therefore a scope restriction. Many firms never win for ordinary competitive reasons: poor fit, high costs, capacity constraints, market exploration, subcontracting positioning, or simple bad luck. The screen’s content comes from persistence inside that heterogeneous loser-side set. If the ranking has enforcement value, it should concentrate adjudication-anchored loser-side adjacency beyond what generic participation volume and market exposure would mechanically produce.

3.2. From Ranking to an Operational Queue

The continuous score is the central object because agencies can translate a rank into different investigative protocols. A team with fixed capacity could inspect the top k firms. A monitoring unit could refresh the rank each quarter and examine bunching below old cutoffs. A case team could use the score to request LANCES bid files for a survivor pool before applying bid-distribution forensics. These implementations differ, but all preserve the same sequencing logic: award records first, costly bid recovery second.

For an auditable binary implementation, we also report a frequent-loser flag. The rule marks zero-win firms whose participation count exceeds the median plus 1.5 times the interquartile range of the always-loser participation distribution. In BEC, the statistical cutoff is 13.5, implemented as $T_i > 14$, which flags 2,735 firms, or 16.2% of the always-loser stratum. The cutoff is fixed from the participation distribution alone, before any CADE label enters the analysis.

The cutoff is an administrative implementation of the rank. Two analysts applying the same rule to the same award records obtain the same first-stage list. Because a static threshold can be gamed, the empirical design treats the binary flag as one implementation among several: agencies can also use continuous ranks, top- k queues, refreshed thresholds, or budget-calibrated survivor pools. The strategic-adaptation analysis later returns to this point.

3.3. Testable Implications

The screen is useful if it survives the threats implied by its limited information. Table 2 organizes the empirical design. The first implication is the core validation claim: persistent zero-win participation should rank adjudication-anchored loser-side cobidders within the always-loser stratum. The main threat is mechanical exposure: frequent participants and firms in the same product markets as CADE defendants have more chances to appear near them even under ordinary competition.

The second implication is a scope check. A loser-side statistic should perform worse against direct CADE defendants than against always-loser cobidders, because it ranks a different role in the procurement environment. The third implication asks whether the

cobidder target has economic content: among frequent losers, cobidders should differ from other frequent losers in directions consistent with loser-side deployment, proximity to defendants, market concentration, and bid-level losing behavior.

The final implications evaluate the enforcement architecture. If award records and bid records carry different information, the award-layer score should add non-redundant signal to bid-distribution forensics. If the screen has institutional value, using it before bid recovery should reduce the bid-microdata footprint while preserving much of the adjudication-anchored recall. And because firms can respond to published administrative rules, the binary threshold should be treated as vulnerable to strategic adaptation.

Table 2: Testable implications and validation threats

Implication	Prediction	Test	Main threat
Loser-side ranking	Persistent zero-win participation ranks always-loser cobidders.	Conservative benchmark, timing audit, exposure-adjusted validation.	Participation volume and market overlap.
Scope asymmetry	Performance is weaker for direct defendants than for co-bidders.	Direct-defendant AUC and temporal direct-defendant check.	Misreading the screen as generic cartel detection.
Economic profile	Cobidders differ from other frequent losers in loser-side deployment and bid patterns.	Firm-level and bid-level profile comparisons.	Product-market composition.
Complementarity	Award-layer signal adds information to bid-layer forensics.	Same-target award, bid, and combined AUCs.	Full-observability benchmark confused with operation.
Gatekeeping	Award-first sequencing reduces bid-microdata recovery with limited recall loss.	Sequential survivor-pool exercise and cost-recall frontier.	Arbitrary budget choice.
Strategic response	Static cutoffs are more fragile than ranks or refreshed queues.	Adaptation diagnostics.	Simulation is not an equilibrium model.

Together, these tests evaluate the screen at the stage for which it is designed: the allocation of costly forensic attention before liability evidence is assembled. The BEC cutoff is an implementation detail; the portable object is the staged architecture that uses cheap administrative records to discipline where costly cartel forensics should begin.

4. Validation Against Cartel-Adjacent Losers

The frequent-loser ranking is useful for enforcement only if it orders the loser-side firms for which bid-layer follow-up is most valuable. The validation target is therefore the

one defined in Section 2.3: always-loser firms observed alongside direct CADE defendants in adjudicated procurement-cartel environments. These cobidders are not treated as legal members of a cartel. They are an adjudication-anchored exposure target for evaluating forensic priority.

The validation exercise has to do more than report a high AUC. The main empirical threat is exposure. Firms that participate often, or that operate in the same product markets as CADE defendants, have more opportunities to appear near those defendants even under ordinary competition. Procurement-screen designs therefore have to separate suspicious conduct from market structure, bidder heterogeneity, and ordinary repeated participation (Porter and Zona, 1993; Bajari and Ye, 2003). Network and large-scale procurement screens sharpen the same concern (Conley and Decarolis, 2016; Kawai and Nakabayashi, 2022). A useful award-layer screen must therefore concentrate cartel-adjacent loser-side firms beyond what participation volume and market exposure alone would imply. A second threat is timing: an enforcement screen should rank candidates before the target information is used. A third threat is legal overreach: a loser-side screen should not be read as a generic detector of cartel membership.

This section evaluates the ranking against those threats. Subsection 4.1 reports the baseline adjudication benchmarks. Subsection 4.2 disciplines the result with participation, exposure, leakage, and timing checks. Subsection 4.3 uses direct CADE defendants as a scope check. Subsection 4.4 summarizes what the validation establishes for the staged enforcement architecture.

4.1. Baseline Adjudication Benchmarks

The baseline asks a simple question before adding stricter discipline: do firms flagged by the award-layer rule appear disproportionately among always-loser firms located near adjudicated cartel anchors? The answer is yes. In the conservative adjudication benchmark, we restrict the CADE portfolio to cases resolved before the end of 2020. This leaves 4 cases, 30 BEC-active firm-defendants, and 210 always-loser cobidders. Among those cobidders, 108 are frequent losers.

This benchmark uses early legal anchors and asks whether the award-layer rank

concentrates the associated loser-side footprint. Within the always-loser stratum, the observed frequent-loser share among conservative cobidders is 3.95%, compared with a participation-stratified placebo mean of 1.24%. The excess ratio is 3.2 $\times$ with $p < 0.001$. The corresponding ROC AUC is 0.748 (95% DeLong CI [0.713, 0.783]) (DeLong et al., 1988). These numbers establish the baseline empirical fact: persistent zero-win participation is not random with respect to adjudication-anchored loser-side exposure.

The second benchmark adds timing discipline. The score is computed from 2009–2016 participation and evaluated against 2017–2019 cobidder labels. This asks whether earlier award records rank firms whose cartel adjacency is revealed later by the adjudication record. The firm-level AUC is 0.864 (95% CI [0.858, 0.870]), below the full-sample AUC 0.924 but well separated from random ranking. Figure 2 reports the rolling-origin version of the same exercise.

The baseline results are deliberately read as ranking evidence. They show that the award-layer statistic is informative about where cartel-adjacent loser-side exposure later appears. The next subsection tightens that reading by asking how much of the signal remains after disciplining participation volume, market exposure, leakage, and strict timing.




Figure 2: Temporal-holdout ROC curves. Scores are computed using participation observed before the test year; the test year supplies the out-of-time cobidder label.

4.2. Exposure and Timing Discipline

The baseline result becomes useful for enforcement when the main sources of mechanical exposure are made explicit. The first is participation volume. In the conservative benchmark, firms are grouped by participation quartiles within the always-loser stratum, cobidder status is permuted within quartiles 1,000 times, and the frequent-loser share among placebo cobidders is recomputed. This keeps the participation distribution fixed while breaking the particular adjacency to CADE defendants. The observed share, 3.95%, remains well above the placebo mean, 1.24%.

The second source is market opportunity. Firms that buy and sell in the same procurement niches as CADE defendants have more chances to appear near those anchors even when conduct is ordinary. The validation design therefore treats exposure adjustment as a required audit rather than as an optional robustness check. The audit compares firms within more similar opportunity sets, using item, modality, year, and buyer information where support permits. Its purpose is to ask whether the award-layer rank continues to concentrate adjudication-anchored cobidders after the markets in which firms could meet CADE anchors are made more comparable.

Timing is disciplined in two layers. The temporal holdout above forms the participation score before the evaluation window and tests later labels. A stricter version applies the same ex ante logic to the screen definition itself: zero-win status, the participation threshold, the frequent-loser flag, and the continuous score are all formed using only pre-window records. In that strict train-period audit, the threshold becomes 7 and the firm-level frequent-loser AUC is 0.767 (95% CI [0.734, 0.800]). The continuous score gives AUC 0.750 (95% CI [0.706, 0.795]). The decline relative to the relaxed holdout is informative: real-time screening has less information than a full-window reconstruction, but the ranking remains useful under a stricter field-like information set.

The leakage audit addresses a different concern: item-level validation can mechanically reuse the same firm histories that construct the score. We therefore treat the in-sample item-level AUC, 0.995, as a mechanical reference point rather than as the main validation estimate. Holding cobidder firms out of score construction within cross-validation folds leaves AUC 0.891 (95% CI [0.887, 0.894]). The surviving component is the signal relevant

for triage: the award-layer rank continues to order cartel-adjacent loser-side exposure when the most direct reuse of firm labels in score construction is weakened.

4.3. Direct-Defendant Scope Check

The direct-defendant exercise aligns the validation target with the screen’s legal-economic role. Direct CADE defendants are the legal anchors that locate adjudicated cartel environments. They need not be persistent zero-win firms: in procurement cartels, defendants may occupy winner-side, allocation, or rotating roles. The award-layer statistic, by construction, ranks the loser side of the procurement record. The empirical prediction is therefore asymmetry: the rank should be informative for cobidders and much less informative for direct defendants.

The data follow that prediction. Against direct defendants in the broader BEC firm universe, AUC is 0.49, close to random classification. Read together with the cobidder results, this pattern supports the staged interpretation. Direct defendants anchor cartel-affected environments; the award layer ranks the loser-side exposure around those anchors; the bid layer then supplies the conduct evidence for case-specific evaluation.

4.4. Validation Summary

Table 3 summarizes the validation architecture. Each row corresponds to a threat identified before the tests are read: legal anchoring, participation volume, exposure, timing, leakage, and scope. The resulting claim is staged: the award-layer ranking identifies a priority queue for the bid-layer stage, where tender-level conduct can be examined with richer evidence.

Taken together, the benchmarks support the screen at the stage for which it is designed. Persistent zero-win participation orders adjudication-anchored loser-side exposure well enough to allocate forensic attention. The next step is economic interpretation: among firms already selected by the frequent- loser rule, do cobidders look like a distinct loser-side type rather than a mechanical collection of high-volume losers?

Table 3: Validation benchmarks for the award-layer ranking

Benchmark	Question	Result
Conservative benchmark	Are frequent losers overrepresented among pre-2020 adjudication-anchored cobidders?	3.2× excess; AUC 0.748
Participation discipline	Does concentration survive volume-preserving placebo assignment?	Placebo mean 1.24% versus observed 3.95%
Exposure discipline	Does the ranking survive more comparable market opportunities to meet CADE anchors?	Exposure-adjusted audit; see Appendix D
Temporal holdout	Does pre-window participation rank later cobidder labels?	AUC 0.864
Leakage audit	Does signal survive when cobidders are held out of score construction?	Out-of-fold AUC 0.891
Direct-defendant scope check	Does the ranking preserve the distinction between legal anchors and loser-side exposure?	Direct-defendant AUC 0.49

5. The Economic Profile of Cartel-Adjacent Losers

The validation results show that the award-layer ranking concentrates adjudication-anchored loser-side exposure. This section asks whether that exposure target has economic content. The comparison is intentionally demanding: among firms already inside the frequent-loser stratum, do cobidders differ from other frequent losers in dimensions relevant for loser-side deployment?

The answer matters for interpretation. If cobidders look like other frequent losers once both groups cross the same participation threshold, the validation target would be easier to read as mechanical exposure. If they differ in deployment breadth, proximity to legal cartel anchors, market specialization, and bid-level losing behavior, the target becomes more meaningful for forensic prioritization. This profile-based reading is consistent with empirical work showing that collusion in auctions can appear through bidder roles, repeated relationships, and cartel organization rather than through a single universal price statistic. Studies of timber, school milk, and highway or procurement cartels show role allocation and repeated relationships in concrete bidding environments ([Baldwin et al., 1997](#); [Porter and Zona, 1999](#)); auction evidence extends the same logic to first-price collusion ([Pesendorfer, 2000](#)). Case-based and dynamic analyses further show how internal

cartel organization shapes observed bidding behavior (Asker, 2010; Clark et al., 2021). The evidence below supports the latter interpretation.

5.1. Comparing Cobidders to Other Frequent Losers

The firm-level comparison holds fixed the award-layer selection margin. Both groups are always-losers, and both cross the frequent-loser threshold. Cobidders are frequent losers observed with direct CADE defendants; the comparison group is frequent losers without that adjacency. The question is whether the adjudication-anchored target marks a different loser-side profile among firms already selected by persistent losing.

The differences line up with the triage interpretation. Cobidders are deployed more broadly: they participate in 136.5 tender-items, compared with 76.7 for other frequent losers ($d = +0.67$). They also face more unique winners, 24.8 versus 13.5 ($d = +1.00$). These are deployment measures, not legal conclusions: they show that cobidders occupy a larger loser-side footprint within the already-selected frequent-loser pool.

They are also closer to the legal anchors used in validation. Among cobidders, 1.5% of loss-bids face a direct CADE defendant, compared with 0.2% for other frequent losers ($d = +0.46$). Their portfolios are more concentrated in product space as well, with item-group HHI 0.380 against 0.288 ($d = +0.39$). The combination of broader deployment and narrower product concentration is exactly the kind of profile that makes the target useful for forensic prioritization.

The profile is not one-dimensional. Cobidders have a lower share of own loss-bids in repeat winner-pair clusters than other frequent losers, 21.6% versus 33.4% ($d = -0.38$). That wrong-signed measure is informative: the cobidder label is not a mechanical proxy for a single winner-pair pattern. The stronger regularities are deployment breadth, proximity to CADE anchors, and product specialization within the frequent-loser stratum.

5.2. Bid-Level Behavior

The LANCES export lets us ask whether the distinction also appears inside submitted bids. This evidence uses the forensic layer; it is deliberately kept outside the award-layer screen. For firms with at least five usable losing bids in convite or pregão, we compute

the gap between each losing bid and the winning bid in the same tender-item,

$$\frac{\text{bid} - \text{winning bid}}{\text{winning bid}},$$

winsorized at $[-0.99, 10]$ before firm-level aggregation. The usable sample contains 182 cobidders and 2,369 frequent-loser non-cobidders.

Two patterns matter. First, cobidders bid closer to winners: their median bid gap is 0.582, compared with 0.809 for other frequent losers ($d = -0.281$). In a cover-bidding environment, losing bids need not be visibly extreme; plausible losing bids can preserve the appearance of rivalry. Second, cobidders show higher within-firm dispersion in bid gaps, 1.207 versus 1.099 ($d = +0.147$), consistent with firms appearing in different losing roles across tenders.

The bid-level distinction is not just participation intensity in another form. In a logit of cobidder status within the frequent-loser stratum, controlling for $\log(1 + \text{tenders_count})$, both signs remain: the median-gap coefficient is -0.708 ($z = -4.14$), and the dispersion coefficient is $+0.442$ ($z = +3.62$). The forensic layer therefore gives independent economic content to the adjudication-anchored loser-side target.

Table 4: Cobidder profile within the frequent-loser stratum

Dimension	Cobidders	FL non-cobidders	Difference
Total participations	136.5	76.7	$d = +0.67$
Unique winners faced	24.8	13.5	$d = +1.00$
Loss-bids facing direct CADE defendant	1.5%	0.2%	$d = +0.46$
Portfolio item-group HHI	0.380	0.288	$d = +0.39$
Repeat winner-pair share	21.6%	33.4%	$d = -0.38$
Median bid gap to winner	0.582	0.809	$d = -0.281$
Within-firm SD of bid gap	1.207	1.099	$d = +0.147$

5.3. Interpretation

The profile gives economic content to the validation target. Within the frequent-loser stratum, cobidders are more broadly deployed, closer to legal cartel anchors, more concentrated in product space, closer to winners in bid gaps, and more dispersed across

their own losing bids. These are precisely the dimensions on which a forensic queue should differ if the award-layer rank is capturing more than generic losing volume.

This evidence also clarifies the division of labor between data layers. The award layer identifies a priority set using information available before bid recovery. The LANCES bid layer then shows that the priority set has economically meaningful loser-side behavior. That is the paper’s staged logic in miniature: cheap records order attention, and richer records evaluate what the screen has prioritized.

The next section turns from interpretation to implementation. It asks whether the award-layer rank can be placed before bid-distribution forensics in an operational sequence that saves bid-microdata recovery while preserving adjudication-anchored recall.

6. From Award-Layer Triage to Bid-Layer Forensics

The preceding sections validate the award-layer ranking on the margin for which it is designed. Section 4 shows that persistent zero-win participation orders adjudication-anchored loser-side exposure; Section 5 shows that the exposed firms have an economically distinctive profile rather than merely high procurement volume. The remaining question is how an enforcement agency should use such a signal. The answer cannot be to stop at the award layer. A thin administrative record can rank priorities, but it cannot by itself evaluate the bid conduct that richer cartel evidence requires.

This section therefore turns from validation to sequencing. The screen is useful because it sits before costly proof-producing work. It is computed from records the agency can observe broadly: who participated, who won, and who kept appearing without winning. Bid-layer forensics, by contrast, uses submitted offers, within-tender bid moments, and distance from the winning bid. Those data are more probative but also more expensive to recover, clean, harmonize, and inspect at scale. The institutional problem is not whether a cheap signal should replace a costly one; it is whether the cheap signal can discipline where the costly one should be opened first.

The section evaluates that architecture in three steps. First, it places the award-layer ranking and an Imhof–Wallimann-style bid-distribution benchmark on the same adjudication-anchored target, so that differences in performance can be read against

differences in information cost. Second, it asks whether the two layers are complements: the award layer measures repeated loser-side exposure across tenders, while the bid layer measures conduct inside tenders after the bid file is opened. Third, it implements the operational gatekeeper. The award layer first selects a survivor pool; only then does the bid layer rank firms inside that pool. The object is a cost-recall frontier for forensic attention, not a horse race among detectors under full observability.

This sequencing is the paper’s central enforcement implication. It turns a legally limited screen into a practical allocation rule: administrative records identify where proof-producing forensics should begin, and bid-level evidence remains the stage at which conduct is evaluated. The value of the method is therefore measured by how much forensic scope it can reduce while preserving adjudication-anchored signal.

6.1. Bid-Layer Benchmark Under Costly Observability

The natural forensic comparator is the bid-distribution tradition. Imhof–Wallimann-style detectors use within-tender bid moments—dispersion, skewness, spread, distance from the winning bid, and related features—to flag patterns consistent with coordinated bidding. This family of screens begins with simple distributional abnormalities ([Abrantes-Metz et al., 2006](#)). Procurement applications develop descriptive bid-rigging indicators ([Imhof et al., 2018](#); [Imhof, 2019](#)) and extend to machine-learning implementations that combine multiple bid moments ([Huber and Imhof, 2019](#); [Wallimann et al., 2023](#)). These features are closer to conduct than the award-layer screen is. They use information from the bids themselves rather than only the administrative trace of participation and awards. For that reason, they are also a more costly evidentiary object: the agency must recover usable bid histories, harmonize submitted offers, construct within-tender moments, and map those moments back to firms.

The relevant comparison is therefore not sophistication. The bid-layer benchmark is more forensic by construction. The question is how much ranking information is already available before the bid file is opened. Table 5 places the award-layer measures and the bid-layer benchmark on the same adjudication-anchored target: CADE-cobidder labels within the always-loser pool for which bid-distribution features can be computed. Holding

the target fixed is important. It means that the table varies the information envelope rather than the legal-economic object being ranked.

The award-layer signal performs strongly despite its thinner information set. Within this common pool ($N = 16,779$), the binary frequent-loser flag reaches AUC 0.903 (95% CI [0.884, 0.923]), and the continuous participation score reaches AUC 0.884 ([0.860, 0.908]). The full bid-layer benchmark reaches AUC 0.888 ([0.865, 0.911]). When the two layers are combined, discrimination rises to 0.962 ([0.954, 0.969]).

These results should be read as an information-cost result. The award-layer statistic does not make bid-distribution forensics unnecessary, and the paper does not ask it to. Rather, a cheap record-layer statistic discriminates the same adjudication-linked loser-side target in the range of a richer bid-layer benchmark. That is enough to justify the sequencing question: if the award layer contains useful information before bid recovery, and if it adds information to the bid layer rather than merely duplicating it, then it can be used to decide where bid-layer forensics should begin.

Table 5: Discrimination at Different Information Costs

Model	Information required	AUC	95% CI
Frequent-loser flag	Award records: participation count and zero wins	0.903	[0.884, 0.923]
Participation intensity	Award records: $\log(1 + \text{tenders_count})$ and zero wins	0.884	[0.860, 0.908]
Bid-layer benchmark	Per-bid offer values and within-tender bid moments	0.888	[0.865, 0.911]
Award layer + bid layer	Participation intensity plus bid-distribution moments	0.962	[0.954, 0.969]

Notes: Five-fold cross-validated AUC against CADE-cobidder labels within the always-loser pool for which the bid-distribution benchmark can be computed ($N = 16,779$). The combined specification uses the same target and adds the award-layer score to the bid-layer moments.

6.2. Complementarity Across Information Layers

Similar discrimination does not by itself establish a useful two-stage architecture. If persistent zero-win participation were only a noisy proxy for bid-distribution moments, the award layer would matter mainly when bid microdata are missing. The stronger test is incremental: once bid-layer moments are observed, does the award layer still contribute information about the adjudication-anchored loser-side target?

It does. Adding the award-layer statistic to the bid-layer benchmark raises discrimination to 0.962 ([0.954, 0.969]). In the identical same-sample audit, where all models are evaluated only on firms with complete bid-layer features, the frequent-loser flag adds +0.035 AUC over the full Imhof–Wallimann benchmark alone (DeLong $p = 0.014$) (DeLong et al., 1988). The combined specifications gain +0.096 – – + 0.098 AUC over the bid-layer benchmark ($p < 0.001$). The increment is not needed to claim that the award layer dominates the bid layer; it shows that the two layers are not measuring the same thing.

The distinction is substantive. Bid moments describe conduct inside tenders after the bid file has been opened: spread, dispersion, rank distance, and related within-tender patterns. Participation intensity describes repeated loser-side exposure across tenders before that file is opened. The award layer is therefore closer to investigative priority; the bid layer is closer to forensic evaluation. An enforcement architecture should preserve that division of labor. The award layer creates a reason to look first in some places rather than others. The bid layer then helps evaluate what is found after looking.

The combined model in Table 5 is thus a diagnostic for complementarity, not the operational baseline. Running the combined model for every firm assumes full bid-layer observability before triage has occurred, which is precisely the resource problem the paper studies. The enforcement use of non-redundancy is sequential: use broad, cheap award records to create a survivor pool, then apply costly bid-distribution forensics within that pool. The next subsection implements that gatekeeper directly.

6.3. Gatekeeping the Forensic Stage

The operational question is whether the cheap layer can reduce the forensic workload while preserving much of the adjudication-anchored signal. This is a different estimand from stand-alone classification performance. A rule that requires bid-level features for every firm may score well, but it has already paid the full cost of forensic observability. A gatekeeper is useful only if it changes the size of the set for which that cost must be paid.

The sequential rule imposes that institutional constraint directly. First, firms are ranked using only the award layer, and the top $K_1 = 2,000$ firms form the survivor

pool. Second, bid microdata are recovered or processed only for those survivors. The bid-distribution benchmark is then applied within that smaller pool to select the final top- k priorities. The joint model, by contrast, gives the classifier both award-layer and bid-layer features for every firm before triage. It is an upper-bound benchmark under full observability, not an implementable first-stage policy when bid recovery is costly.

Table 6 reports the resulting cost-recall frontier. The evaluation pool contains 11,676 always-loser firms with computable bid-distribution features and 193 CADE-cobidder positives. At the top 500 flags, award-layer screening alone finds 66 positives with precision 0.132. The bid-layer benchmark alone finds 73 positives with precision 0.146, but it requires bid microdata for the entire pool. Joint scoring finds 112 positives with precision 0.224, again under full bid-layer observability. The sequential rule finds 96 positives with precision 0.192 while opening bid microdata for only 2,000 firms.

The same trade-off appears at the top 1,000 cutoff. Joint scoring recovers 142 positives, or recall 0.736, after observing bid-layer features for all 11,676 firms. The sequential rule recovers 131 positives, or recall 0.679, while requiring bid microdata for only 2,000 firms. Relative to joint scoring, the sequential architecture reduces the bid-microdata footprint by 83% at a recall cost of 8%.

Table 6: Gatekeeping the Bid-Layer Forensic Stage

Rule	k	TP	Precision	Recall	Bid-microdata firms
Award-layer only	500	66	0.132	0.342	0
Bid-layer only	500	73	0.146	0.378	11,676
Joint scoring, award + bid	500	112	0.224	0.580	11,676
Sequential: award $\rightarrow$ bid, $K_1 =$ 2,000	500	96	0.192	0.497	2,000
Joint scoring, award + bid	1,000	142	0.142	0.736	11,676
Sequential: award $\rightarrow$ bid, $K_1 =$ 2,000	1,000	131	0.131	0.679	2,000

Notes: Evaluation pool: 11,676 always-loser firms with computable bid-distribution features; positives are 193 CADE-cobidders. The bid-microdata column reports how many firms require per-bid offer values before the rule can be applied.

The sequential rule is not meant to dominate the full-information joint model. Its value is that it approaches much of the full-information signal without requiring full-information

forensic scope. That is the enforcement margin on which the paper’s contribution rests. In an agency setting, scarce analyst time, data cleaning, subpoena capacity, and case-development attention cannot be applied with equal intensity to every participant. A useful screen does not decide liability; it decides which files deserve the next, more expensive form of evidentiary attention.

6.4. Timing the Gatekeeper

The gatekeeper exercise also needs a timing discipline. A ranking that uses the full sample may perform well because the feature history is long and close to the adjudication target. A stricter audit asks whether the same sequencing logic remains useful when the scores are formed before the target window. This is not a field deployment in real time; the labels remain adjudication-based. The restriction is on the information used to rank firms before bid-layer forensics is allocated.

The temporal audit constructs award-layer and bid-layer features using the 2009–2016 window and evaluates against later adjudication-anchored cobidder exposure. The holdout pool contains 8,257 always-loser firms with complete train-window bid features and 142 positives. This design asks whether the staged architecture survives when the ranking layer is separated in time from the validation target.

It does. At the top 500 cutoff, the award-only rule loses 47% of its in-sample true positives, joint scoring loses 24%, and the sequential rule loses only 9%. In levels, the sequential rule recovers 87 positives at top 500, compared with 85 under joint scoring. At top 1,000, it recovers 114 positives, compared with 111 under joint scoring. The corresponding precisions are 0.174 and 0.114 for the sequential rule.

The temporal result reinforces the sequencing interpretation rather than expanding it. It does not show that a fixed cutoff is universally optimal, nor that an agency can infer agreement from the award layer alone. It shows that conditioning costly bid-layer forensics on an earlier award-layer screen remains competitive under timing discipline while preserving the smaller forensic footprint. The conclusion of Section 6 is therefore operational: cheap administrative records can structure the order in which expensive proof-producing evidence is assembled, while the richer bid layer remains the place where

conduct is evaluated.

7. Scope, Limits, and Price Corroboration

The paper’s main claim does not require a causal price effect. The screen is designed to allocate costly proof-producing attention, not to estimate cartel markups, damages, or welfare losses. Price evidence is still useful, but it plays a narrower role. It asks whether the environments prioritized by the award-layer screen look economically non-neutral, and it helps define where that interpretation must stop. Cartel detection, liability, and overcharge measurement are related but distinct tasks: overcharge measurement has its own evidentiary tradition ([Bryant and Eckard, 1991](#); [Connor, 2007](#)), while cartel economics and detection treat screening as a separate enforcement problem ([Marshall and Marx, 2012](#); [Harrington, 2008](#)).

This distinction matters for the paper’s legal-economic position. A positive association between frequent-loser presence and negotiated prices would be consistent with the idea that the screen points toward economically meaningful procurement environments. But it would not establish agreement, cover bidding, antitrust liability, or a damages base. Conversely, heterogeneous or even sign-reversing price estimates would not invalidate the screen if the screen’s proper object is forensic priority rather than price estimation. The price record therefore disciplines the interpretation instead of carrying the paper’s identification burden.

The section uses price results in that limited way. It first reports the broad descriptive price imprint associated with frequent-loser presence. It then explains why overlap and segment results narrow the scope of that imprint, especially where the sign reverses outside high-value procurement. The final implication is institutional: the award-layer signal is best used as a trigger for richer evidence collection, not as a stand-alone measure of harm. That boundary is what makes the screening claim stronger for enforcement design rather than weaker.

7.1. Price Evidence as Scope

The price regressions are reported to assess economic non-neutrality, not to estimate harm. If the award-layer screen ranks environments where cartel-adjacent loser-side participation is more likely, it is informative to ask whether those environments also differ in negotiated unit prices. The answer can only play a corroborative role. The estimates below are descriptive associations, not cartel markups, causal effects of frequent-loser presence, or inputs into a damages calculation.

In the broad item-level regressions, frequent-loser presence is positively associated with log negotiated unit price. The coefficient is 0.0677*** with item and year fixed effects and 0.0636*** after adding procuring-unit fixed effects. The association is larger in pregão, 0.0933***, than in convite, 0.0382**. Across the main descriptive specifications, the implied price range is +3.6% – +7.7%.

The modality split is useful because it disciplines the screening interpretation. If frequent losers were only mechanical quorum fillers under the three-bidder rule, the price imprint should be stronger in convite. Instead, the larger association appears in pregão, the electronic-auction environment where repeated loser-side positioning can be deployed without the same formal invitation structure. This contrast supports the view that the screen is not merely picking up a procedural artifact. It still does not identify a causal procurement-price effect.

Table 7: Descriptive Price Imprint

Exercise	Estimate	Standard error	Use in the paper
Broad sample, item and year FE	0.0677***	0.0230	Economic non-neutrality
Broad sample, item, year, PBU FE	0.0636***	0.0215	Positive after buyer controls
Pregão subsample	0.0933***	0.0255	Stronger in electronic auctions
Convite subsample	0.0382**	0.0186	Smaller under quorum-prone procedure

Notes: Dependent variable is log negotiated unit price. Estimates are descriptive associations in the BEC item-level sample and are not interpreted as causal price effects, cartel markups, or damages estimates.

The price imprint therefore corroborates the economic relevance of the priority set, but it does not carry the paper’s central claim. The central claim remains about allocating the next unit of proof-producing effort. The next subsection uses the overlap and segment results to make that boundary explicit.

7.2. *Why the Sign Reversal Matters*

The overlap results define where the price evidence stops. The broad-sample estimate averages over the environments into which frequent losers actually sort. The overlap exercises instead restrict or reweight toward cells where frequent-loser-present and frequent-loser-absent items have comparable support on observables. These are different descriptive objects. Neither is a causal effect, and the difference between them is informative: it shows why the paper should not convert the screen into a general claim about price effects. The distinction follows the standard separation between descriptive comparisons and causal estimands (Imbens and Rubin, 2015); coefficient stability and sensitivity logic reinforce the same caution (Oster, 2019; Cinelli and Hazlett, 2020).

Under overlap-cell ATT reweighting, the aggregate coefficient reverses to -0.097 ; under propensity-score trimming it becomes -0.307 . Only 1.06% of treated items lack a within-cell counterfactual, so the reversal is not mainly a dropped-observation result. It is a weighting result. The broad price imprint and the overlap estimate put weight on different parts of the procurement distribution, and those parts carry different economic content.

The segment decomposition makes the boundary sharper. In the lowest three tender-value quartiles, the overlap estimates are negative. In the highest tender-value quartile, the estimate remains positive, $+0.041$ ($p = 0.045$). That is the segment where a cartel interpretation is most plausible ex ante because fixed coordination costs are spread over larger contracts. Even there, the result is scope evidence, not a markup estimate. It tells the reader where the price imprint is economically concentrated; it does not identify harm.

The direct-CADE overlap estimate is also null, -0.061 ($p = 0.45$), which reinforces the same interpretation. Price regressions do not recover the legal object of the case, and they do not supply a damages base. They show that the award-layer screen points to economically non-neutral environments in parts of the distribution, while the overlap reversal prevents price evidence from bearing more weight than it can support.

Table 8: Where Price Evidence Stops

Exercise	Estimate	p / SE	Interpretation
Overlap-cell ATT	-0.097	$p < 0.001$	Aggregate sign reversal under common support
PS-trimmed ATT	-0.307	0.020	Stronger negative estimate after trimming
Tender-value Q4, overlap ATT	+0.041	$p = 0.045$	Positive only in high-value segment
Direct-CADE overlap items	-0.061	$p = 0.45$	No separate price effect for legal anchors

Notes: Dependent variable is log negotiated unit price. The overlap rows restrict or reweight comparisons toward cells with treated and untreated support. Estimates are used to discipline the scope of price interpretation, not to estimate causal effects, overcharges, or damages.

7.3. Implications for Legal-Economic Use

The legal-economic use of the screen follows from these boundaries. The price evidence is consistent with economic relevance, but it does not transform an award-layer statistic into proof of agreement or harm. The correct institutional use is staged: award records rank where costly bid-level and documentary evidence should be assembled; bid records, case documents, and legal process then determine whether the conduct supports an antitrust case.

This distinction also clarifies portability. The BEC estimates should not be exported as numerical external validity for other jurisdictions. What travels is the information problem. Many procurement systems observe participation and awards more cheaply than they can recover, clean, and evaluate bid histories. In such settings, a useful screen requires stable participant identities, observable winners and losers, repeated procurement, sufficiently comparable items, and a feasible path from the initial ranking to richer forensic follow-up. Without that path to proof, a cheap signal has little enforcement value.

7.4. Adaptive Deployment

The screen should also be deployed as an adaptive queue rather than as a static legal trigger. Any administrative rule can generate strategic response once firms learn which observable patterns attract attention. In this setting, firms could rotate loser roles, engineer occasional wins, fragment identities, or avoid a known participation threshold.

That possibility does not undermine the sequencing architecture, but it does rule out a naive use of the binary cutoff as a permanent enforcement boundary. Strategic response is a standard concern for cartel enforcement and screening because firms may adapt to observable enforcement rules once those rules are understood ([Green and Porter, 1984](#); [Haltiwanger and Harrington, 1991](#)). Screening rules face the same adaptation problem in enforcement settings ([Harrington, 2008](#); [Sánchez Graells, 2019](#)).

The implementation implication is straightforward. Agencies should use continuous ranks or capacity-constrained top- k queues, refresh thresholds, monitor bunching below operational cutoffs, and combine award-layer screening with bid-layer forensics whenever bid files can be recovered. A firm that moves to avoid one observable pattern may create another: discontinuous participation, suspicious identity changes, or shifts in bid-level behavior. The architecture is therefore strongest when the award-layer queue is periodically updated and linked to forensic follow-up rather than treated as a one-time blacklist.

This is also the right legal posture. The screen is not a rule of liability and should not be used as one. It is a way to order scarce investigative attention when proof is costly and the first observable record is legally thin. Used adaptively, it gives an agency a disciplined place to begin without pretending that the beginning is the end of the evidentiary process.

8. Conclusion

Cartel enforcement often begins before the agency has the evidence it would need to prove a case. Procurement records then pose a problem that is prior to liability: how should scarce investigative capacity be allocated when the cheap administrative record is broad but legally thin, and the richer bid-level record is costly to recover and evaluate? This paper answers that question by treating cartel screening as evidence allocation under costly proof.

The empirical design separates the award layer from the bid layer. The award layer observes participation and outcomes; it can rank persistent zero-win participation before bid histories are opened. The bid layer observes submitted offers and within-tender

patterns; it is the natural place for proof-producing forensic evaluation. The contribution is not to collapse these layers. It is to show how they can be sequenced: administrative records identify where investigation should begin, and bid-level evidence then evaluates conduct inside the prioritized set.

The results support that architecture. Persistent zero-win participation is enriched among adjudication-anchored cartel-adjacent losers, survives timing and leakage discipline, and does not operate as a generic direct-defendant detector. The cobidder profile gives the target economic content: within the frequent-loser stratum, cartel-adjacent firms are more broadly deployed, closer to legal anchors, more product-concentrated, and distinct in bid-level behavior. The bid-distribution exercises then show complementarity rather than substitution. Award-layer ranking adds information to bid-layer moments, and a sequential gatekeeper preserves much of the adjudication signal while reducing the bid-microdata footprint by 83%.

The price evidence reinforces the same legal-economic posture. Flagged environments have a positive broad price imprint, but overlap and segment results prevent that imprint from becoming a damages or overcharge claim. The screen is valuable because it orders forensic priority, not because it decides harm. Its proper use is institutional and staged: rank where costly evidence should be assembled, then let bid records, documents, and legal process determine what the evidence proves.

That staged use is the broader lesson. Agencies routinely face a gap between what they can observe cheaply and what they must know to prove coordination. Waiting for full proof before prioritizing investigation wastes the cheap information already in the record; treating cheap suspicion as proof overstates what the record can bear. The approach in this paper occupies the middle ground. It turns a thin administrative signal into a disciplined queue for costly legal-evidentiary work.

References

Abrantes-Metz, R. M., L. M. Froeb, J. F. Geweke, and C. T. Taylor (2006). A variance screen for collusion. *International Journal of Industrial Organization* 24(3), 467–486.

- Asker, J. (2010). A study of the internal organization of a bidding cartel. *American Economic Review* 100(3), 724–762.
- Athey, S., J. Levin, and E. Seira (2011). Comparing open and sealed bid auctions: Evidence from timber auctions. *Quarterly Journal of Economics* 126(1), 207–257.
- Bajari, P. and L. Ye (2003). Deciding between competition and collusion. *Review of Economics and Statistics* 85(4), 971–989.
- Baker, J. B. (2003). The case for antitrust enforcement. *Journal of Economic Perspectives* 17(4), 27–50.
- Baldwin, L. H., R. C. Marshall, and J.-F. Richard (1997). Bidder collusion at forest service timber sales. *Journal of Political Economy* 105(4), 657–699.
- Becker, G. S. (1968). Crime and punishment: An economic approach. *Journal of Political Economy* 76(2), 169–217.
- Bryant, P. G. and E. W. Eckard (1991). A price-fixing probability index. *Review of Economics and Statistics* 73(4), 749–752.
- Chassang, S. and J. Ortner (2022). Robust screens for non-competitive bidding in procurement auctions. *Econometrica* 90(1), 315–346.
- Cinelli, C. and C. Hazlett (2020). Making sense of sensitivity: Extending omitted variable bias. *Journal of the Royal Statistical Society: Series B* 82(1), 39–67.
- Clark, R., J.-F. Houde, and J. Kastl (2021). The dynamics of bidder collusion: Evidence from a long-lived cartel. *Journal of Political Economy* 129(8), 2353–2391.
- Conley, T. G. and F. Decarolis (2016). Detecting bidder groups in collusive auctions. *American Economic Journal: Microeconomics* 8(2), 1–38.
- Connor, J. M. (2007). Price-fixing overcharges: Legal and economic evidence. In *Research in Law and Economics*, Volume 22, pp. 59–153. Elsevier.
- Coviello, D. and S. Gagliarducci (2017). Tenure in office and public procurement. *American Economic Journal: Economic Policy* 9(3), 59–105.
- Decarolis, F. (2014). Awarding price, contract performance, and bids screening: Evidence from procurement auctions. *American Economic Journal: Applied Economics* 6(1), 108–132.
- Decarolis, F., R. Fisman, P. Pinotti, and S. Vannutelli (2025). Rules, discretion, and corruption in procurement: Evidence from italian government contracting. *Journal of Political Economy Microeconomics* 3(2), 213–254.

- Decarolis, F., L. M. Giuffrida, E. Iossa, V. Mollisi, and G. Spagnolo (2020). Bureaucratic competence and procurement outcomes. *Journal of Law, Economics, and Organization* 36(3), 537–597.
- DeLong, E. R., D. M. DeLong, and D. L. Clarke-Pearson (1988). Comparing the areas under two or more correlated receiver operating characteristic curves: a nonparametric approach. *Biometrics* 44(3), 837–845.
- Green, E. J. and R. H. Porter (1984). Noncooperative collusion under imperfect price information. *Econometrica* 52(1), 87–100.
- Haltiwanger, J. and J. E. Harrington (1991). The impact of cyclical demand movements on collusive behavior. *RAND Journal of Economics* 22(1), 89–106.
- Harrington, J. E. (2008). Detecting cartels. In P. Buccirossi (Ed.), *Handbook of Antitrust Economics*, pp. 213–258. MIT Press.
- Hovenkamp, H. (2005). *Federal Antitrust Policy: The Law of Competition and Its Practice*. Thomson West.
- Huber, M. and D. Imhof (2019). Machine learning with screens for detecting bid-rigging cartels. *International Journal of Industrial Organization* 65, 277–301.
- Imbens, G. W. and D. B. Rubin (2015). *Causal Inference for Statistics, Social, and Biomedical Sciences*. Cambridge University Press.
- Imhof, D. (2019). Detecting bid-rigging cartels with descriptive statistics. *Journal of Competition Law & Economics* 15(4), 427–467.
- Imhof, D., Y. Karagök, and S. Rutz (2018). Screening for bid rigging—does it work? *Journal of Competition Law & Economics* 14(2), 235–261.
- Kawai, K. and J. Nakabayashi (2022). Detecting large-scale collusion in procurement auctions. *Journal of Political Economy* 130(5), 1585–1629.
- Marshall, R. C. and L. M. Marx (2012). *The Economics of Collusion: Cartels and Bidding Rings*. Cambridge, MA: MIT Press.
- Oster, E. (2019). Unobservable selection and coefficient stability: Theory and evidence. *Journal of Business & Economic Statistics* 37(2), 187–204.
- Pesendorfer, M. (2000). A study of collusion in first-price auctions. *Review of Economic Studies* 67(3), 381–411.
- Porter, R. H. and J. D. Zona (1993). Detection of bid rigging in procurement auctions. *Journal*

- of Political Economy* 101(3), 518–538.
- Porter, R. H. and J. D. Zona (1999). Ohio school milk markets: An analysis of bidding. *RAND Journal of Economics* 30(2), 263–288.
- Sánchez Graells, A. (2019). ‘screening for cartels’ in public procurement: Cheating at solitaire to sell fool’s gold? *Journal of European Competition Law & Practice* 10(4), 199–211.
- Stigler, G. J. (1970). The optimum enforcement of laws. *Journal of Political Economy* 78(3), 526–536.
- Wallimann, H., D. Imhof, and M. Huber (2023). A machine learning approach for flagging incomplete bid-rigging cartels. *Computational Economics* 62(4), 1669–1720.
- Wils, W. P. J. (2007). Leniency in antitrust enforcement: Theory and practice. *World Competition* 30(1), 25–64.
- Wils, W. P. J. (2016). The use of leniency in EU cartel enforcement: An assessment after twenty years. *World Competition* 39(3), 327–388.